


# Three different things share one valence axis — including a classifier that was never taught valence

each row places the nine emotions where that source puts them; thin lines mark the language model's positions



## Re-teaching the classifier the axis hurts ( $-0.029$ , $p < .01$ ); averaging 10 checkpoints helps

balanced accuracy on FACED-9 (the 9-emotion EEG benchmark above, 123 people)



*Nothing left to teach — the axis is already there. Averaging seeds removes only what still differs between them.*