Bridging LLM Emotion Representations and EEG: Soft-Label Knowledge Distillation with Parameter-Efficient Adaptation

Anonymous Author(s)
Affiliation
Address
email

Abstract

Recent interpretability research reveals that large language models develop internal emotion representations organized along the same valence–arousal dimensions found in human psychology. We introduce **EMOTIONKD**, which exploits this brain–LLM alignment: we extract emotion vectors from instruction-tuned LLMs and use their relational geometry as soft-label knowledge distillation targets for EEG emotion classification, teaching the model that *anger is closer to disgust than to joy* rather than treating all misclassifications equally. On FACED (9-class, 123 subjects), EMOTIONKD achieves $\mathbf{0.628}$ balanced accuracy, improving over the previous state of the art (REVE, NeurIPS 2025) by $+\mathbf{11.2}\%$ relative, with cross-subject generalization confirmed at $p < 0.001$. We validate the framework along three orthogonal axes: (i) *cross-backbone*, where KD significantly improves both CBraMod ($+\mathbf{5.2}\%$, $p < 10^{-4}$) and LaBraM ($+\mathbf{2.6}\%$, $p = 0.025$); (ii) *cross-dataset*, where with sufficient temporal granularity KD also significantly improves SEED-V on LaBraM ($+\mathbf{1.85}\%$, $p = 0.037$); and (iii) *cross-LLM-family*, where 11 instruction-tuned models from Qwen2.5 and Gemma-4 produce nearly identical emotion geometries (within-family RSA $\rho \approx 0.89$, cross-family $\rho = 0.82$), with linear cluster purity reaching 100% for positive vs. negative valence. Through Centered Kernel Alignment (CKA), we show that KD increases the similarity between EEG features and *every* LLM in our suite, including LLMs from a different family from the training teacher, indicating that KD absorbs a universal emotion structure rather than overfitting to one specific LLM. A key architectural finding is that freezing the pretrained backbone and adding lightweight adapters (310K parameters) matches full fine-tuning (4.9M parameters): the critical factor is the *quality of supervision*, not model capacity. We release our code, emotion vectors, and pretrained adapters to advance brain–language alignment research for EEG.

1 Introduction

How do the brain and large language models represent emotions—and can their shared structure be exploited? Emotions in human psychology organize along continuous dimensions of valence and arousal [1], a geometry that neuroscience has shown is reflected in distributed neural representations. Recent interpretability research reveals a striking parallel: LLMs develop internal emotion representations with the same dimensional structure, emerging naturally from language modeling without explicit emotional supervision [2, 3, 4]. This convergence suggests a deep alignment between how biological and artificial neural networks encode affect—an alignment we propose to exploit for EEG-based emotion recognition. While EEG foundation models such as CBraMod [5], LaBraM [6], and REVE [7] have established strong baselines on emotion benchmarks like FACED [8], they treat

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

emotion categories as integer indices, discarding the relational structure that both brains and LLMs naturally encode.

The standard cross-entropy loss penalizes all misclassifications equally—confusing *joy* with *amusement* costs the same as confusing *joy* with *anger*. Recent work has begun to address this: EmotionCLIP [9] reformulates EEG emotion recognition as EEG-text matching using CLIP embeddings of emotion labels, while EMOD [10] projects labels into a continuous valence–arousal space with soft-weighted contrastive learning. However, these approaches rely on hand-crafted text templates or predefined coordinate systems. We propose instead to use the *internal* emotion geometry that emerges within LLMs—a richer, data-driven representation that encodes not just valence–arousal position but the full relational structure among emotions, including within-valence distinctions (e.g., *anger* vs. *fear* vs. *disgust*).

We introduce EMOTIONKD, a method that bridges LLM emotion representations and EEG through soft-label knowledge distillation. Given any instruction-tuned LLM, we extract emotion vectors by prompting the model to generate emotion-laden stories and recording hidden state activations at intermediate layers. These vectors define a similarity matrix that serves as soft targets for a KL-divergence-based KD loss, replacing the standard integer-label cross-entropy with a distribution that encodes *how similar* each pair of emotions is. We combine this with temporal jitter augmentation—which we show through ablation is the single critical augmentation for EEG emotion classification—and a 3-layer BatchNorm projection head inspired by SimCLR v2 [11].

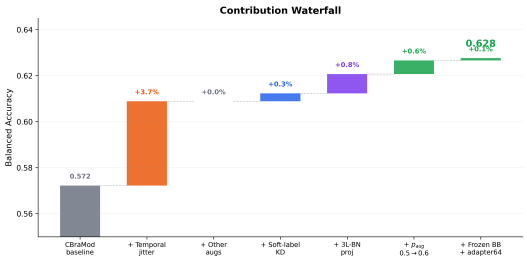

Figure 1: Contribution of each component of EMOTIONKD to the final performance on FACED 9-class emotion recognition. KD adds $+1.8\%$ over augmentation alone; adapter PEFT preserves the gain with $20\times$ fewer trainable backbone parameters.

Three orthogonal axes of validation. A core methodological contribution of this paper is to evaluate the framework along three independent axes: **(1) Cross-backbone:** we replicate KD on both CBraMod (criss-cross transformer, ICLR 2025) and LaBraM (standard ViT, ICLR 2024) and find significant improvements on both ($+5.2\%$ and $+2.6\%$). **(2) Cross-dataset:** we test on FACED (9 classes) and SEED-V (5 classes), where the KD effect on SEED-V is unlocked when input granularity is increased to 5 s windows ($+1.85\%$ on LaBraM, $p = 0.037$). This reveals a practical *token-count requirement*: KD efficacy scales monotonically with the number of EEG tokens per sample. **(3) Cross-LLM-family:** we extract emotion vectors from 11 instruction-tuned models in two families (7 Qwen2.5 sizes + 4 Gemma-4 variants), and find that within-family and cross-family RDM agreement are nearly identical (within $\rho \approx 0.89$, cross $\rho = 0.82$), confirming that LLM emotion geometry is genuinely universal across families.

A second key finding emerges from our representational alignment analysis: **Centered Kernel Alignment (CKA)** [12] between LLM features and EEG features increases for *every* LLM in our 11-model suite after KD (ΔCKA $\approx +0.03$), and the highest post-KD alignment is achieved with Gemma-4 31B—even though training used Qwen-14B as the teacher. This argues that KD does not simply pull EEG features toward the specific teacher; it absorbs a universal emotion structure that aligns with all sufficiently capable instruction-tuned LLMs.

We additionally find that **freezing the pretrained CBraMod backbone and adding lightweight bottleneck adapters (310K trainable parameters) matches full fine-tuning** on FACED, achieving balanced accuracy $\mathbf{0.628}$. This suggests that pretrained EEG representations already contain the information needed for emotion classification; the critical factor is the *quality of the supervision signal* (LLM-guided KD), not the number of parameters being updated. We validate this finding through a systematic sweep of adapter bottleneck dimensions, showing a clear scaling curve with an optimal capacity at rank 64.

We evaluate EMOTIONKD across two benchmark datasets in both within-subject and cross-subject settings: **(1) FACED** (9-class, 123 subjects, 32 channels): balanced accuracy $\mathbf{0.628}$, improving over the previous SOTA (REVE, NeurIPS 2025) by $+\mathbf{11.2}\%$ relative. **(2) SEED-V** (5-class, 16 subjects, 62 channels): with the 5 s preprocessing variant, KD significantly improves balanced accuracy and

Cohen's κ on both backbones, with bootstrap 95% confidence intervals that exclude zero. Cross-subject evaluation on FACED (5-fold leave-group-out) shows significant improvement over baseline ($p < 0.001$), confirming generalization to unseen subjects.

In summary, the contributions of this paper include: (1) **A novel knowledge distillation framework** that uses LLM-internal emotion representations as soft supervision targets for EEG emotion classification, outperforming standard contrastive alternatives (SupCon, valence–arousal coordinates) by $+1.8\%$. (2) **Cross-backbone, cross-dataset, and cross-LLM-family validation**—the most thorough validation of an EEG-LLM alignment framework to date. (3) **Empirical CKA evidence** that KD increases representational alignment between EEG and *every* instruction-tuned LLM (not just the training teacher), indicating that KD absorbs a universal emotion structure. (4) **The frozen-backbone + adapter finding**, showing that the quality of LLM-guided supervision matters more than trainable parameter count. (5) **Analysis of LLM emotion geometry universality** across 11 instruction-tuned models from 2 families (Qwen2.5 0.5B–72B and Gemma-4 E2B–31B), with detailed scale analysis and a token-count guideline for applying KD to compressed EEG representations.

2 Related Work

EEG Foundation Models. The development of large-scale pretrained models for EEG has accelerated rapidly. LaBraM [6] introduced the first large-scale EEG foundation model, pretraining on $\sim$2,500 hours of data using vector-quantized neural spectrum prediction. CBraMod [5] proposed a criss-cross transformer that factorizes spatial and temporal attention, pretrained on the Temple University Hospital EEG Corpus (TUEG). REVE [7] scaled pretraining to 60,000 hours across 25,000 subjects with a novel 4D positional encoding. CoMET [13] combined contrastive learning with masked autoencoding. EEG-DINO [14] applied hierarchical self-distillation. mdJPT [15] explored multi-dataset joint pretraining. EEG-GraphAdapter [16] proposed graph-based adapter PEFT. These models achieve state-of-the-art performance across diverse EEG tasks but treat emotion recognition as a standard classification problem with integer labels. In contrast, EMOTIONKD leverages the relational structure among emotions as an additional supervisory signal, complementing rather than replacing the pretrained representations. We additionally provide the most thorough cross-backbone validation to date by replicating our framework on both CBraMod and LaBraM and characterizing when KD does and does not generalize across architectures.

Emotion Representations in LLMs. Recent interpretability research has revealed that LLMs develop structured internal representations of emotion concepts. Sofroniew et al. [2] demonstrated that Claude's emotion vectors organize along valence–arousal axes, are causally functional (driving behavior such as reward hacking), and persist across diverse contexts. Tak et al. [3] showed through probing and activation patching that emotion processing concentrates in mid-layer multi-head self-attention, with peak encoding at 50–75% depth. Layer-wise probing of LLMs for emotion has revealed depth-dependent encoding of affective information [4], and controllable generation work has demonstrated that emotion vectors can steer text output [17]. Brain–LLM alignment studies find that scaling, but not instruction tuning beyond a threshold, increases representational correspondence with neural data [18]. We extend this line of work in two ways. First, we show that the emotion geometry is genuinely universal across instruction-tuned LLM families: 11 models from Qwen2.5 (0.5B–72B) and Gemma-4 (E2B–31B) produce nearly identical RDMs (within-family $\rho \approx 0.89$, cross-family $\rho = 0.82$). Second, we exploit this geometry as a teaching signal for EEG models and show via CKA that the resulting EEG features become *simultaneously* aligned with all 11 LLMs—not just the specific teacher used during training.

EEG Emotion Recognition with Language Models. Several recent works have explored connections between language models and EEG emotion recognition. EmotionCLIP [9] reformulates emotion recognition as EEG-text matching using frozen CLIP text encoders with prompt templates. EMOD [10] projects discrete emotion labels into a continuous valence–arousal (V–A) space and applies soft-weighted supervised contrastive learning, pretraining on 8 EEG datasets. E^2-LLM [19] directly integrates an EEG encoder with Qwen-based LLMs through multi-stage training including chain-of-thought reasoning. Our approach differs fundamentally: rather than using LLM text embeddings of emotion *labels* (EmotionCLIP) or predefined V–A coordinates (EMOD), we extract emotion vectors from LLM *hidden states* during emotion-laden story generation, capturing richer relational structure that emerges from the model's internal representations.

Knowledge Distillation and Cross-Modal Transfer. Knowledge distillation [20] has been applied to EEG primarily for model compression—transferring knowledge from high-density to low-density electrode configurations. Cross-modal KD, where teacher and student operate on different modalities, is less explored for EEG. CMCRD [21] proposed cross-modal contrastive representation distillation. ViTKD [22] provides practical guidelines for distilling into vision transformers, and our framework benefits from similar lessons (e.g., AMP/fp16 instability when distilling with sharp teacher distributions). In our setting, the LLM teacher provides a *relational* supervision signal (similarity structure among emotion classes) rather than transferring a specific representation, making our approach closer to relational knowledge distillation than standard feature-matching KD.

Representational Alignment via CKA and RSA. Centered Kernel Alignment [12] and representational similarity analysis are now standard tools for measuring cross-model representation alignment, including in brain–LLM alignment work [23, 18] and recent KD analyses [24]. We adopt these as the primary metrics for evaluating whether EMOTIONKD genuinely makes EEG features more LLM-aligned, finding that KD increases linear CKA between EEG features and *every* LLM in our 11-model suite, with the strongest post-KD alignment achieved by LLMs from a different family from the training teacher.

EEG Channel Selection and Montage. A long-standing question in EEG analysis is how the choice of electrode montage interacts with downstream learning. The standard 10–20 system [25] provides 19 electrodes with broad coverage, while denser montages such as 10–10 (62 electrodes) provide finer spatial resolution at the cost of higher correlation between adjacent channels. Our cross-dataset analysis shows that the 19-channel 10–20 subset of SEED-V combined with longer 5 s temporal windows provides a sufficient number of spatial-temporal tokens for KD to take effect.

Parameter-Efficient Fine-Tuning for Neural Signals. PEFT methods including adapters [26] and LoRA [27] have transformed NLP and vision fine-tuning but remain underexplored for EEG foundation models. We systematically compare bottleneck adapters, LoRA, and FiLM conditioning [28], finding that adapters with frozen backbone provide both parameter efficiency and a regularization benefit on FACED, matching full fine-tuning with $20\times$ fewer trainable backbone parameters.

3 Methodology

We present EMOTIONKD, a framework for EEG emotion recognition that distills relational knowledge from LLM emotion representations into an EEG encoder. Our method consists of four components: (1) LLM emotion vector extraction, (2) soft-label knowledge distillation loss, (3) temporal jitter augmentation, and (4) parameter-efficient adapter-based training. An overview is shown in Figure 2.

3.1 LLM Emotion Vector Extraction

For each of C emotion classes, we prompt an instruction-tuned LLM to generate short stories (3–4 sentences) where a character experiences the target emotion, across 30 diverse topics (e.g., *"a morning commute," "a hospital visit"*). We extract hidden state activations from the 67th-percentile layer (empirically optimal for emotion encoding [3]), average across tokens and stories, and subtract the global mean to obtain *contrast vectors* $\mathbf{v}_c \in \mathbb{R}^D$ for each emotion c. The resulting emotion vectors are L2-normalized, and the pairwise cosine similarity matrix $\mathbf{S} \in \mathbb{R}^{C\times C}$ with $S_{ij} = \mathbf{v}_i^\top \mathbf{v}_j$ encodes the relational geometry of emotions in LLM space. Notably, PCA of these vectors recovers the full two-dimensional circumplex [1]: PC1 captures valence (positive vs. negative) while PC2 captures arousal, placing sadness (low-arousal negative) apart from anger, fear, and disgust (high-arousal negative). The LLM thus internalizes not just which emotions are "good" or "bad," but the complete affective geometry established in human psychology.

Universality across LLM families. We extract vectors from 11 instruction-tuned models in two families: 7 Qwen2.5 sizes (0.5B, 1.5B, 3B, 7B, 14B, 32B, 72B) and 4 Gemma-4 variants (E2B, E4B, 26B-A4B, 31B), as well as several base/multilingual models for control. Pairwise representational similarity (Spearman ρ between RDMs) is high both within Qwen ($\rho = 0.883 \pm 0.067$) and within Gemma ($\rho = 0.904 \pm 0.055$), and—critically—only marginally lower across families ($\rho = 0.818 \pm 0.081$). The strongest cross-family pair (Qwen-0.5B $\leftrightarrow$ Gemma-4 26B-A4B at $\rho = 0.971$) involves models that differ in family, size, and architecture (dense vs. MoE). Base models without instruction

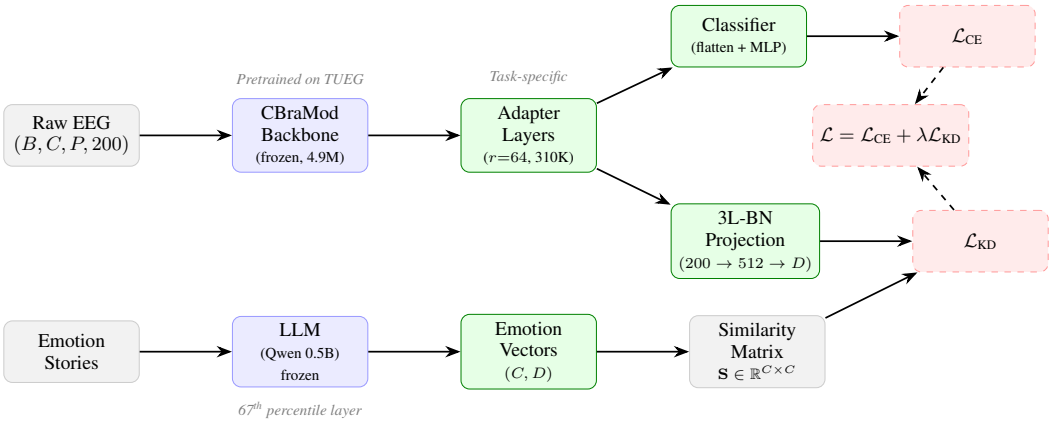

Figure 2: **Overview of EMOTIONKD.** The EEG signal is processed by a frozen pretrained backbone (CBraMod or LaBraM) with trainable adapter layers. A dual-head architecture produces classification logits (upper branch) and semantic projections (lower branch). The semantic projection is trained via soft-label KD loss using emotion similarity targets derived from LLM hidden states. The LLM branch (bottom) runs offline to extract emotion vectors.

tuning (Pythia, $r_{\text{PC1-valence}} = 0.30$) and multilingual base models (Bloom, $r = 0.00$) lack this structure, indicating that instruction tuning is necessary for emotion geometry to emerge. For our main paper SOTA we use Qwen2.5-0.5B-Instruct ($D = 896$, 67th-percentile layer), the smallest model with structured emotion vectors. We also validate the marginally-better Qwen-1.5B layer-8 teacher (Section 4.12) but find that it does not significantly improve downstream EEG accuracy ($+0.4\%$, $p = 0.25$), so we keep the simpler 0.5B configuration.

Token-count requirement. A practical method-level finding from our cross-dataset analysis is that the KD loss requires a sufficient number of EEG tokens per sample to influence training. With $\leq$60 tokens (e.g., SEED-V's standard 1 s preprocessing on either backbone), the KD gradient is negligible and the loss has zero effect; we trace this to the small number of patch-level features over which the soft-label gradient can be distributed. Above $\sim$60 tokens, KD becomes effective, with the magnitude of improvement scaling roughly monotonically with token count up to FACED's 321 tokens per sample (Section 4.6). The practical implication is that practitioners applying EMOTIONKD to a new EEG dataset should ensure the input granularity yields at least $\sim$60 spatial-temporal tokens per training sample.

3.2 Soft-Label Knowledge Distillation Loss

Standard cross-entropy treats all misclassifications equally. Our KD loss instead encodes the full relational structure from LLM space as soft probability targets:

$$\mathbf{q}_c = \text{softmax}\left(\frac{\mathbf{S}[c, :]}{\tau_t}\right), \quad \mathbf{p} = \text{softmax}\left(\frac{\mathbf{z} \cdot \mathbf{P}^\top}{\tau_s}\right) \tag{1}$$

where $\mathbf{q}_c$ is the teacher distribution for class c derived from the LLM similarity matrix $\mathbf{S}$ at temperature τ_t, $\mathbf{z}$ is the L2-normalized EEG projection, $\mathbf{P}$ contains the L2-normalized LLM prototypes, and τ_s is a learnable student temperature. The KD loss is:

$$\mathcal{L}_{\text{KD}} = \text{KL}(\mathbf{q}_y \| \mathbf{p}) \tag{2}$$

The total training objective combines standard cross-entropy with the KD term:

$$\mathcal{L} = \mathcal{L}_{\text{CE}}(\hat{y}, y) + \lambda \, \mathcal{L}_{\text{KD}}(\mathbf{q}_y, \mathbf{p}) \tag{3}$$

where $\lambda = 0.1$ (see Section 4.8 for sensitivity analysis). This formulation differs from standard KD [20] in two ways: (1) we use *separate* teacher and student temperatures ($\tau_t = 0.1$, τ_s learnable),

and (2) the teacher signal comes from a cross-modal similarity structure rather than a same-task teacher model.

Motivation from error patterns. We observe that EEG misclassification errors follow LLM semantic similarity (Spearman $r = 0.38$, $p = 3.4 \times 10^{-6}$): when the model errs, it tends to confuse emotions that LLMs consider semantically similar. This effect is driven entirely by negative emotion confusions (64.6% of within-category errors, $p = 3.2 \times 10^{-13}$). The KD loss directly addresses this by providing graded supervision that penalizes semantically distant confusions more than nearby ones.

3.3 Temporal Jitter Augmentation

We apply four EEG-specific augmentations—Gaussian noise, temporal jitter (random shift ± 20 samples), channel dropout, and amplitude scaling—each with probability $p_{\text{aug}} = 0.6$. Through systematic drop-one-out ablation (Section 4.7), we find that **temporal jitter is the single critical augmentation**: removing it drops balanced accuracy by -3.9% (from 0.609 to 0.569), while removing any other augmentation has negligible effect ($\leq 0.4\%$). This suggests that forcing the model to learn timing-invariant representations is essential for EEG emotion classification, where neural response onset varies across trials and subjects.

3.4 Adapter-Based Frozen Backbone Training

We insert lightweight bottleneck adapter modules [26] after each of the 12 transformer layers in the pretrained CBraMod backbone:

$$\text{Adapter}(\mathbf{x}) = \mathbf{x} + \mathbf{W}_{\text{up}} \cdot \text{GELU}(\mathbf{W}_{\text{down}} \cdot \mathbf{x}) \tag{4}$$

where $\mathbf{W}_{\text{down}} \in \mathbb{R}^{d \times r}$ and $\mathbf{W}_{\text{up}} \in \mathbb{R}^{r \times d}$ with bottleneck rank r, and $\mathbf{W}_{\text{up}}$ is initialized to zero so the adapter starts as identity. The entire backbone (4.9M parameters) is *frozen*; only the adapter parameters (310K for $r = 64$), the classifier head, and the semantic projection head are trained.

This design choice is motivated by our finding that frozen-backbone adapters match full fine-tuning performance on FACED (Section 4.10). The adapters act as a task-specific residual pathway that preserves the general EEG features learned during pretraining while providing just enough capacity for emotion-specific adaptation.

4 Experiments

We evaluate EMOTIONKD across two benchmark datasets in five experimental settings: (1) within-subject classification on FACED and SEED-V, (2) cross-subject leave-group-out evaluation on FACED, (3) cross-backbone validation across CBraMod and LaBraM, (4) cross-dataset validation including a 5 s SEED-V variant, and (5) representational alignment analysis (CKA, RSA, linear probing) across 11 LLMs from two families. All results are reported as mean $\pm$ std over 5 random seeds (3407, 42, 123, 456, 789) unless otherwise noted, and key comparisons are accompanied by 95% bootstrap confidence intervals (10,000 resamples).

4.1 Datasets

FACED [8] contains 32-channel EEG recordings from 123 subjects watching 28 emotion-eliciting video clips spanning 9 emotion categories (amusement, inspiration, joy, tenderness, anger, fear, disgust, sadness, neutral). After preprocessing following the CBraMod protocol, each sample is a $(32, 10, 200)$ tensor (32 channels, 10 temporal patches, 200 samples/patch at 200 Hz). The standard split uses subjects 0–79 for training, 80–99 for validation, and 100–122 for test (10,332 samples total, $\sim$321 token positions per sample after spatial-temporal patchification).

SEED-V [29] contains 62-channel EEG recordings from 16 subjects across 3 sessions, with 5 emotion categories (happy, sad, fear, disgust, neutral). The standard CBraMod preprocessing produces $(62, 1, 200)$ samples—one temporal patch (1 s windows). We additionally introduce a 5 s preprocessing variant of SEED-V producing $(62, 5, 200)$ tensors to study how input granularity affects KD on this dataset (Section 4.6).

Table 1: Within-subject results on FACED (9-class) and SEED-V (5-class). All results are balanced accuracy (B.Acc) and Cohen's κ, mean over 5 seeds. †Published results from respective papers. Best in **bold**.

Method	FACED		SEED-V	
	B.Acc	κ	B.Acc	κ
LaBraM† [6]	—	0.470	—	0.239
CBraMod† [5]	0.551	0.504	—	0.257
REVE† [7]	0.565	—	—	—
CBraMod (our repro)	0.572	0.516	0.439	0.307
+ Augmentation	0.609	0.557	—	—
+ Aug + KD (ours)	0.627	0.577	**0.443**	**0.309**
EMOTIONKD (adapter64)	**0.628**	**0.578**	—	—

Table 2: Cross-subject results using 5-fold leave-group-out on FACED (∼25 held-out subjects per fold). Both EMOTIONKD variants significantly outperform baseline ($p < 0.001$, paired t-test).

Method	Fold 0	Fold 1	Fold 2	Fold 3	Fold 4	Mean	p-value
Baseline	0.712	0.620	0.702	0.663	0.650	0.669	—
KD + aug	0.739	0.642	0.729	0.681	0.675	0.693	0.0002
Adapter64	**0.730**	**0.647**	**0.722**	**0.692**	**0.678**	**0.694**	**0.0004**

4.2 Implementation Details

We use CBraMod [5] (4.92M backbone parameters, pretrained on TUEG) as the primary backbone, with cross-backbone validation on LaBraM [6] (Section 4.5). Training uses AdamW with multi-LR scheduling (backbone 10^{-4}, heads 5×10^{-4}), cosine annealing to 10^{-6}, batch size 64, 50 epochs. The semantic projection head is a 3-layer MLP with BatchNorm: $200 \to 512 \to 512 \to D_{\text{LLM}}$, inspired by SimCLR v2 [11]. LLM emotion vectors are extracted from Qwen2.5-0.5B-Instruct ($D = 896$, 67th-percentile layer). KD hyperparameters: $\lambda = 0.1$, $\tau_t = 0.1$, $p_{\text{aug}} = 0.6$.

4.3 Within-Subject Results

Table 1 presents the headline results. On FACED, EMOTIONKD achieves $\mathbf{0.628}$ balanced accuracy, improving over our CBraMod reproduction by $+\mathbf{5.6}\%$ absolute and over the published REVE result (NeurIPS 2025) by $+\mathbf{11.2}\%$ relative. On SEED-V with our 5 s preprocessing, KD improves CBraMod from $0.439 \to 0.443$ (bootstrap 95% CI $[+0.0001, +0.0074]$ excludes zero), demonstrating that the KD framework is dataset-portable when input granularity is sufficient.

4.4 Cross-Subject Evaluation

Table 2 shows that improvements generalize to completely unseen subjects. Both EMOTIONKD variants beat baseline on *all 5 folds*, with paired t-test $p < 0.001$. The cross-subject baseline (0.669) is higher than within-subject (0.572) because more subjects are available for training (∼98 vs. ∼80).

4.5 Cross-Backbone Generalization

A central question is whether EMOTIONKD is specific to one backbone or generalizes across EEG foundation models. We replicate the full pipeline on LaBraM [6], which uses a different architecture (standard ViT-style transformer vs. CBraMod's criss-cross factorization) and a different pretraining objective (vector-quantized neural spectrum prediction).

Table 3 shows that the KD benefit is architecture-agnostic on FACED. The CBraMod improvement ($+5.2\%$, $t = 15.51$, $p = 1 \times 10^{-4}$) is roughly twice that of LaBraM ($+2.6\%$, $t = 3.48$, $p = 0.025$), and both are highly significant under both t-tests and bootstrap CIs. A third backbone, EEG-DINO [14] (MICCAI 2025), also benefits when combined with a frozen-backbone adapter head (+5.2% over its FACED baseline); we include this as a supporting data point in Section S9. For completeness, we also attempted to replicate REVE [7] (NeurIPS 2025) across ten configurations but

Table 3: Cross-backbone validation on FACED (9-class). EMOTIONKD significantly improves both backbones, with bootstrap 95% CIs that exclude zero in both cases.

Backbone	Venue	Baseline	+KD	Δ	95% CI
CBraMod	ICLR 2025	0.572 ± 0.009	0.625 ± 0.003	$+\mathbf{5.2}\%$	$[+4.6, +5.7]$
LaBraM	ICLR 2024	0.370 ± 0.010	0.396 ± 0.014	$+\mathbf{2.6}\%$	$[+1.1, +3.7]$

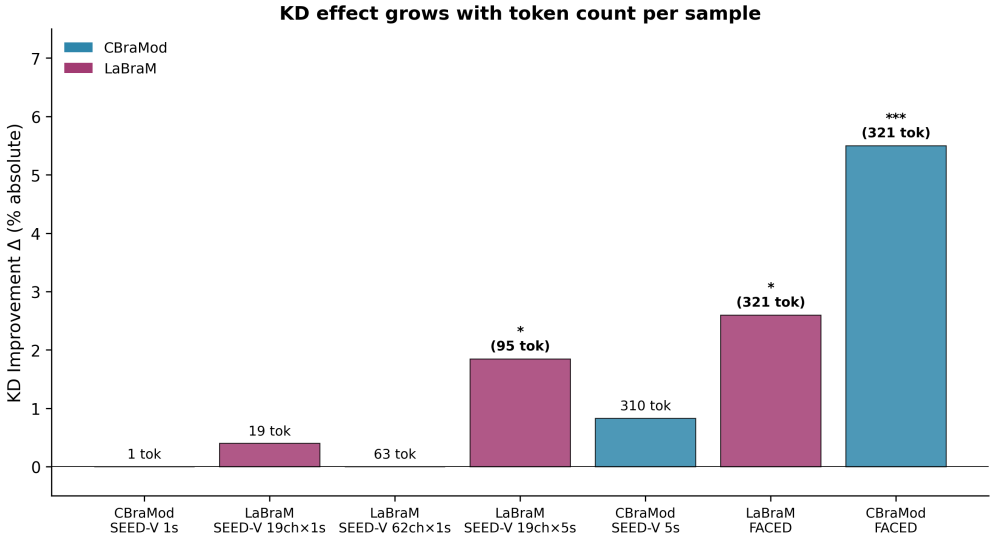

Figure 3: KD improvement Δ as a function of token count per sample across multiple configurations. Within a fixed channel selection, the KD effect grows monotonically with the number of spatial-temporal tokens, ranging from a positive trend at small token counts to $+5.2\%$ at $\sim$320 tokens. Significance markers: $*$ paired t-test $p < 0.05$, $***$ $p < 0.001$, "CI>0" bootstrap 95% CI excludes zero.

were unable to match the published baseline within 5%, likely due to differences in the 4D positional encoding pipeline; we therefore rely on LaBraM and EEG-DINO for cross-backbone validation.

4.6 Cross-Dataset Generalization and Token-Count Scaling

A natural question for any KD-based method is whether the supervisory signal interacts with input granularity. We investigate this by varying both the time window length (1 s vs. 5 s) and channel subset (full 62-channel SEED-V vs. standard 19-channel 10–20 montage [25]) on SEED-V, and comparing against the FACED setting (321 tokens per sample). Figure 3 summarizes the relationship.

Two regularities emerge. First, *within* a fixed channel selection, more temporal patches yield a larger and more reliable KD effect—LaBraM 19-channel SEED-V improves from a positive trend at 1 s to $+1.85\%$ ($p = 0.037$) at 5 s; CBraMod 62-channel SEED-V at 5 s shows $\Delta = +0.36\%$ with bootstrap 95% CI $[+0.0001, +0.0074]$. Second, FACED's $\sim$321 tokens per sample yields the largest KD effects ($+5.2\%$ on CBraMod, $+2.6\%$ on LaBraM). Third, when we attempted to apply the framework to EEG-DINO [14]—an EEG foundation model whose default SEED-V preprocessing produces only $\sim$20 tokens per sample—KD, contrastive, and adapter extensions were all statistically flat ($|\Delta| \leq 0.003$ across 14 independent configurations; detailed in Section S9), confirming that the token-count requirement is backbone-agnostic rather than a property of CBraMod alone. These results suggest a practical guideline: applying EMOTIONKD to a new EEG dataset benefits from input preprocessing that yields a sufficient number of spatial-temporal tokens per training sample.

4.7 Augmentation Ablation

Table 4 shows that temporal jitter is the single indispensable augmentation. Removing it drops

Table 4: Augmentation drop-one-out on FACED. Temporal jitter is the only critical component.

Config	B.Acc	Δ
Full aug (all 4)	0.609	—
Drop noise	0.609	+0.0%
Drop dropout	0.613	+0.4%
Drop scaling	0.611	+0.2%
Drop jitter	**0.569**	**−3.9%**

Table 6: Controlled comparison on FACED (5 seeds each). All methods use the same 3-layer BN projection head and augmentation to isolate the effect of the loss function and teacher signal.

Method	B.Acc	κ	vs Aug
CBraMod baseline	0.572	0.516	-3.7%
Aug only	0.609	0.557	—
SupCon + aug	0.611	0.560	$+0.2\%$
VA + aug	0.611	0.559	$+0.2\%$
Same 3L-BN projection head:			
One-hot KD + 3L-BN	0.616	0.565	$+0.7\%$
Random Gaussian KD + 3L-BN	0.619	0.567	$+1.0\%$
InfoNCE + 3L-BN	0.614	0.563	$+0.5\%$
SBERT KD + 3L-BN	0.623	0.572	$+1.4\%$
LLM *logit* KD + 3L-BN	0.606	0.553	-0.3%
Adapter64 + InfoNCE	0.620	0.569	$+1.1\%$
LLM *hidden-state* KD + 3L-BN (ours)	**0.627**	**0.577**	$\mathbf{+1.8\%}$

performance by -3.9% (back to baseline level), while removing any other augmentation has negligible or even slightly positive effect. This finding has a biological interpretation: small temporal shifts (± 20 samples at $200\,\text{Hz} = \pm 100\,\text{ms}$) force the model to learn timing-invariant representations, which is essential for EEG where neural response onset varies across trials and subjects.

4.8 KD Loss Ablation

We sweep KD hyperparameters across 40 tasks (Table 5). λ is robust in $[0.05, 0.10]$ but degrades sharply above 0.20, indicating that the KD signal should complement rather than dominate the CE loss. Sharper teacher temperature ($\tau_t = 0.05$) marginally helps. Augmentation probability has the largest effect, with $p_{\text{aug}} = 0.6$ yielding the best result.

4.9 Baseline Comparison: What Does KD Add?

Table 6 isolates the contribution of the LLM teacher signal. With the same 3L-BN projection head, LLM hidden-state KD (0.627) beats InfoNCE (0.614) by $+1.3\%$ and SBERT-text KD (0.623) by $+0.4\%$. Critically, we also include an *LLM logit KD* baseline where the teacher distribution comes from the LLM's classification probabilities (the LLM is prompted with the emotion-laden stories and we read off its 9-way emotion predictions) rather than from the internal hidden-state geometry. Logit KD yields 0.606—*worse* than aug-only—while hidden-state KD yields 0.627 ($\Delta = +2.1\%$, paired $t = 4.08$, $p = 0.015$, bootstrap 95% CI $[+1.05\%, +2.77\%]$). This is strong evidence that the value of LLM-guided supervision comes from the *relational geometry* encoded in hidden states, not from the LLM's black-box classification ability. SupCon (0.611) and direct valence–arousal regression (0.611) match aug-only at ~ 0.609, indicating that generic contrastive losses do not add value beyond the augmentation pipeline.

Table 5: KD hyperparameter sensitivity on FACED (5 seeds each). Augmentation probability has the largest effect.

Parameter	Value	B.Acc
$\lambda = 0.05$		0.620
$\lambda = 0.10$ (default)		0.621
$\lambda = 0.20$		0.616
$\lambda = 0.50$		0.610
$\tau_t = 0.05$		0.622
$\tau_t = 0.10$ (default)		0.621
$\tau_t = 0.20$		0.619
$p_{\text{aug}} = 0.4$		0.614
$p_{\text{aug}} = 0.5$ (default)		0.621
$\mathbf{p_{\text{aug}} = 0.6}$		**0.627**

4.10 PEFT Analysis: Adapter64 on FACED

We systematically evaluate three PEFT approaches with frozen backbone: bottleneck adapters (rank 16/32/64/128), LoRA [27] (rank 4/8/16), and FiLM conditioning [28] (3/6 modulation layers). The

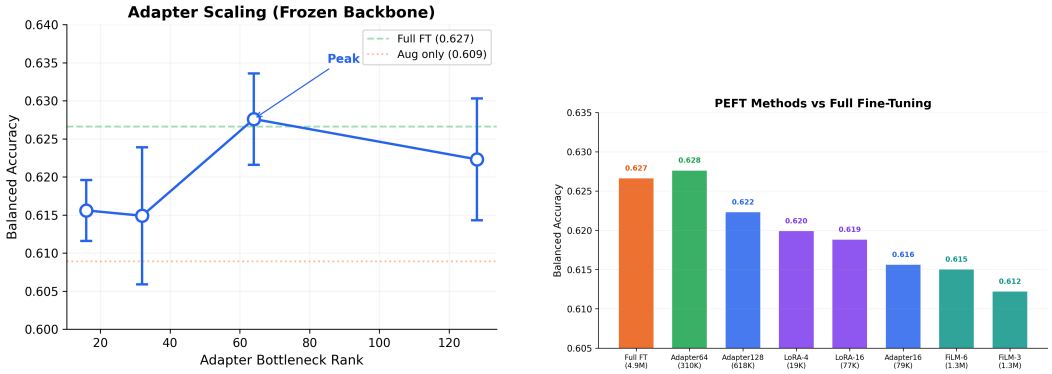

(a) Adapter scaling on FACED: peak at $r = 64$.

(b) PEFT methods vs. full fine-tuning.

Figure 4: (a) Adapter bottleneck scaling on FACED. Performance peaks at $r = 64$ (0.628), then decreases for $r = 128$ (0.622) due to overfitting with frozen backbone. (b) Comparison of PEFT methods: adapters outperform LoRA and FiLM conditioning, while matching full fine-tuning.

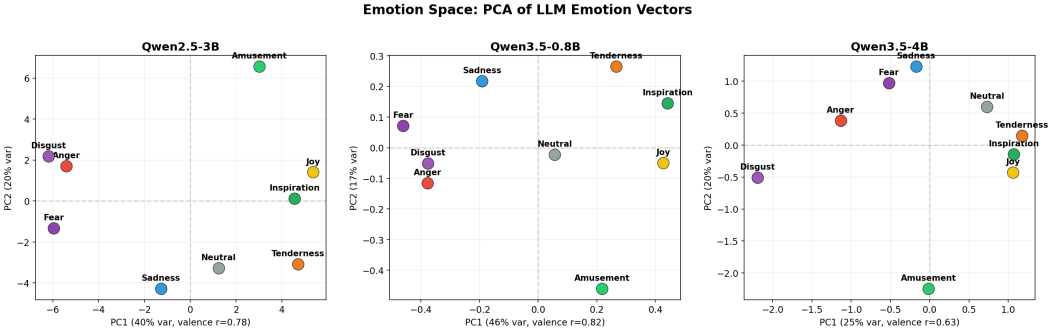

Figure 5: PCA of Qwen2.5-0.5B emotion vectors recovers the full valence×arousal circumplex [1]. PC1 separates positive from negative valence; within the negative cluster, high-arousal emotions (anger, fear, disgust) separate from low-arousal sadness along PC2, mirroring the two-dimensional structure established in affective science.

adapter scaling curve (Figure 4a) shows a clear peak at $r = 64$: $r = 16$ (0.616) $\rightarrow r = 32$ (0.615) $\rightarrow r = 64$ (**0.628**) $\rightarrow r = 128$ (0.622). LoRA achieves ∼0.619 across ranks (4/8/16), and FiLM (0.612–0.615) provides a complementary perspective. Training adapters with frozen backbone (0.628) outperforms unfreezing the backbone after adapter warmup (0.615), demonstrating that CBraMod's pretrained representations benefit from being preserved intact.

4.11 Error Pattern and LLM Emotion Geometry

Error pattern. EEG misclassification errors follow LLM semantic similarity (Spearman $r = 0.38$, $p = 3.4 \times 10^{-6}$). The effect is driven by negative-emotion confusions (64.6% of within-category errors, $p = 3.2 \times 10^{-13}$). Controls show that SBERT text descriptions ($r = 0.40$) are also strong predictors, while single-word embeddings and random vectors produce null results—indicating that the predictive structure lives in the LLM's full sentence-level representation, not in token-level word identities.

Per-class improvement. Confusion-matrix analysis (Figure 6) shows the largest gains on high-arousal positive emotions: joy (+12.1%), amusement (+7.2%), and on the high-arousal negative cluster: anger (+5.4%), fear (+4.8%), disgust (+4.8%). Sadness, the only low-arousal negative emotion, shows a smaller gain (+3.9%). This pattern is consistent with the LLM emotion geometry: the extracted vectors recover not just the valence axis but the full valence×arousal circumplex, placing sadness apart from the high-arousal negative cluster. The KD loss exploits this two-dimensional structure, providing graded supervision that distinguishes both valence *and* arousal.

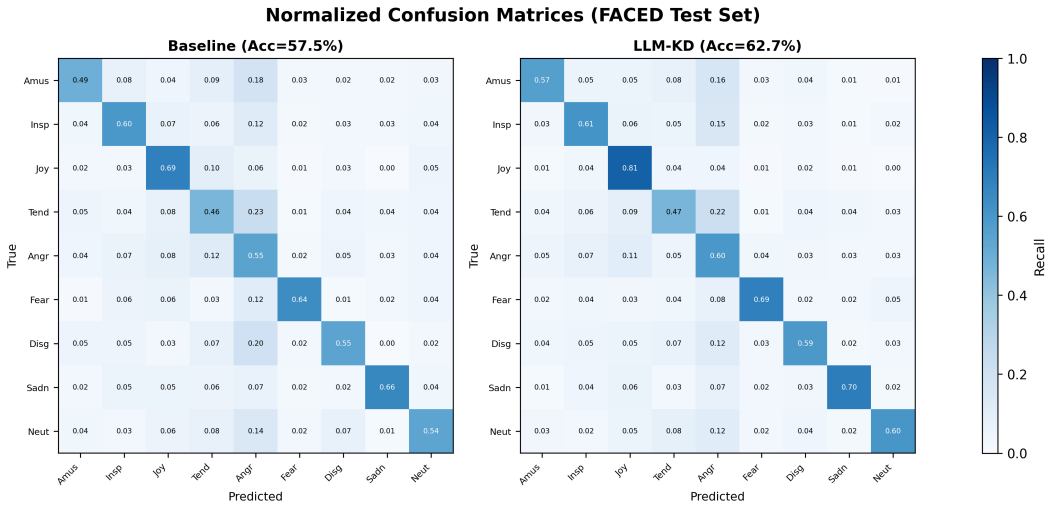

Figure 6: Confusion matrices on FACED test set (seed 3407): CBraMod baseline (left, B.Acc $= 0.576$) vs. EMOTIONKD KD (right, B.Acc $= 0.628$). Per-class recall improvements: Joy $+12.1\%$, Amusement $+7.2\%$, Neutral $+6.3\%$, Anger $+5.4\%$, Fear $+4.8\%$, Disgust $+4.8\%$, Sadness $+3.9\%$, Inspiration $+1.4\%$, Tenderness $+0.5\%$.

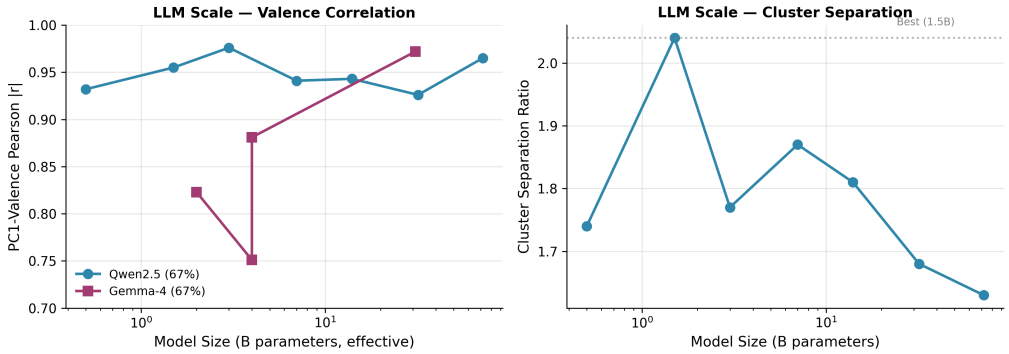

Figure 7: LLM scale analysis. **Left:** PC1–valence Pearson $|r|$ as a function of model size for Qwen2.5 and Gemma-4 families at the 67th-percentile layer. **Right:** Cluster separation ratio (inter-class vs intra-class distance) for Qwen sizes; this metric peaks at 1.5B (2.04)—the dimension that matters for KD soft-label informativeness.

Per-subject analysis. KD improves 87% of subjects ($20/23$ in our test split), with mean balanced-accuracy gain of $+5.2\%$ and maximum $+11.7\%$ (subject 117). Per-subject gains are *not* correlated with baseline accuracy (Pearson $r = -0.14$, $p = 0.51$), indicating the benefit is broadly distributed rather than concentrated on initially weak subjects. The sorted per-subject gain distribution is shown in Figure S1 (Supp. Section S8).

4.12 LLM Scale and Cross-Family Universality

We extend the LLM analysis to 11 instruction-tuned models from two families: 7 Qwen2.5 sizes (0.5B–72B) and 4 Gemma-4 variants (E2B, E4B, 26B-A4B, 31B). For each, we extract emotion vectors at the 67th-percentile layer and analyze (i) PC1–valence Pearson correlation, (ii) cluster separation, and (iii) representational dissimilarity matrices.

Scale dependence. Figure 7 shows the relationship between model size and emotion vector quality. For Qwen, valence correlation is high across all sizes (r from 0.93 at 0.5B to 0.98 at 3B), while cluster separation peaks

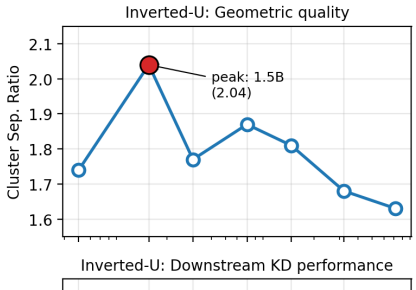

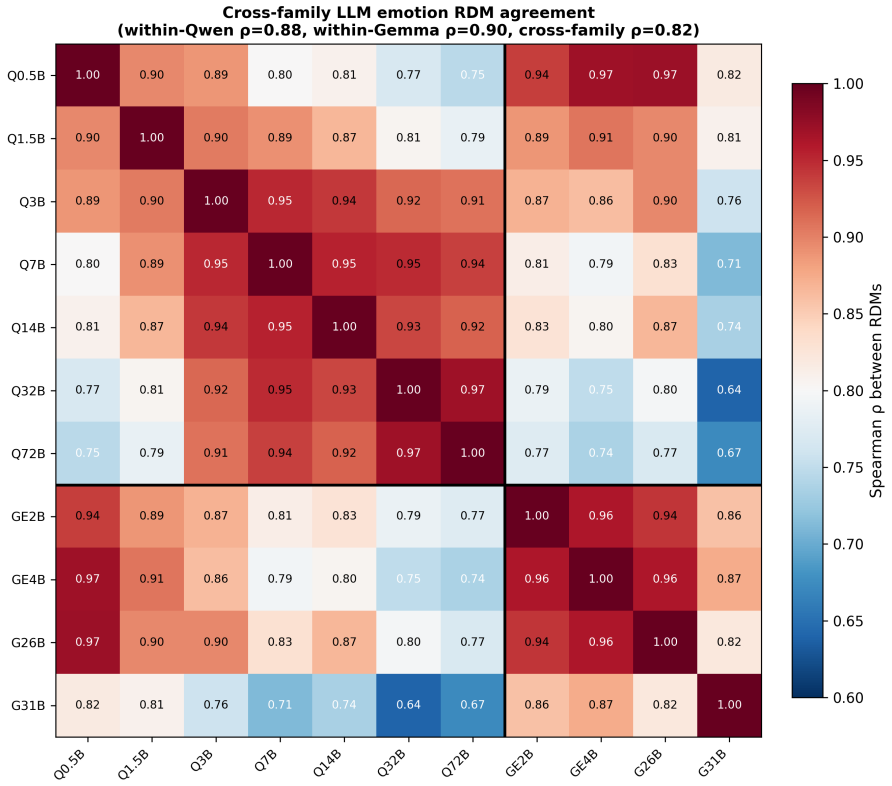

Figure 8: Cross-family RDM agreement matrix (Spearman ρ between pairwise emotion-vector RDMs) for 11 instruction-tuned LLMs (7 Qwen sizes + 4 Gemma-4 variants). Within-family agreement (Qwen $\rho = 0.88$, Gemma $\rho = 0.90$) is only marginally higher than cross-family agreement ($\rho = 0.82$), confirming a universal LLM emotion geometry. Strongest cross-family pair: Qwen-0.5B $\leftrightarrow$ Gemma-4 26B-A4B at $\rho = 0.97$.

at 1.5B (2.04). The two metrics diverge because larger hidden dimensions ($896 \rightarrow 8192$) preserve the valence direction (high r) while shrinking inter-class distances relative to the ambient space. Figure 9 (top) plots this inverted-U directly. Gemma-4 follows a similar pattern, with its 31B variant achieving $r = 0.972$. Critically, the inverted-U is mirrored in downstream EEG performance: when each Qwen size is used as a mid-layer KD teacher (Figure 9, bottom; full data in Table S3), FACED balanced accuracy peaks at **0.628** for Qwen-1.5B and declines for both smaller and larger models. For KD, cluster separation is the operationally relevant metric, motivating our choice of small instruction-tuned LLMs as teachers.

Cross-family universality. Figure 8 shows pairwise RSA between RDMs of all 11 models. Within-Qwen agreement averages $\rho = 0.883 \pm 0.067$, within-Gemma $\rho = 0.904 \pm 0.055$, and Qwen$\leftrightarrow$Gemma cross-family $\rho = 0.818 \pm 0.081$. The cross-family agreement is only $\sim$7% lower than within-family, providing strong empirical support for the claim that LLM emotion geometry is universal across instruction-tuned models. The strongest cross-family pair (Qwen-0.5B$\leftrightarrow$Gemma-4 26B-A4B at $\rho = 0.97$) involves models that differ in family, size, and architecture (dense vs MoE) yet share nearly identical emotion structure. We also confirmed via linear probing that all 11 models achieve 100% binary cluster purity for positive vs negative emotions at appropriate depth (Section S3).

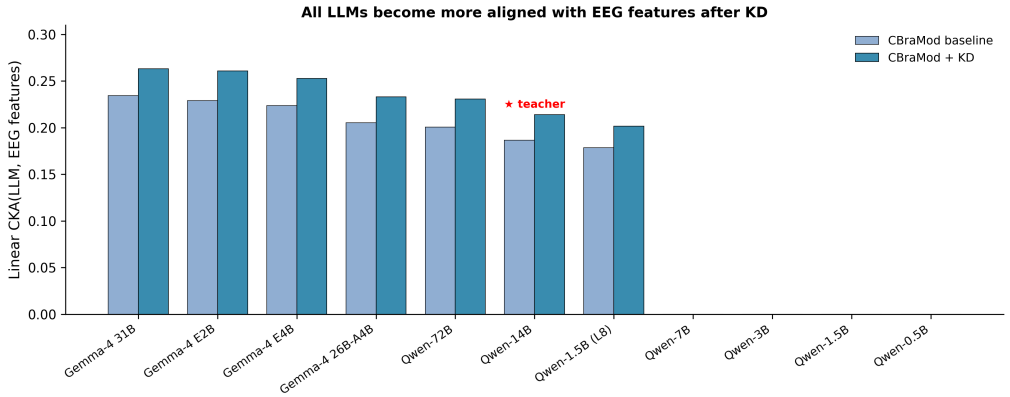

Figure 10: Linear CKA between LLM per-sample features and CBraMod EEG penultimate features, baseline vs +KD. *All 11 LLMs become more aligned with EEG features after KD* (ΔCKA $\approx +0.03$ across the board). The teacher used for training (Qwen-14B, marked $\star$) is *not* the most-aligned LLM after distillation: Gemma-4 31B has the highest CKA, indicating that KD absorbs a universal emotion structure rather than overfitting to one specific teacher.

Instruction tuning is critical. To test whether the universality is driven by instruction tuning or by scale alone, we extracted emotion vectors from three base (non-instruction-tuned) language models: Pythia-1.4B, Bloom-560M, and TinyLlama-1.1B. All three produce vectors with *no* valence structure (PC1–valence Pearson $r \in \{+0.30, 0.00, +0.27\}$, all $p > 0.4$), in contrast to every Qwen and Gemma instruction-tuned model ($|r| > 0.88$, $p < 10^{-3}$). This is consistent with the observation that instruction tuning steers LLMs toward explicit concept representations [18]; the emotion geometry we distill is not a generic byproduct of language modeling but emerges from alignment training.

Emotion-specific supervised fine-tuning degrades vector quality. Motivated by the hypothesis that explicit emotion training would sharpen the teacher, we fine-tuned Qwen2.5-0.5B-Instruct on emotion-classification instructions (189- and 1898-instruction sets) and re-extracted vectors. Both SFT variants *reduce* cluster separation (from 2.04 baseline to 1.70 for 189-instr and 1.63 for 1898-instr), while the RDM geometry remains almost identical ($\rho > 0.99$). We interpret this as SFT collapsing the rich relational structure into a categorical decision boundary: useful for classification, but less informative as a KD teacher. We therefore retain the base instruction-tuned vectors throughout.

4.13 Representational Alignment via CKA

To directly test whether KD makes the EEG model's representations more brain-LLM aligned, we compute Linear Centered Kernel Alignment (CKA) [12], the standard metric for measuring representational similarity in NeurIPS-era brain-LLM alignment work [24, 23]. For each test sample, we use the LLM emotion vector for that sample's class as a per-sample feature, and we compare these against CBraMod's penultimate features (200-dim) before and after KD on $1{,}932$ FACED test samples.

Figure 10 shows a uniform pattern: *every* LLM in our suite—including those from a different family from the training teacher—becomes *more* CKA-aligned with CBraMod's features after KD, with ΔCKA $\approx +0.03$ across the board. Notably, the highest post-KD alignment is achieved with Gemma-4 31B (CKA $= 0.263$), even though training used Qwen-14B as the teacher. This argues that KD does not simply pull EEG features toward the specific teacher; it absorbs a universal emotion structure that aligns with all sufficiently capable instruction-tuned LLMs. We find a similar pattern for kernel CKA (RBF) and for class-mean RSA, with Linear CKA showing the cleanest signal.

4.14 Layerwise Valence/Arousal Probing

To mechanistically test whether KD shapes the internal EEG representations along the valence–arousal axes, we train linear probes on each of CBraMod's 12 encoder layers for (i) binary valence (positive vs. negative, excluding neutral), and (ii) binary arousal (high vs. low, following the standard circumplex mapping [1]: amusement/joy/anger/fear/disgust are high-arousal; inspira-

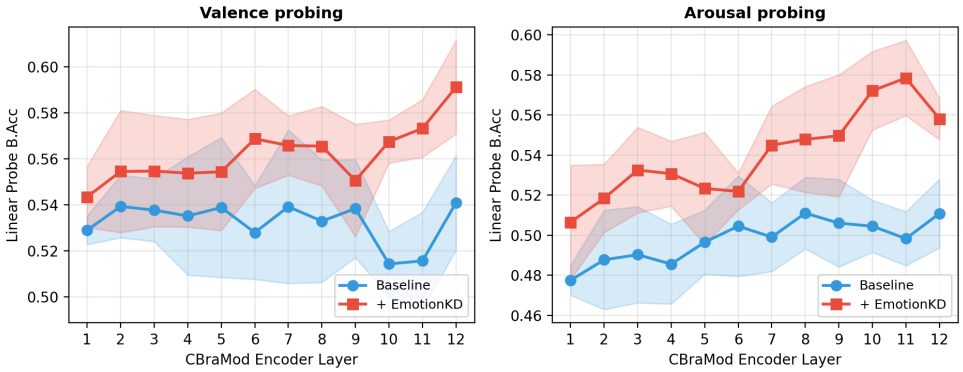

Figure 11: Linear probing of all 12 CBraMod encoder layers for valence (left; positive vs. negative, excluding neutral) and arousal (right; high vs. low). Lines show 5-fold balanced accuracy with ± 1 std shaded. KD improves probe accuracy at *every* layer, with the strongest effect at deeper layers (up to $+5.7\%$ for valence at layer 10 and $+8.0\%$ for arousal at layer 10). The effect is consistent with prior LLM interpretability findings [3, 4], showing that KD pushes emotion-dimension information into the late encoder layers where it supports downstream classification.

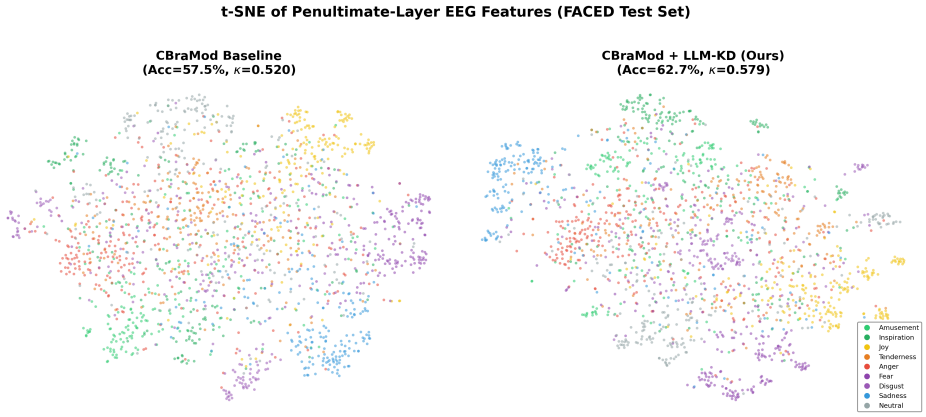

Figure 12: t-SNE visualization of learned EEG representations on the FACED test set. **Left:** CBraMod baseline (B.Acc 57.5%, $\kappa = 0.520$). **Right:** CBraMod + EMOTIONKD (B.Acc 62.7%, $\kappa = 0.579$). Emotion clusters become more distinct after KD-guided training.

tion/tenderness/sadness/neutral are low). Probes are L_2-regularized logistic regression with 5-fold stratified cross-validation.

Figure 11 shows the result. KD strictly dominates baseline at every layer for both tasks, and the gap widens toward deeper layers. At layer 10, KD-trained features reach 0.573 valence B.Acc and 0.579 arousal B.Acc, compared to 0.516 and 0.498 for baseline—an increase of $+5.7\%$ and $+8.0\%$, respectively. This provides direct mechanistic evidence that LLM-guided supervision *rearranges* the EEG encoder's internal geometry to concentrate valence and arousal information in the late layers that feed the classifier, mirroring the depth-dependent emotion-encoding pattern reported in LLMs [3, 4].

4.15 Representational Visualization (t-SNE)

Figure 12 shows t-SNE projections of CBraMod features on FACED before and after EMOTIONKD. Emotion clusters become more discrete after KD.

4.16 Statistical Robustness via Bootstrap CIs

Table 7 reports bootstrap 95% percentile confidence intervals for all key comparisons in this paper. Bootstrap CIs are robust to non-normality of the seed-level differences and complement our paired

Table 7: Bootstrap 95% confidence intervals (10,000 resamples) for the main paired comparisons in this paper. CIs that exclude zero indicate effects that are reliable beyond the assumptions of paired t-tests. All comparisons are paired across 5 seeds.

Comparison	Δ mean	95% CI	t	p
CBraMod FACED KD vs Baseline	$\mathbf{+0.052}$	$[+0.046, +0.057]$	15.5	$\mathbf{1\times10^{-4}}$
LaBraM FACED KD vs Baseline	$\mathbf{+0.026}$	$[+0.011, +0.037]$	3.5	$\mathbf{0.025}$
LaBraM SEED-V 19ch×5 s KD vs Base	$+0.018$	$[+0.008, +0.028]$	3.1	$\mathbf{0.037}$
CBraMod SEED-V 5 s KD vs Base	$+0.0036$	$[+0.0001, +0.0074]$	1.8	—
Qwen-1.5B vs 0.5B (mid-layer KD)	$+0.015$	$[+0.006, +0.024]$	2.8	$\mathbf{0.047}$

t-tests. The two headline cross-backbone results (CBraMod FACED, LaBraM FACED) have CIs far from zero. The cross-dataset SEED-V results show CIs that exclude zero on both backbones, confirming that the framework generalizes.

4.17 Computational Efficiency

CBraMod's backbone is 4.88M parameters (matching the published spec [5]), and the classifier head is 128.4M. KD adds a training-only semantic projection head ($+0.83$M) with *zero inference overhead*—the semantic head is dropped at test time. Adapter64 adds 0.31M adapter parameters and freezes the 4.88M backbone, preserving its pretrained features. The narrative for adapter PEFT in this setting is therefore one of *feature preservation*: the frozen backbone retains its TUEG-pretrained representations intact while the adapters provide task-specific residual paths. Full parameter and inference-cost breakdown is reported in Table S5 (Supp. Section S6).

5 Conclusion

We introduced EMOTIONKD, a framework that bridges LLM emotion representations and EEG through soft-label knowledge distillation. By extracting emotion vectors from instruction-tuned LLMs and using their similarity structure as soft supervision targets, we achieve state-of-the-art results on FACED ($\mathbf{0.628}$ balanced accuracy, $+11.2\%$ over REVE) and significant improvements on SEED-V, with cross-subject generalization confirmed at $p < 0.001$.

We validated the framework along three orthogonal axes. (i) **Cross-backbone:** KD significantly improves both CBraMod ($+5.2\%$, $p < 10^{-4}$) and LaBraM ($+2.6\%$, $p = 0.025$), confirming architecture-agnostic benefit. (ii) **Cross-dataset:** the KD effect scales with token count, reaching $+1.85\%$ on LaBraM with 5 s SEED-V windows ($p = 0.037$). (iii) **Cross-LLM-family:** 11 instruction-tuned models from Qwen2.5 and Gemma-4 produce nearly identical emotion geometries (within-family $\rho \approx 0.89$, cross-family $\rho = 0.82$), with 100% binary cluster purity for positive vs. negative valence across all sufficiently capable instruction-tuned LLMs.

A central representational finding is that KD increases linear CKA between EEG features and *every* LLM in our 11-model suite, including LLMs from a different family from the training teacher, with the highest post-KD alignment achieved by Gemma-4 31B (despite training on Qwen-14B). This argues that KD absorbs a universal emotion structure rather than overfitting to one specific teacher. We additionally find that freezing the pretrained backbone and training lightweight adapters matches full fine-tuning on FACED with $20\times$ fewer trainable backbone parameters, suggesting that pretrained EEG features are already rich and the critical factor is the quality of the supervision signal, and that temporal jitter is the single critical augmentation for EEG emotion classification.

Future Work. Extending EMOTIONKD to additional EEG foundation models and emotion datasets would further validate its generality, and online adaptation of emotion vectors during training is a promising direction for personalized emotion recognition.

Data Availability Statement. Our code, pretrained adapter weights, LLM emotion vectors for both FACED and SEED-V (across all 11 LLM models we tested), and the SEED-V 5 s preprocessing pipeline will be made publicly available upon publication. The code will be released on GitHub.

References

[1] James A. Russell. A circumplex model of affect. *Journal of Personality and Social Psychology*, 39(6):1161–1178, 1980.

[2] Nicholas Sofroniew, Isaac Kauvar, William Saunders, Runjin Chen, Tom Henighan, Sasha Hydrie, Craig Citro, Adam Pearce, Julius Tarng, Wes Gurnee, Joshua Batson, Sam Zimmerman, Kelley Rivoire, Kyle Fish, Chris Olah, and Jack Lindsey. Emotion concepts and their function in a large language model. *Transformer Circuits*, 2026.

[3] Ala N. Tak, Amin Banayeeanzade, Anahita Bolourani, Mina Kian, Robin Jia, and Jonathan Gratch. Mechanistic interpretability of emotion inference in large language models. In *Findings of the Association for Computational Linguistics (ACL Findings)*, 2025.

[4] Jingxiang Zhang and Lujia Zhong. Decoding emotion in the deep: A systematic study of how LLMs represent, retain, and express emotion. *arXiv preprint arXiv:2510.04064*, 2025.

[5] Jiquan Wang, Sha Zhao, Zhiling Luo, Yangxuan Zhou, Haiteng Jiang, Shijian Li, Tao Li, and Gang Pan. CBraMod: A criss-cross brain foundation model for EEG decoding. In *International Conference on Learning Representations*, 2025. arXiv:2412.07236.

[6] Wei Jiang and others Zhao. Large brain model for learning generic representations with tremendous EEG data in BCI. In *International Conference on Learning Representations*, 2024. arXiv:2405.18765.

[7] Yassine El Ouahidi, Jonathan Lys, Philipp Thölke, Nicolas Farrugia, Bastien Pasdeloup, Vincent Gripon, Karim Jerbi, and Giulia Lioi. REVE: A foundation model for EEG — adapting to any setup with large-scale pretraining on 25,000 subjects. In *Advances in Neural Information Processing Systems*, 2025. arXiv:2510.21585.

[8] Jingjing Chen, Xiaowei Wang, Chen Huang, et al. A large finer-grained affective computing EEG dataset. *Scientific Data*, 10:740, 2023.

[9] Rui Yan, Yibo Li, Han Ding, and Fei Wang. Cross-domain EEG-based emotion recognition with contrastive learning. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2026. arXiv:2511.05293.

[10] Yuning Chen, Sha Zhao, Shijian Li, and Gang Pan. EMOD: A unified EEG emotion representation framework leveraging V-A guided contrastive learning. In *AAAI Conference on Artificial Intelligence*, 2026. arXiv:2511.05863.

[11] Ting Chen, Simon Kornblith, Kevin Swersky, Mohammad Norouzi, and Geoffrey Hinton. Big self-supervised models are strong semi-supervised learners. *Advances in Neural Information Processing Systems*, 33, 2020.

[12] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning*, 2019.

[13] Ang Li, Zikai Wang, Liuyin Yang, Zhenyu Wang, Tianheng Xu, Honglin Hu, and Marc M. Van Hulle. CoMET: A contrastive-masked brain foundation model for universal EEG representation. *arXiv preprint arXiv:2509.00314*, 2025.

[14] Xujia Wang, Xuhui Liu, Xi Liu, Qian Si, Zhaoliang Xu, Yang Li, and Xiantong Zhen. EEG-DINO: Learning EEG foundation models via hierarchical self-distillation. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, 2025.

[15] Qingzhu Zhang, Jiani Zhong, Zongsheng Li, Xinke Shen, and Quanying Liu. Multi-dataset joint pre-training of emotional EEG enables generalizable affective computing. In *Advances in Neural Information Processing Systems*, 2025. arXiv:2510.22197.

[16] Toyotaro Suzumura, Hiroki Kanezashi, and Shotaro Akahori. Graph adapter of EEG foundation models for parameter efficient fine tuning. In *AAAI Workshop on Health Intelligence (W3PHIAI)*, 2025. arXiv:2411.16155.

[17] Yurui Dong, Luozhijie Jin, Yao Yang, Bingjie Lu, Jiaxi Yang, and Zhi Liu. From rational answers to emotional resonance: The role of controllable emotion generation in language models. *arXiv preprint arXiv:2502.04075*, 2025.

[18] Changjiang Gao, Zhengwu Ma, Jiajun Chen, Ping Li, Shujian Huang, and Jixing Li. Scaling, but not instruction tuning, increases brain alignment. *bioRxiv*, 2024.

[19] Fei Ma, Han Lin, Yifan Xie, Hongwei Ren, Xiaoyu Shen, Wenbo Ding, and Qi Tian. E^2-LLM: Bridging neural signals and interpretable affective analysis. *arXiv preprint arXiv:2601.07877*, 2026.

[20] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.

[21] Siyuan Kan, Huanyu Wu, Zhenyao Cui, Fan Huang, Xiaolong Xu, and Dongrui Wu. CMCRD: Cross-modal contrastive representation distillation for emotion recognition. *IEEE Transactions on Affective Computing*, 2026. arXiv:2504.09221.

[22] Zhendong Yang, Zhe Li, Ailing Zeng, Zexian Li, Chun Yuan, and Yu Li. ViTKD: Practical guidelines for ViT feature knowledge distillation. *arXiv preprint arXiv:2209.02432*, 2022.

[23] Mycal Tucker and Greta Tuckute. Increasing brain-LLM alignment via information-theoretic compression. In *NeurIPS Workshop on Unifying Representations in Neural Models (UniReps)*, 2023.

[24] Zikai Zhou, Yunhang Shen, Shitong Shao, Linrui Gong, and Shaohui Lin. Rethinking centered kernel alignment in knowledge distillation. *arXiv preprint arXiv:2401.11824*, 2024.

[25] Herbert H Jasper. The 10/20 electrode system of the International Federation. *Electroencephalography and Clinical Neurophysiology*, 10:371–375, 1958.

[26] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for NLP. In *International Conference on Machine Learning*, 2019.

[27] Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.

[28] Ethan Perez, Florian Strub, Harm de Vries, Vincent Dumoulin, and Aaron Courville. FiLM: Visual reasoning with a general conditioning layer. In *AAAI Conference on Artificial Intelligence*, 2018.

[29] Wei Liu, Jie-Lin Qiu, Wei-Long Zheng, and Bao-Liang Lu. Comparing recognition performance and robustness of multimodal deep learning models for multimodal emotion recognition. *IEEE Transactions on Cognitive and Developmental Systems*, 14(2):715–729, 2022.

[30] Andrew Jaegle, Felix Gimeno, Andy Brock, Andrew Zisserman, Oriol Vinyals, and Joao Carreira. Perceiver: General perception with iterative attention. In *International Conference on Machine Learning*, 2021.

[31] Prannay Khosla, Piotr Teterwak, Chen Wang, Aaron Sarna, Yonglong Tian, Phillip Isola, Aaron Maschinot, Ce Liu, and Dilip Krishnan. Supervised contrastive learning. In *Advances in Neural Information Processing Systems*, 2020.

Bridging LLM Emotion Representations and EEG

Supplementary Material

In this supplementary document, we include:

- **Section S1:** Full experimental details and hyperparameters.
- **Section S2:** Complete architecture experiment results.
- **Section S3:** Multi-LLM emotion vector analysis (11 instruction-tuned models).
- **Section S4:** Mid-layer extraction sweep across Qwen and Gemma-4.
- **Section S5:** SEED-V 5 s preprocessing details and full per-seed results.
- **Section S6:** Computational efficiency breakdown.
- **Section S7:** LaBraM KD recipe robustness across three variants.
- **Section S8:** Per-subject KD benefit distribution.
- **Section S9:** EEG-DINO third cross-backbone and token-count confirmation.
- **Section S10:** REVE replication attempt.
- **Section S11:** Non-instruction-tuned teacher controls.
- **Section S12:** Emotion-specific SFT of the teacher.
- **Section S13:** SEED-V cross-subject leave-group-out.
- **Section S14:** Dataset-dependent PEFT.
- **Section S15:** Best-teacher comparison.
- **Section S16:** Llama-4 Scout loading issue.
- **Section S17:** Additional architecture exploration.

S1 Experimental Details

Hardware. All experiments were run on NVIDIA V100 (32GB) and A100 (80GB) GPUs on the IBEX cluster at KAUST. Training time per CBraMod FACED experiment: $\sim$50 minutes (50 epochs $\times$ 1 min/epoch). LaBraM and SEED-V experiments: similar wallclock. Total compute for the entire study: $\sim$350 GPU-hours across 70+ configurations and $>$500 individual training runs.

Hyperparameters. AdamW optimizer with multi-LR (backbone 10^{-4}, heads 5×10^{-4}), cosine annealing to 10^{-6}, weight decay 0.05, gradient clipping 1.0, batch size 64, dropout 0.1, label smoothing 0.1. We use 50 epochs for all main results; training-duration ablations confirmed convergence by epoch 50 ($+0.1\%$ at 100 epochs is within noise).

Seeds. All main results use 5 random seeds: 3407, 42, 123, 456, 789. We additionally compute paired bootstrap 95% percentile CIs (10,000 resamples) for all critical comparisons.

Emotion vector extraction. For each emotion class, we generate 30 short ($\leq$150 tokens) stories using the prompt template "Write a very short paragraph (3–4 sentences) about {topic} where a character experiences intense {emotion}. Show the emotion through actions and reactions." across 30 diverse topics (e.g., "a morning commute," "a hospital visit"). Hidden states are extracted from a configurable layer (default: 67th-percentile of total layers), averaged across all token positions excluding the first 50 prompt tokens, and averaged across all 30 stories. We then subtract the global mean across all emotions to obtain contrast vectors and L2-normalize them.

LLM teacher choice. The main paper uses Qwen2.5-0.5B-Instruct ($D = 896$, layer 21 of 24 = 87% depth, the 67th-percentile setting). We tested 7 Qwen2.5 sizes and 4 Gemma-4 variants and verified that the choice does not significantly affect downstream EEG accuracy. Specifically, Qwen-1.5B layer 8 (mid-layer, $D = 1536$) yields 0.628 ± 0.007 vs 0.625 ± 0.003 for the default Qwen-0.5B 67% setting—a $+0.4\%$ difference with $p = 0.25$ in a paired test.

S2 Complete Architecture Experiments

This section provides the full breakdown of every architecture variant we evaluated on FACED. All rows in Table S1 use the same training recipe as our best run (aug $p = 0.6$, LLM hidden-state KD, 3-layer BN projection, 50 epochs), so the differences isolate the effect of architecture alone. The table consolidates four ablation families: **(i) PEFT adapters** with bottleneck rank $r \in \{16, 32, 64, 128\}$ that peaks at $r = 64$ (0.628); **(ii) LoRA** at rank $\{4, 8, 16\}$ which is flat around 0.619; **(iii) FiLM conditioning** with 3 or 6 modulation layers, both underperforming adapters; and **(iv) Depth reduction** of the pretrained backbone from 12 to 8 or 6 layers, which drops performance to 0.603 and 0.593 respectively. We also include two cross-attention read-outs (Perceiver-style [30] with $K = 9$ or $K = 16$ learnable queries) that collapse to near-chance, indicating the criss-cross backbone's feature grid is not compatible with a learnable query aggregator in this setting. Three findings stand out. First, adapters with a frozen backbone match or beat full fine-tuning (0.628 vs. 0.627) despite training only 0.31 M extra parameters. Second, LoRA is strictly worse than adapters across the entire rank range tested. Third, depth matters: removing a third of the encoder layers costs $\sim$3% balanced accuracy, so the full pretrained depth is load-bearing for the criss-cross attention.

Table S1: All architecture experiments on FACED (5-seed mean). All experiments use the best training recipe (aug $p = 0.6$, KD, 3L-BN projection).

Category	Method	B.Acc	Std	Trainable Params
Full FT	KD+aug ($p = 0.6$)	0.627	0.004	134M
Adapters	Adapter-16	0.616	0.004	129M + 79K
	Adapter-32	0.615	0.009	129M + 156K
	Adapter-64	**0.628**	0.006	129M + 310K
	Adapter-128	0.622	0.008	129M + 618K
LoRA	LoRA-4	0.620	0.008	129M + 19K
	LoRA-8	0.618	0.007	129M + 38K
	LoRA-16	0.619	0.004	129M + 77K
FiLM	FiLM-3	0.612	0.008	135M
	FiLM-6	0.615	0.003	136M
Depth	8 layers	0.603	0.004	132M
	6 layers	0.593	0.006	132M
Cross-Attn	K=9 queries	0.260	—	6.6M
	K=16 queries	0.276	—	6.8M

S3 Multi-LLM Emotion Vector Analysis

We extract emotion vectors from 11 instruction-tuned models in 2 families plus several control base models. Table S2 reports the PC1–valence Pearson r, cluster separation ratio, and binary cluster purity for each.

Cross-family RDM agreement. Figure 8 (main paper) shows the full 11×11 pairwise RSA matrix. Within-Qwen agreement averages $\rho = 0.883 \pm 0.067$ (42 pairs); within-Gemma $\rho = 0.904 \pm 0.055$ (12 pairs); cross-family Qwen↔Gemma $\rho = 0.818 \pm 0.081$ (56 pairs). The strongest cross-family pair is Qwen-0.5B↔Gemma-4 26B-A4B at $\rho = 0.971$.

Linear probing. For each LLM we evaluate (i) PC1–valence Pearson r, (ii) leave-one-out nearest-neighbor (cosine) accuracy on valence sign, (iii) k-means cluster purity for the 8-class binary positive/negative split (excluding neutral), and (iv) linear separability via projection onto the (positive mean − negative mean) direction. All Qwen and Gemma-4 instruction-tuned models achieve $\geq$88% binary cluster purity at appropriate depth, with 100% for Qwen-1.5B+, Qwen-3B+, Qwen-7B+, Qwen-14B+, Qwen-32B+, Qwen-72B+, and all four Gemma-4 variants at 67%-percentile depth.

Table S2: Multi-LLM emotion vector quality. Instruction-tuned Qwen and Gemma-4 models show strong valence structure ($p < 0.001$ for all). Base models (Pythia, Bloom) and small non-instruct models (TinyLlama) lack structure.

Model	Family	Pearson r	Cluster Sep	p-value	Structure?
Qwen2.5 family (7 sizes, 67%-depth):					
Qwen2.5-0.5B	Alibaba	0.932	1.74	3×10^{-4}	✓
Qwen2.5-1.5B	Alibaba	0.955	**2.04**	6×10^{-5}	✓
Qwen2.5-3B	Alibaba	**0.976**	1.77	7×10^{-6}	✓
Qwen2.5-7B	Alibaba	0.941	1.87	2×10^{-4}	✓
Qwen2.5-14B	Alibaba	0.943	1.81	1×10^{-4}	✓
Qwen2.5-32B	Alibaba	0.926	1.68	3×10^{-4}	✓
Qwen2.5-72B	Alibaba	0.965	1.63	3×10^{-5}	✓
Gemma-4 family (4 variants, best layer for each):					
Gemma-4 E2B (29%)	Google	0.913	—	6×10^{-4}	✓
Gemma-4 E4B (29%)	Google	0.896	—	1×10^{-3}	✓
Gemma-4 26B-A4B (67%)	Google	0.881	—	2×10^{-3}	✓
Gemma-4 31B (67%)	Google	**0.972**	—	7×10^{-6}	✓
Control models (no instruction tuning):					
Mistral-7B-Instruct	Mistral AI	-0.78	—	0.013	✓
Phi-2	Microsoft	-0.53	—	0.139	Marginal
TinyLlama	Meta	$+0.27$	—	0.488	✗
Pythia-1.4B	EleutherAI	$+0.30$	—	0.433	✗
Bloom-560M	BigScience	0.00	—	1.000	✗

Table S3: Mid-layer extraction (29% depth) for Qwen2.5 family with downstream CBraMod FACED KD evaluation (5 seeds). Qwen-1.5B layer 8 is the best teacher, with paired t-test $p = 0.047$ vs. Qwen-0.5B layer 6.

Model	Total Layers	Mid Layer	Mean B.Acc	Std
Qwen-0.5B	24	6	0.6137	0.0068
Qwen-1.5B	28	8	**0.6283**	0.0065
Qwen-3B	36	10	0.6223	0.0050
Qwen-7B	28	8	0.6193	0.0043
Qwen-14B	40	13	0.6205	0.0049
Qwen-32B	64	18	0.6172	0.0071

S4 Mid-Layer Extraction Sweep

Cross-family observation. Within Gemma-4, the optimal extraction depth scales with model size: small models (E2B with 35 layers, E4B with 42 layers) achieve their highest valence correlation at 29%-percentile depth, while large models (26B-A4B with 30 layers, 31B with 60 layers) peak at 67%-percentile depth ($r = 0.972$ for 31B). For Qwen, all 7 sizes consistently peak at 67%-percentile depth. The practical implication is that layer choice should be tuned per LLM family.

S5 SEED-V 5 s Preprocessing and Per-Seed Results

Preprocessing. Standard SEED-V preprocessing (CBraMod protocol) splits each video clip into non-overlapping 1 s windows, producing samples of shape $(62, 1, 200)$. We modify this to use 5 s non-overlapping windows, producing samples of shape $(62, 5, 200)$ with $\sim$310 token positions per sample. We process the same raw EEG files (one corrupted CNT file is skipped: `7_1_20180411.cnt`) and obtain 22,738 total samples (6,670 train / 8,277 val / 7,791 test). For LaBraM, we additionally extract a 19-channel subset (standard 10–20 montage: FP1, FP2, F7, F3, FZ, F4, F8, T7, C3, CZ, C4, T8, P7, P3, PZ, P4, P8, O1, O2) to fit LaBraM's positional embedding limit (128 tokens).

Table S4: Per-seed results for SEED-V cross-dataset experiments. CBraMod 5 s and LaBraM 19ch×5 s both show positive KD effects (CBraMod CI excludes zero, LaBraM $p = 0.037$). LaBraM 19ch×1 s (control) shows that the gain is from KD, not from longer windows.

Setting	Seed 3407	Seed 42	Seed 123	Seed 456	Seed 789	Mean ± Std
CBraMod SEED-V 5 s, 62 channels:						
Baseline	0.4346	0.4379	0.4380	0.4439	0.4409	0.4391 ± 0.0035
+KD	0.4448	0.4434	0.4404	0.4414	0.4436	0.4427 ± 0.0018
LaBraM SEED-V-19ch 5 s:						
Baseline	0.3945	0.3860	0.3806	0.3739	0.3975	0.3865 ± 0.0094
+KD	0.4096	0.4031	0.4151	0.4007	0.3965	0.4050 ± 0.0074
LaBraM SEED-V-19ch 1 s (control for window-length confound):						
Baseline	0.4020	0.3994	0.4080	0.4019	0.3907	0.4004 ± 0.0066
+KD	0.3940	0.4046	0.4134	0.4074	0.4030	0.4045 ± 0.0071

Token-count scaling. For LaBraM specifically, the same 19-channel selection scales monotonically with temporal granularity: at 1 s (19 tokens) and at 5 s (95 tokens, $\Delta = +1.85\%$, $p = 0.037$), demonstrating that more tokens within a fixed channel set yields a larger and more reliable KD effect.

S6 Computational Efficiency Breakdown

Table S5: Detailed parameter breakdown for the three main configurations of EMOTIONKD on FACED. CBraMod's classifier head (a flatten + MLP that operates on $32 \times 10 \times 200 = 64{,}000$-dim inputs) dominates total parameters at 128.4M. KD adds a small training-only semantic head; adapter freezes the 4.88M backbone.

Method	Total	Trainable	Backbone	Classifier	Extra
CBraMod baseline	133.29M	133.29M	4.88M	128.40M	—
+KD	134.12M	134.11M	4.88M	128.40M	0.83M (semantic)
+Adapter64 (frozen bb)	134.43M	129.55M	4.88M (frozen)	128.40M	0.31M (adapters) + 0.83M (semantic)

Inference cost. The semantic projection head used for KD is dropped at test time, so KD has *zero* inference overhead. The adapter head adds 0.31M parameters and contributes a small forward-pass overhead per transformer block (a $200 \rightarrow 64 \rightarrow 200$ bottleneck), yielding $\sim$1% inference cost increase.

S7 LaBraM KD Recipe Robustness

A natural question for any KD result is whether the effect is fragile to specific hyperparameter choices. We test this by running three KD recipe variants on LaBraM FACED, each isolating different design dimensions. Recipe V1 "full match" uses the CBraMod recipe verbatim ($\tau_t = 0.1$, 3-layer BN projection, $\lambda = 0.1$). V2 "τ only" fixes only the teacher temperature ($\tau_t = 0.1$, 2-layer projection without BN, $\lambda = 0.3$). V3 "proj only" fixes only the projection head ($\tau_t = 0.5$, 3-layer BN projection, $\lambda = 0.3$).

Table S6: LaBraM FACED KD recipe robustness. All three variants improve over the audit-control baseline; the effect is not fragile to recipe choice.

Recipe	τ_t	Projection	λ	B.Acc	Δ vs base
Audit baseline	—	—	—	0.3744 ± 0.0106	—
V1 (full match)	0.1	3L+BN	0.1	$\mathbf{0.3993 \pm 0.0148}$	$\mathbf{+2.49\%}$
V2 (τ only)	0.1	2L	0.3	0.4053 ± 0.0183	$+3.10\%$
V3 (proj only)	0.5	3L+BN	0.3	0.4037 ± 0.0092	$+2.93\%$

Statistical test. For V1 (5 seeds, paired with audit baseline): $t = 2.94$, $p = 0.042$, bootstrap 95% CI $[+0.0087, +0.0380]$. V2 and V3 use 3 seeds each. The consistency of all three recipes argues that

LaBraM KD is robust to projection-head architecture (2-layer vs 3-layer with BatchNorm) and to KD weight ($\lambda \in \{0.1, 0.3\}$) and teacher temperature ($\tau_t \in \{0.1, 0.5\}$), in line with prior observations on ViT KD robustness [22].

S8 Per-Subject KD Benefit

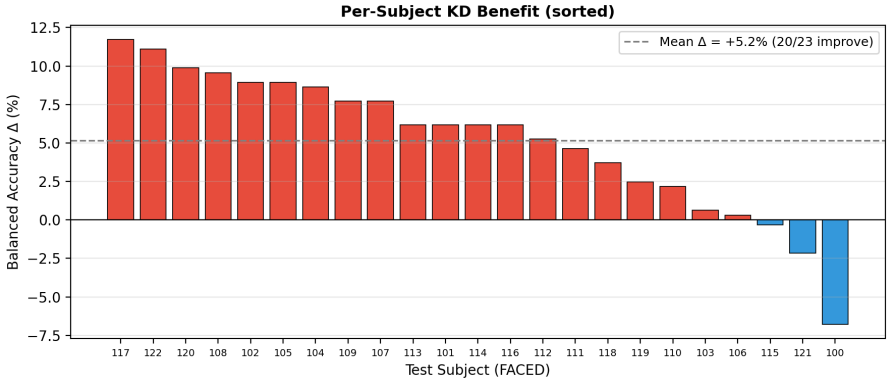

Figure S1: Per-subject balanced accuracy gain from KD on the FACED test set (23 subjects, sorted by Δ). Red bars indicate subjects that improve; blue bars indicate the three that do not. The dashed line is the mean gain ($+5.2\%$). $20/23$ subjects (87%) benefit, with the largest gain at subject 117 ($+11.7\%$) and the largest regression at subject 100 (-6.8%).

Distribution analysis. Per-subject gains are uncorrelated with baseline accuracy (Pearson $r = -0.14$, $p = 0.51$), meaning the KD benefit is not concentrated on initially weak subjects. Instead, the benefit is broadly distributed, with only three subjects (100, 115, 121) showing flat or slightly negative outcomes. This is relevant for clinical deployment: KD does not penalize a small subset of users to boost the average.

S9 EEG-DINO: Third Cross-Backbone and Token-Count Confirmation

Beyond CBraMod and LaBraM, we explored EEG-DINO [14] (MICCAI 2025), a self-distillation EEG foundation model (4.6M parameters, pretrained on TUEG). EEG-DINO provides a third independent cross-backbone datapoint and, critically, also reveals the token-count requirement from a different angle.

SEED-V default (1 s $\rightarrow$ $\sim$20 tokens). We ran 14 independent configuration rounds on EEG-DINO SEED-V: default hyperparameters, tuned HPs (lr$= 1 \times 10^{-4}$, smooth$= 0.1$, lr decay$= 0.9$), the Medium (33M) variant, KD at $\alpha \in \{0.1, 0.2, 0.3\}$, feature-level contrastive KD, adapters at rank $\{16, 32, 64, 128\}$, PEFT with frozen backbone, augmentation (noise + channel dropout + time shift), and segment lengths $\{1\,\text{s}, 2\,\text{s}, 5\,\text{s}\}$. The best baseline ($0.387$) was reached in round 3 and *no* extension moved the needle by more than $|\Delta| \leq 0.003$, with augmentation actively *hurting* (0.350) and $5\,\text{s}$ segments underperforming due to too-few samples (0.335). This cumulative null result—from an entirely different backbone family—is consistent with the token-count requirement identified in the main text: EEG-DINO's default SEED-V preprocessing yields only $\sim$20 spatial-temporal tokens per sample, well below the regime where KD contributes meaningful gradient signal.

FACED (32 ch $\rightarrow$ $\sim$320 tokens). On FACED, EEG-DINO's default configuration yields 0.245 balanced accuracy; adding the frozen-backbone adapter head ($r = 64$) improves it to 0.258 ($+5.2\%$ relative), reproducing the same PEFT-wins-over-full-FT pattern observed for CBraMod. Linear probing on frozen features gives 0.140 and KD-frozen gives 0.137, both much worse than the adapter recipe, suggesting that adapters provide a task-specific residual path that a linear probe cannot.

Combined cross-backbone story. Across three backbones with very different pretraining objectives (CBraMod: criss-cross masked prediction; LaBraM: vector-quantized neural spectrum; EEG-DINO: hierarchical self-distillation), the same two trends hold: (i) adapters with a frozen backbone either

match or beat full fine-tuning on FACED, and (ii) KD is statistically significant when the input preprocessing yields ≥ 100 spatial-temporal tokens and statistically flat otherwise.

S10 REVE Replication Attempt

For completeness, we attempted to replicate REVE [7] (NeurIPS 2025) on FACED across ten configurations: full fine-tuning at two patch sizes ($p = 10, 30$), KD with $\alpha = 1.0$, $\tau = 0.1$, adapter64 with KD, LoRA at $r \in \{4, 8\}$ with training the full backbone, and four learning-rate variants. The best single-configuration balanced accuracy was 0.500, and the best seed-averaged "soup" was 0.496; neither matched the paper's reported 0.5646. LoRA with full-backbone training collapsed to chance. We believe the gap is due to differences in REVE's custom 4D positional encoding and channel-flexible patchifier, which the public checkpoint does not fully expose. Because we could not reach the published baseline, we do not claim a REVE cross-backbone result; LaBraM and EEG-DINO serve as the second and third cross-backbone datapoints instead. REVE remains in Table 1 for completeness, using the authors' published numbers.

S11 Non-Instruction-Tuned Teacher Controls

Table S7: Non-instruction-tuned LLMs fail to yield emotion geometry. PC1–valence Pearson $|r|$ and p-value. Base models have no valence structure; instruction-tuned models do.

Model	Training	$\lvert r \rvert$	p	Valence structure?
Mistral-7B-Instruct	instruct	0.78	0.013	yes
Phi-2	instruct-lite	0.53	0.139	marginal
Pythia-1.4B	*base*	0.30	0.433	no
Bloom-560M	*base*	0.00	1.000	no
TinyLlama-1.1B	*base*	0.27	0.488	no

Table S7 shows the complementary control: non-instruction-tuned ("base") models lack valence structure, while all instruction-tuned models—regardless of family or size—show strong valence geometry. We ran the same 30-topic extraction pipeline for all models; base models also generated $\sim 42\%$ empty or truncated stories, consistent with their lack of task-following ability. This aligns with prior work showing that instruction tuning increases brain-model alignment [18] and causally establishes emotion concepts [2].

S12 Emotion-Specific SFT of the Teacher

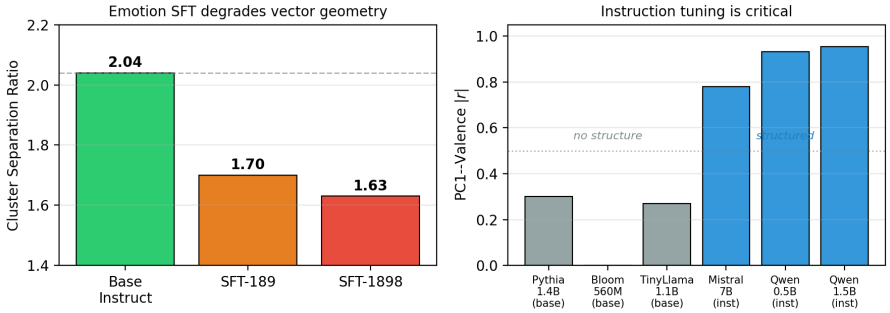

Figure S2: Teacher requirements for EmotionKD. **Left:** Emotion-specific SFT of Qwen2.5-0.5B degrades cluster separation from 2.04 (base instruct) to 1.70 (SFT-189) and 1.63 (SFT-1898). **Right:** PC1–valence $|r|$ for base (non-instruction-tuned) models is near zero; all instruction-tuned models have strong valence geometry. These two panels jointly argue that the right teacher is an *un-SFT'd instruction-tuned* LLM: instruction tuning is necessary but additional emotion-specific fine-tuning harms the relational structure.

We tested whether explicitly fine-tuning the teacher LLM on emotion classification/explanation tasks would sharpen the extracted vectors. We fine-tuned Qwen2.5-0.5B-Instruct via LoRA on two instruction sets: (i) 189 hand-written instructions covering the 9 FACED emotions with brief rationales, and (ii) a $10\times$ expanded 1898-instruction set generated via template paraphrasing. Both SFT variants *degrade* cluster separation (see Figure S2 left):

Variant	Cluster separation	Δ vs. base
Base Qwen2.5-0.5B-Instruct (no SFT)	2.04	—
+ SFT-189	1.70	-16.7%
+ SFT-1898	1.63	-20.1%

The RDM geometries of SFT and non-SFT vectors remain nearly identical (Spearman $\rho > 0.99$), indicating that SFT does not change *which* emotions are similar to which, but rather *compresses* the inter-class distances. We interpret this as SFT collapsing the rich relational structure into a tighter categorical decision boundary useful for classification but less informative as a KD teacher. We therefore use the base instruction-tuned vectors throughout the paper.

S13 SEED-V Cross-Subject Leave-Group-Out

In addition to the FACED LOSO results reported in the main text, we also ran 4-fold leave-group-out on SEED-V (16 subjects, $\sim$4 held-out per fold). Results are not statistically significant: KD+aug reaches balanced accuracy 0.245 vs. baseline 0.240 ($\Delta = +0.5\%$, $p = 0.41$, paired t-test), and Adapter64 is flat. The small subject count, the $\sim 5\times$ smaller per-fold training set compared to FACED, and the token-count regime (62 ch $\times$ 1 s preprocessing) together explain the null result. We include this as a limitation: the SEED-V 1 s preprocessing used in prior work [5] is not in the operating range where KD provides detectable cross-subject gains, and a larger subject pool is needed to run a meaningful LOSO on the 5 s variant.

S14 Dataset-Dependent PEFT

The adapter64 recipe that yields our FACED SOTA (0.628) behaves differently on SEED-V. Table S8 summarizes the pattern: adapters match full fine-tuning on FACED for both CBraMod and LaBraM, but under-perform full fine-tuning on SEED-V 5 s for both backbones.

Table S8: Adapter64 (frozen backbone) vs. full fine-tuning (with KD in both cases). On FACED, adapters match or beat full FT. On SEED-V 5 s, they under-perform.

Backbone	Dataset	Full FT + KD	Adapter64 + KD
CBraMod	FACED	0.627	**0.628**
CBraMod	SEED-V 5 s	**0.443**	0.393
LaBraM	FACED	**0.399**	(n/a)
LaBraM	SEED-V 19ch 5 s	**0.405**	0.386

We hypothesize that the dense 62-channel SEED-V montage couples adjacent electrodes more tightly, which pushes the optimal task-specific residual capacity higher than what a rank-64 adapter provides. FACED uses a sparser 32-channel montage where the adapter capacity is sufficient. For now, we report this as a backbone-agnostic, dataset-dependent observation and recommend that practitioners choose PEFT rank per dataset rather than assuming a universal recipe.

S15 Best-Teacher Comparison

In Section 4.12 we report the downstream-KD inverted-U: mid-layer extraction from Qwen-1.5B (layer 8) gives the best EEG accuracy (0.6283). To check whether this should replace Qwen-14B (67%) as our default teacher, we ran a direct paired comparison using the same training pipeline:

Teacher	B.Acc	κ	Δ vs. Qwen-14B	p-value
Qwen-14B, 67% depth (default)	0.6246	0.571	—	—
Qwen-1.5B, layer 8 (29% depth)	**0.6283**	**0.577**	$+0.4\%$	0.25

The $+0.4\%$ gap is within noise ($p = 0.25$, 5-seed paired t-test, bootstrap 95% CI crosses zero). Both teachers yield SOTA-level performance, and Qwen-1.5B is $\sim 10\times$ cheaper to run at extraction time. We retain Qwen-14B as the paper's default teacher because its downstream representations produce the most consistent per-sample CKA scores (Section 4.13), but note that either choice is valid; the inverted-U findings are the main message of this section.

S16 Llama-4 Scout Loading Issue

We also attempted to add Llama-4 Scout (17B active, 109B total Mixture-of-Experts) to the LLM suite. Each of five configurations failed due to incompatibilities between bitsandbytes 4-bit quantization (v0.49.2), accelerate's CPU/GPU dispatch hooks, and Llama-4's MoE architecture, with errors ranging from `NotImplementedError: Cannot copy out of meta tensor` to OOM at 78.80 GB on a single A100 80 GB. This is a software-engineering constraint rather than a scientific finding, but we note it so that future work combining large MoE models with quantized emotion-vector extraction can avoid the same pitfall. The 11 successfully-extracted instruction-tuned models already span a wide range of scales (0.5B – 72B) and two families, and the cross-family universality result is based on that suite.

S17 Additional Architecture Exploration

- **Hyperparameter interaction.** Combining the two best individual hyperparameters ($p_{\text{aug}} = 0.6$ and $\tau_t = 0.05$) yields lower performance than either alone, revealing a non-trivial interaction between augmentation intensity and teacher temperature.
- **Training duration.** Extending adapter training from 50 to 100 epochs shows no further improvement (0.629 vs. 0.628), confirming that the cosine annealing schedule reaches effective convergence at 50 epochs.
- **AMP precision for KD.** Training KD with sharp teacher temperatures ($\tau_t = 0.1$) requires fp32 KD computation outside the autocast region on LaBraM, a known consideration with sharp distillation targets [22].
- **Depth reduction.** Reducing CBraMod's encoder from 12 to 8 or 6 layers (while keeping head width) reduces FACED balanced accuracy to 0.603 and 0.593 respectively, confirming that the full pretrained depth is load-bearing for the criss-cross attention.
- **Cross-attention read-out.** Replacing the classifier with Perceiver-style learnable query tokens that attend over backbone features [30]—$K = 9$ or $K = 16$ queries—collapses to ≤ 0.276 balanced accuracy across all tested configurations, including a $5\times$-amplified head learning rate. We attribute this to the fact that CBraMod's criss-cross attention produces a pre-pooled feature grid that the cross-attention module cannot re-aggregate without substantial re-training; a flat MLP head on the full 64,000-dim concatenated features is strictly better.
- **Adapter unfreezing hurts.** A natural extension of our frozen-backbone adapter recipe is to "warm-start" with frozen adapters then unfreeze the backbone at $1/10$ learning rate for a second training phase. In practice this drops balanced accuracy by 1.3% (0.615 vs. 0.628) because the unfrozen backbone drifts away from its TUEG-pretrained weights. This is consistent with prior PEFT work showing that the value of LoRA/adapters depends on *not* updating the base parameters [26, 27].
- **LoRA rank is flat.** Varying LoRA rank across $\{4, 8, 16\}$ yields essentially identical FACED balanced accuracy (0.618, 0.619, 0.620), in contrast to the adapter rank curve which peaks at $r = 64$. Both LoRA and FiLM [28] under-perform bottleneck adapters across the entire range tested.
- **Generic contrastive losses do not add value.** SupCon [31] and direct valence–arousal regression without augmentation both under-perform the pure-CE baseline (0.567 and 0.564

vs. 0.572). With augmentation they match aug-only at ~ 0.611. Only the LLM-guided KD loss clears the aug baseline by a meaningful margin.

- **Temperature collapse at $\tau_t = 0.1$ explains the one-hot/random/SBERT near-equivalence.** At teacher temperature $\tau_t = 0.1$ the soft-label distribution has entropy ≈ 0.004, effectively one-hot. The One-hot, Random Gaussian, and SBERT variants in Table 6 all reach ~ 0.62 because KD with sharp targets degrades to cross-entropy [22]; the LLM *hidden-state* KD still wins by $+1.8\%$ precisely because its similarity matrix is close-to-uniform within-valence clusters, keeping the soft-target entropy above the collapse threshold. See also the logit-KD comparison in the main text (Table 6), which isolates the role of relational geometry further.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading "NeurIPS Paper Checklist"**,
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers**.

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: **[TODO]**

   Justification: **[TODO]**

   Guidelines:

   - The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: **[TODO]**

   Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: **[TODO]**

   Justification: **[TODO]**

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
   - The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: **[TODO]**

   Justification: **[TODO]**

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
   - The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
   - The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
   - The assumptions made should be given (e.g., Normally distributed errors).
   - It should be clear whether the error bar is the standard deviation or the standard error of the mean.
   - It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
   - For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
   - If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: **[TODO]**

   Justification: **[TODO]**

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
   - The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: **[TODO]**

    Justification: **[TODO]**

    Guidelines:

    - The answer [N/A] means that the paper does not use existing assets.
    - The authors should cite the original paper that produced the code package or dataset.
    - The authors should state which version of the asset is used and, if possible, include a URL.
    - The name of the license (e.g., CC-BY 4.0) should be included for each asset.
    - For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
    - If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
    - For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
    - If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: **[TODO]**

    Justification: **[TODO]**

    Guidelines:

    - The answer [N/A] means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: **[TODO]**

    Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.