
A Shared Valence Axis Across Modern LLMs and Human EEG: The Saturation Regularity

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A common machine-learning technique adds an *auxiliary loss* that pushes a net-
2 work’s internal features toward a known target direction in feature space—for
3 example, the difference between positive and negative sentiment, or between
4 calming and disturbing emotion. We test 25 such recipes (knowledge distillation,
5 representational-similarity analysis, contrastive losses, low-rank adapters, channel-
6 targeted brain-region losses) on FACED [Chen et al., 2023], a public dataset of
7 EEG recordings from 123 subjects watching 28 emotional video clips, with 9
8 emotion classes. None help: 16 of the 25 produce statistically significant accuracy
9 decrements, and 0 produce gains. We trace this to a regularity we call *saturation*:
10 once a network has converged on the emotion classification task from labels alone,
11 its features already encode the target concept; the auxiliary loss then adds gradient
12 noise rather than information. Saturation is distinct from overfitting (memorising
13 the training set); it is a property of how task and auxiliary supervision interact once
14 both target the same direction in feature space. The same picture suggests a positive
15 prescription: ensemble across the seed-specific differences in the features, rather
16 than supervise the part of the features that has already saturated. Following this pre-
17 scription on FACED yields a new state of the art at **0.6948** balanced accuracy—a
18 +10.5% relative improvement over EMOD [Chen et al., 2025], the previous SOTA
19 on FACED, at 0.6287. Surprisingly, the language-model-derived class prototypes
20 used by the knowledge-distillation step in our recipe are interchangeable with
21 random orthonormal vectors ($\Delta \leq 0.003$ balanced accuracy): the SOTA gain
22 comes from the *shape* of the loss (a 9-class structure imposed on the features),
23 not from the *content* of the prototypes. To probe what the network has saturated
24 on—and to demonstrate it is a real semantic axis, not a noise dimension—we
25 extract one direction explicitly. We write nine short emotion-evocative stories (one
26 per FACED class), pass them through a language model (Qwen3-4B), and take the
27 first principal component of the language model’s late-layer features across the
28 nine stories. We call this direction the *valence axis*. The same direction predicts
29 movie-review sentiment zero-shot on SST-2 [Socher et al., 2013] at AUC 0.832,
30 predicts the EEG response averaged across 123 FACED subjects at $r = 0.87$, and
31 emerges spontaneously in 36 EEG classifiers trained only on the 9 emotion labels.
32 All claims replicate on a second public EEG dataset, SEED-V [Liu et al., 2022b]
33 (Appendix S8).

1 Introduction

34 A common technique in machine learning is to point a network at a target direction in feature
35 space—a known semantic concept such as sentiment polarity, the truthfulness of a statement, or
36 the valence of an emotion—and add an *auxiliary loss* that pulls the network’s penultimate features
37

38 toward this direction. The recipes vary widely. Knowledge distillation (KD) [Hinton et al., 2015]
 39 aligns a student network’s features with those of a teacher network. Representational similarity
 40 analysis (RSA) [Kriegeskorte et al., 2008] aligns the geometry of pairwise distances between two
 41 networks. Contrastive losses [Khosla et al., 2020, van den Oord et al., 2018] pull same-class features
 42 together and push different-class features apart. Parameter-efficient fine-tuning (PEFT) adapters
 43 such as LoRA [Hu et al., 2022] restrict the network’s update to a low-rank direction in weight space.
 44 Channel-targeted topographic losses [Padakanti et al., 2025] bias which input channels matter most.
 45 The shared assumption across all of these methods is that the auxiliary loss provides information the
 46 task labels alone cannot.




Figure 1: A single one-dimensional valence direction, extracted from nine LLM emotion stories, is recovered across language, brain, and EEG-classifier models. (a) Zero-shot probe AUC / $|r|$ of the V-axis (PC1 of nine Qwen3-4B story-class centroids at the penultimate layer) on standard text-sentiment benchmarks; no labels are seen during axis construction. (b) Stimulus-level scatter of LLM V-axis projection vs. cohort EEG response on FACED ($n=28$ stimuli, 123 subjects); inset shows the random-direction permutation null. (c) 36 EEG-classifier checkpoints (CBraMod and EMOD families) plotted as their class-PC1 V-axis $|r|$ (no V-axis loss used in training) against FACED 9-class balanced accuracy; linear fit $r=+0.885$. The same one-dimensional direction explains text sentiment, predicts evoked EEG, and tracks classifier accuracy without sharing data across the three settings. [K: this fig is just sitting here without reference on its own without any grounding. as if hey just look at my nice plot and figure out by yourself.]

47 We test 25 such techniques on FACED [Chen et al., 2023], a public dataset of EEG recordings from
 48 123 subjects watching 28 emotional video clips (32 EEG channels at 250 Hz, 9 emotion classes).
 49 None help. Sixteen produce statistically significant accuracy decrements (change in balanced accuracy
 50 $\Delta\text{BACC} \leq -0.013$, the fraction-correct averaged uniformly across the 9 classes; $p < 0.05$, paired
 51 across five seeds); zero produce gains. Five of the 25 recipes show *monotonic destruction*: the harder
 52 we push the auxiliary loss, the worse accuracy gets, with no positive setpoint to find. The transition
 53 from neutral-to-harmful is sharp: an auxiliary loss that has no measurable effect at a weak baseline
 54 (balanced accuracy ≤ 0.62) becomes a statistically significant negative once the baseline strengthens
 55 (BACC ≥ 0.66). The same intervention can be a no-op on one recipe and a meaningful regression on
 56 the next, separated by a few points of headline accuracy.

57 We call this regularity *saturation*, and it is not the same thing as overfitting. Overfitting is a network
 58 memorising its training set. Saturation is what happens when the network has *already* learned
 59 the concept the auxiliary loss was meant to teach, from the task labels alone. Once the EEG

60 classifier has converged on the 9-class emotion task, its penultimate features are organised so that
 61 each emotion’s trials cluster around a stable mean position; the average direction in feature space
 62 along which positive and negative emotions sit on opposite sides—the *class-mean component*—is
 63 already the target concept. The auxiliary loss can only deform this already-correct direction. The
 64 remaining variation, which lives in the directions *within* each class (the *within-class residual*), is
 65 what discriminates between trials of the same emotion and what an additional seed could plausibly
 66 improve. The auxiliary loss does not reach this within-class residual. Pushing harder on the loss
 67 therefore moves the saturated component further along the target direction without adding any new
 68 discriminative signal, and accuracy drops.

69 The same picture suggests a positive prescription: rather than supervise the saturated component,
 70 ensemble across multiple training seeds whose within-class residuals differ, recovering the gain that
 71 auxiliary supervision cannot. Following this prescription on FACED-9 yields a new state of the art at
 72 **0.6948** balanced accuracy on a 10-checkpoint softmax ensemble (best val-selected single checkpoint
 73 **0.6755**), a +10.5% relative improvement over EMOD [Chen et al., 2025], the previous FACED-9
 74 SOTA at 0.6287. Surprisingly, the class prototypes used by the knowledge-distillation step in our
 75 recipe can be replaced by random orthonormal vectors without measurable degradation (ΔBACC
 76 ≤ 0.003). The SOTA gain is therefore *structural*: the loss imposes a useful 9-class shape on the
 77 features regardless of which direction the prototypes point in.

78 This raises a tension that the rest of the paper resolves directly: if random prototypes give the same
 79 gain as language-model-derived ones, what is a language-model-derived direction doing in this
 80 paper at all? It serves as a probe of *which* concept the network has saturated on, demonstrating that
 81 saturation operates on a real semantic axis rather than on noise. The construction uses nine sentences.
 82 We write one short emotion-evocative story per FACED class, pass each through a language model
 83 (we use Qwen3-4B as the lead model), and take the first principal component of the model’s late-layer
 84 features across the nine resulting story representations. We call this one-dimensional direction the
 85 *valence axis* (or V-axis for short), oriented so that the centroid of the Joy story projects to a positive
 86 value. The same direction predicts movie-review sentiment zero-shot on SST-2 [Socher et al., 2013],
 87 a standard benchmark of $\sim 10\text{k}$ labelled film reviews, at AUC 0.832 from a single Qwen3-4B run. It
 88 predicts the EEG response averaged across 123 FACED subjects: a single ridge regression on all 160
 89 channel-band features matches the V-axis projection at $r = 0.87$ over the 28 video clips ($p < 10^{-9}$).
 90 And it emerges spontaneously in EEG classifiers: across 36 checkpoints trained only on the 9 emotion
 91 labels, the alignment between the network’s class-mean component and the V-axis predicts balanced
 92 accuracy at $r = +0.885$. The probe demonstrates the concept is real (the saturation operates on a
 93 genuine semantic axis); the SOTA prescription shows that, once the network has saturated, only the
 94 loss’s structure matters, not the direction it points in.

95 **Contributions.** (1) The *saturation regularity*: across 25 representation-alignment recipes on
 96 FACED-9, 16 produce statistically significant accuracy decrements and 0 produce gains, with
 97 monotonic destruction in 5 families and a sharp transition in $[0.62, 0.66]$ balanced accuracy. A
 98 direct mechanism check shows that the auxiliary loss moves the saturated class-mean component by
 99 $\Delta|r| = +0.01$ to $+0.36$ but moves the within-class residual by $\sim 10^{-7}$, a numerical zero (§3). (2)
 100 A residual-diversity ensemble that turns the mechanism into a positive prescription: per-checkpoint
 101 within-class residual encoding predicts leave-one-out ensemble contribution at $r = +0.74$, and a
 102 10-checkpoint softmax ensemble reaches FACED-9 SOTA at **0.6948** (§4, §8). (3) The *V-axis probe*:
 103 a one-dimensional valence direction extracted from 9 language-model emotion stories that simulta-
 104 neously predicts text sentiment, predicts cohort EEG response, and emerges spontaneously in EEG
 105 classifiers (§5–§7). All three contributions robustness-tested on a second EEG dataset, SEED-V [Liu
 106 et al., 2022b] (Appendix S8).

107 Limitations and scope statements are catalogued in §9.

108 2 Related Work

109 **Concept-direction supervision.** The interventions we test are drawn from seven years of
 110 representation-alignment work. Knowledge distillation [Hinton et al., 2015, Tian et al., 2020, Park
 111 et al., 2019] trains a student’s penultimate features against a fixed teacher target. Representational
 112 similarity analysis [Kriegeskorte et al., 2008, Sundaram et al., 2024] aligns representational geome-
 113 tries through pairwise similarity matrices. Supervised contrastive losses [Khosla et al., 2020, van den

114 Oord et al., 2018, Radford et al., 2021] pull same-class features together and push different-class
 115 apart. Parameter-efficient adaptation [Hu et al., 2022, Liu et al., 2022a] inserts low-rank or scaling
 116 modules whose updates are localised to a chosen geometry. The shared assumption across these
 117 methods is that the auxiliary signal provides information the task does not. This paper studies a
 118 regime where it does not.

119 **Concept directions in language models.** Arditi et al. [Arditi et al., 2024] showed that refusal in
 120 instruction-tuned LMs is mediated by a single direction in residual-stream activations: ablating it
 121 eliminates refusal, adding it induces it. Earlier work in this family includes [Li et al., 2023, Zou
 122 et al., 2023, Kim et al., 2018]. Concurrent work from Anthropic interpretability [Sofroniew et al.,
 123 2026] identified emotion-concept features in production-scale LMs through sparse autoencoders
 124 and showed that ablating these features changes downstream behaviour, complementing the cohort-
 125 level brain alignment we report here. We use the same PCA-on-class-centroids construction to
 126 extract a one-dimensional valence direction from nine LLM emotion stories, and use it as a probe
 127 of which concept the network has saturated on rather than as a target for steering. The transfers we
 128 report (SST-2 zero-shot, cohort EEG, EEG-classifier convergence) are direction-existence claims, not
 129 feature-importance claims for biological emotion processing.

130 **Brain-LM alignment.** Toneva and Wehbe [Toneva and Wehbe, 2019] first showed that mid-layer
 131 transformer activations predict fMRI responses to narrative text; Schrimpf et al. [Schrimpf et al.,
 132 2021], Goldstein et al. [Goldstein et al., 2022], and Caucheteux and King [Caucheteux and King,
 133 2022] extended the picture across MEG, ECoG, and large model families. Huh et al. formalised
 134 the emerging picture as the *Platonic Representation Hypothesis* [Huh et al., 2024]. We touch this
 135 literature through the probe section: an LLM-derived direction predicts cohort EEG, and EEG
 136 classifiers find the LLM-side direction in return. Our load-bearing claim is the saturation regularity,
 137 not the alignment itself.

138 **Frontal-alpha asymmetry.** Davidson’s frontal-alpha asymmetry [Davidson, 1992, Coan and Allen,
 139 2004, Davidson, 2004] remains the dominant electrophysiological account of approach/withdrawal
 140 emotion, developed primarily on static stimuli. Our cohort signal on video-evoked emotion is
 141 posterior-dominant, with frontal asymmetries that survive in direction at smaller magnitude—an
 142 additive scope statement for video paradigms rather than a refutation of FAA.

143 **Ensemble theory.** The two-tier mechanism we identify—a saturated class-mean basin plus a seed-
 144 specific within-class residual—is a representation-level form of bias–variance [Geman et al., 1992,
 145 Krogh and Vedelsby, 1995] and linear-mode connectivity [Frankle et al., 2020, Wortsman et al.,
 146 2022]. Our SOTA prescription operationalises this geometry: ensembling recovers signal in the
 147 residual subspace, which auxiliary supervision on the saturated basin cannot reach.

148 **Baselines.** The FACED-9 baselines we compare against are CBraMod [Wang et al., 2025]
 149 (0.572), EmotionKD [Liu et al., 2023] (0.628), and EMOD [Chen et al., 2025] (0.6287), with
 150 REVE [El Ouahidi et al., 2025], LaBraM [Jiang et al., 2024], and EEGPT [Wang et al., 2024] as
 151 foundation-model references.

152 3 The Saturation Regularity

153 We test whether auxiliary supervision toward a known concept direction helps a strong EEG-emotion
 154 classifier. The setup uses an EMOD [Chen et al., 2025] backbone (the previous FACED-9 SOTA),
 155 augmented to balanced accuracy ≈ 0.66 via the recipe described in §8. Each of the 25 interventions
 156 adds an auxiliary loss that pulls the network’s penultimate features toward a target direction. The
 157 directions span LLM-derived class prototypes (the V-axis of §5), CLIP-text embeddings, and frontal-
 158 channel masks matched to the analytical valence ceiling. The loss families span knowledge distillation,
 159 RSA, contrastive losses, PEFT adapters, curriculum schedules, multi-LLM ensembles, channel-
 160 targeted topographic losses, and anger-weighted variants (full 25-row table in Appendix S10).

161 **Result.** None help. Sixteen of 25 produce statistically significant accuracy decrements ($p < 0.05$,
 162 paired across five seeds); zero produce gains. Five families show *monotonic destruction*: the
 163 harder we push the auxiliary loss, the worse accuracy gets, with no positive setpoint. Even the

164 anger-weighted target that analytically maximises the V-axis ceiling on FACED is among the worst
 165 ($\Delta\text{BACC} = -0.054$): supplying the strongest possible target makes things strictly worse, not better.

166 **The transition is sharp and unidirectional in baseline strength.** An intervention whose effect
 167 lies within seed noise on a weak baseline ($\text{BACC} \leq 0.62$) becomes a statistically significant negative
 168 on the strong recipe ($\text{BACC} \geq 0.66$). The transition sits cleanly in the band $[0.62, 0.66]$ on FACED-9
 169 (we make no claim about universal cut-offs).

170 **Saturation: an empirical regularity.** *Let f_θ be a classifier whose penultimate*
 171 *features encode a target concept direction v at strength $\rho(\theta)$. We observe an*
 172 *empirical threshold ρ^* above which adding any auxiliary loss $\mathcal{L}_v$ at $\lambda > 0$ produces*
 173 *$\mathbb{E}[\Delta\text{BACC}] \leq 0$ across the 25 families tested (16/25 at $p < 0.05$, 0/25 positive).*

174 **Saturation is not overfitting.** Overfitting is memorisation of the training set. Saturation is a
 175 property of how task and auxiliary supervision interact once both target the same direction in feature
 176 space. The mechanism check below makes the distinction concrete.

177 **Mechanism check.** The auxiliary loss does push the network’s class-mean features further along
 178 the target direction (alignment increases by $\Delta|r| = +0.01$ to $+0.36$ across the 25 families), but it
 179 leaves the within-class residual—the part of the features that actually distinguishes one trial of an
 180 emotion from another—at a numerical zero ($\sim 10^{-7}$). The basin the network was using to classify
 181 gets reshaped along the loss direction; the load-bearing orthogonal residual subspace receives no
 182 compensating gradient. Twenty-five interventions reduce to one principle: *a saturated representation*
 183 *is an unhelpful supervision target.* Full table, monotonic-destruction curves, transition table, and the
 184 Path-B failure (mixing V-axis-trained checkpoints into the SOTA pool, $\Delta \in [-0.0193, -0.0145]$)
 185 are in Appendix S10–S11.

186 4 Mechanism: Saturated Basin, Load-Bearing Residual

187 §3 reported that V-axis supervision moves the class-mean component of the network’s features by up
 188 to $+0.36$ but moves the within-class residual by a numerical zero. We take that decomposition as the
 189 operative geometry and ask what each subspace actually does for the trained network.

190 **Different seeds find the same axis.** Across 36 EEG-emotion checkpoints from two foundation-
 191 model backbones (CBraMod [Wang et al., 2025] and EMOD [Chen et al., 2025]) and six recipe
 192 variants, balanced accuracy correlates with V-axis encoding strength at $r=+0.885$ in the class-
 193 mean subspace ($p=7.8 \times 10^{-13}$) and $r=+0.738$ in the orthogonal residual (Figure 2). A 1000-draw
 194 matched-norm random-direction null places the observed correlation at the 93.5th percentile, so
 195 the convergence is V-axis-specific. The class-mean basin is essentially saturated across seeds
 196 ($|r| \in [0.60, 0.77]$): every well-trained checkpoint finds the same direction.

197 **The residual is what averaging recovers.** Within-class residual encoding predicts ensemble
 198 contribution: across the 10 checkpoints used to build our SOTA ensemble (§8), per-checkpoint
 199 residual strength predicts leave-one-out contribution at $r=+0.74$ ($p=0.014$). Bootstrap CIs and
 200 per-deletion ranges are in Appendix S12.

201 **Residual Contribution Regularity.** *For a saturated EEG classifier with within-*
 202 *class residual encoding of the V-axis at strength ρ_k , the leave-one-out ensemble*
 203 *contribution scales linearly with ρ_k ($p=0.014$, $n=10$, FACED 9-class). $n=10$ is*
 204 *small; we report this as an empirical regularity rather than a law.*

205 **Direct test: ablate the residual at inference time.** A directional ablation [Arditi et al., 2024]—
 206 projecting the LLM-derived V-axis residual out of the classifier’s penultimate features at inference
 207 time, without retraining—drops accuracy on every one of the 10 SOTA-pool checkpoints (mean
 208 $\Delta\text{BACC} = -0.0157$, mean $z \approx 7.7$ above a matched random-direction null). The residual carries the
 209 load. This decomposition—saturated basin, load-bearing residual—is the operative geometry that §3
 210 predicted from the failure of V-axis supervision and that §8 turns into a positive prescription.




Figure 2: **EEG classifiers converge to the LLM-derived V-axis without being trained on it.** (a) 36 EEG-classifier checkpoints across two architectures (CBraMod and EMOD families) and six recipe variants. x -axis: absolute correlation between class-PC1 of the penultimate features and the LLM V-axis (no V-axis loss used in training). y -axis: FACED 9-class balanced accuracy. Linear fit $r = +0.885$ ($p = 7.8 \times 10^{-13}$, $n = 36$). (b) Matched-norm random-direction null: 1000 random unit directions in feature space, each producing its own across-checkpoint correlation. The V-axis sits at the 93.5th percentile (one-sided $p = 0.066$): the convergence is V-axis-specific. The within-class residual alignment carries the statistically robust signal ($r = +0.738$, $p = 3 \times 10^{-7}$; §4).

5 Probing the Saturated Concept: A Valence Direction in LLMs

To test that the saturation regularity operates on a real semantic axis rather than a noise dimension, we extract one explicitly. The construction uses nine sentences. Given a language model, we author one short emotion-evocative story per FACED class and generate 50 paraphrases per story by independent LLM calls. For each class we average the language model’s last-token activation (at the penultimate layer) over all 50 paraphrases; the resulting vector is the class *centroid*—a single point in feature space that summarises the model’s representation of that emotion. Principal component analysis (PCA) on the nine centroids gives the V-axis as the first principal component, oriented so the Joy centroid projects positive. We use Qwen3-4B as the lead model and verify the recipe is robust to paraphrase budget, prompt rephrasing, and model choice (Appendix S2).

The same direction across modalities. A single Qwen3-4B run produces a direction that simultaneously recovers human affect across three settings that share no training data: language sentiment, brain responses, and visual scenes (Table 1). All three numbers come from *projecting* onto the LLM-derived axis without any target-modality fine-tuning—a zero-shot transfer from nine LLM sentences to (i) a 9k-sentence sentiment benchmark, (ii) the cohort-averaged EEG of 123 subjects, and (iii) a 900-image affective-norm dataset.

Modality	Benchmark / signal	Score	Reference
Text	SST-2 sentiment AUC (zero-shot)	0.832	matches sup. LR on 5k examples
Brain	Cohort EEG ridge r ($n=123$ subj.)	0.87	$p < 10^{-9}$, see §6
Vision	OASIS [Kurdi et al., 2017] valence r (CLIP-text)	0.869	beats sup. ridge

Table 1: Cross-modal external validity of a single Qwen3-4B V-axis. Nine LLM sentences produce a direction that recovers human affect on text sentiment, video-evoked EEG, and natural-image valence—without any target-modality training. Full sentiment table (IMDB, TweetEval, Rotten Tomatoes), lexicon recoveries (Hu & Liu, Warriner), and multilingual results are in Appendix S3.

The LLM-side direction is a strong sentiment classifier: 0.832 AUC on SST-2 matches a supervised logistic regression on 5k in-domain examples (0.837), beats an SBERT prototype-cosine baseline (0.793; Reimers and Gurevych, 2019), and trails a 355M RoBERTa-large-MNLI zero-shot entailment

230 classifier (0.912; Liu et al., 2019) by 0.08—all from nine sentences of supervision and no labelled
231 target-modality data.

232 **The same direction across families.** We extract the V-axis from 14 language models from 560M to
233 32B parameters: top-tier alignment ($r > 0.85$ against a behavioural valence reference) holds for all six
234 Qwen3 sizes plus Mistral-7B; Llama-4-Scout, Gemma-27B, and Gemma-4 sit in a middle tier; three
235 older small models (Pythia, TinyLlama, BLOOM) fall below detection threshold (per-LLM table
236 in Appendix S3). Within the Qwen3 family, the V-axis is essentially scale-invariant (off-diagonal
237 $r = +0.995$ across 15 pairs); across families, correlations are looser ($r = +0.585$ over 48 pairs). The
238 direction is a property of *modern* LM training rather than of parameter count.

239 **Specificity.** The recipe is concept-generic (20 further concepts produce a working axis, 17 exceed
240 AUC 0.95; Appendix S4). Nonce-word ablation drops SST-2 to chance and random-Gaussian
241 directions give chance lexicon recovery, so the V-axis is a real semantic direction, not a noise artefact
242 (Appendix S6). An Arditi-style same-model ablation [Arditi et al., 2024] on Qwen3-4B’s own features
243 collapses zero-shot SST-2 from 0.832 to chance (0.505): the V-axis is the direction Qwen3-4B uses
244 for sentiment, not a post-hoc projection.

245 6 The Probe Predicts Cohort EEG

246 We use FACED [Chen et al., 2023]: 123 subjects watching 28 emotional video clips, 32 EEG
247 channels at 250 Hz. From the cohort EEG averaged over subjects and time, we fit a single *ridge*
248 *regression* (a linear model with L_2 regularisation) predicting the V-axis projection from the 160
249 channel-band features. The ridge prediction matches the LLM-derived V-axis at $r = +0.87$ ($n = 28$
250 stimuli, $p < 10^{-9}$; Figure 3). Three explicit controls support the brain-LM correlation. (i) Replacing
251 the V-axis with a matched-norm random direction drops the same ridge to $r = +0.07$ ($p = 0.73$): the
252 alignment is V-axis-specific, not a generic property of the ridge. (ii) Split-half subject reliability on
253 per-stimulus valence ratings is $r = +0.99$: the cohort signal we predict is highly reliable. (iii) The
254 same EEG ridge predicts the human-rated valence of each clip at $r = +0.86$, matching the V-axis
255 prediction to within 0.01—the brain signal we extract is at the same strength a panel of human raters
256 would produce. Multi-comparison correction and other controls are in Appendix S7.

257 The same picture survives across all 14 LLMs (Qwen3 family mean $r = +0.796 \pm 0.011$ irrespective
258 of size) and on the held-out SEED-V dataset [Liu et al., 2022b], where the FACED-derived V-axis
259 ranks the five SEED-V emotions at $r = +0.96$ without retraining. Appendix S8 re-derives the V-axis
260 from scratch on SEED-V and re-tests every claim.

261 7 Brain Topography: A Posterior-Visual Scope Statement

262 The cohort signal is posterior-dominant: region-mean correlation is strongest at occipital electrodes
263 ($|r| = 0.21$) and weakest at frontal ($|r| = 0.16$), with the largest single cell at PO3/ γ ($r = +0.48$;
264 Figure 4). The signal is carried by an *anger-versus-warm-positive* contrast on 9 of the 28 video
265 clips: cohort $r = +0.87$ at PO3/ γ for Anger, Amusement, and Tenderness, against $r = -0.02$ on the
266 remaining 19 mid-valence stimuli. Removing Anger alone shrinks the cohort correlation to +0.33.
267 Full per-region tables, the Simpson’s-paradox per-subject caveat, time-resolved late-positive-potential
268 (LPP) window dynamics, and theta-gamma phase-amplitude coupling tests are in Appendix S9.

269 **This is not a refutation of frontal-alpha asymmetry (FAA).** Davidson’s FAA hypothesis [David-
270 son, 1992, Coan and Allen, 2004] was developed on static stimuli; the cohort signal we report is on
271 dynamic 28-second video. Right-minus-left frontal-alpha asymmetry remains positive in our data
272 ($+0.006$ to $+0.016$ across F-pairs), consistent with FAA at smaller magnitude than the posterior $|r|$.
273 We claim only that, on dynamic video, the cohort signal aligned with the LLM-derived direction is
274 carried by visual-content processing of high-arousal clips—not that the V-axis subsumes FAA or
275 that posterior visual cortex is the seat of valence. Within-subject FAA is a separate question this
276 cohort-level analysis does not address.




Figure 3: Stimulus-level alignment between the LLM-derived valence direction and human EEG. Each point is one of the 28 FACED video stimuli; x is the LLM V-axis projection of the textual stimulus description, y is the single ridge-regression prediction of that projection from all 160 cohort channel-band features (32 channels $\times$ 5 frequency bands, averaged over 123 subjects and over each clip). Cohort Pearson $r = +0.87$ ($n=28$, $p < 10^{-9}$). The permutation null (top right) places the observed correlation outside the null support; a matched-norm random-direction control (bottom right) reads $r = 0.07$. The V-axis predicts EEG responses at population level using a direction extracted from nine LLM emotion stories with no EEG supervision.

8 Two-Tier Ensemble: The Positive Prescription

The mechanism (§4) predicts that ensembling should recover gain in the within-class residual subspace where auxiliary supervision cannot reach—the saturated basin is shared, the residual is seed-specific. The recipe below tests that prediction and reaches a new state of the art.

Method	BACC
CBraMod [Wang et al., 2025]	0.572
EmotionKD [Liu et al., 2023]	0.628
EMOD [Chen et al., 2025]	0.6287
EMOD+aug+KD+ $d=6$ (5-seed)	.658 $\pm$.010
Best single (val-sel)	0.6755
10-ckpt ensemble (SOTA)	0.6948

(a) Headline. +0.0661 abs. (+10.5%) over EMOD, the previous FACED-9 SOTA.

Recipe step	BACC	Δ
EMOD ($d=3$) replication	.619 $\pm$.004	—
+ augmentation	.634 $\pm$.004	+.015
+ aug + KD	.644 $\pm$.007	+.025*
+ depth doubling ($d=6$)	.658 $\pm$.007	+.014
+ longer training ($e=150$)	.658 $\pm$.010	.000
5-seed e100 ensemble	.6798	+.022
10-ckpt e100+e150 (SOTA)	.6948	+.037

(b) Recipe cascade (*super-add., $p = 3 \times 10^{-4}$).

Table 2: New FACED 9-class SOTA. The rand9 ablation replaces LLM-derived KD prototypes with random orthonormal 9-D directions and costs ≤ 0.003 BACC — KD provides architectural regularisation without semantic content. Visual cascade (LF7) in Appendix S13.

The cascade starts at the EMOD replication (0.6194 ± 0.004 , within seed noise of the published 0.6287 [Chen et al., 2025]) and walks up in six steps. Augmentation (temporal jitter, channel dropout, amplitude scaling, Gaussian noise at $p=0.6$) contributes the largest single jump (+0.0149).




Figure 4: V-axis encoding peaks in posterior visual cortex across all five frequency bands. (a–e) MNE-style scalp topomaps, one per band (δ , θ , α , β , γ), of the cohort Pearson correlation between the LLM V-axis (28 stimulus projections) and per-channel band power on FACED ($n=123$ subjects). Diverging colourmap: positive (red) vs. negative (blue) correlation; symmetric limits across bands. Gold circles mark the top channel per band; text under each topomap names that channel and its r value (strongest cell P03/ γ , $r=+0.48$). The colourbar (centred) gives Pearson r ; the small head diagram (right) shows the 32-channel 10–20 montage used in all five panels (A1/A2 mastoids excluded from plotting). (f) Region-mean $|r|$ aggregated over all five bands: occipital (0.21) > parietal (0.18) > central (0.18) > frontal (0.16). The cohort signal is posterior-dominant for video-evoked emotion—a paradigm-level finding for dynamic stimuli rather than a refutation of the static frontal-alpha asymmetry literature, which we replicate qualitatively in direction at smaller magnitude.

284 Knowledge distillation through 9-D LLM-derived class prototypes adds +0.0096; a rand9 ablation
 285 (identical recipe with the 9 LLM prototypes replaced by random orthonormal directions) costs at
 286 most 0.003 BACC, paired $p > 0.5$, so KD imposes a 9-direction frame on the features rather than
 287 transferring the LLM’s emotion concept. Depth-doubling to $d=6$ adds +0.0142; longer training
 288 ($e=150$) does not move single-seed but supplies the diversity axis the ensemble exploits.

289 **Two-tier ensemble theory.** The 10 $d=6$ checkpoints share an identical class-PC1 basin ($|r| \in$
 290 $[0.60, 0.77]$ across all 10 seeds): the first principal component of the per-class mean features points in
 291 the same direction. Ensemble gain is concentrated in the orthogonal within-class residual ($r=+0.743$,
 292 $p=0.014$): each seed encodes the residual differently, so averaging cancels seed-specific noise without
 293 disturbing the basin. Mixing $e=100$ and $e=150$ ties at 0.6581 on single-seed BACC but provides
 294 the diversity axis the ensemble exploits: the mixed pool reaches 0.6948 versus 0.6798 for $e=100$
 295 alone. Headroom-monotonicity (near-zero on MNIST), the 10-checkpoint plateau, and the failure
 296 of cross-architecture mixing confirm intra-architecture training-length diversity as the operative
 297 mechanism (Appendix S13).

298 9 Discussion: Limitations and Future Work

299 **Limitations and future work.** Saturation is empirical, not analytical: the boxed statement describes
 300 rather than proves, the FACED-9 transition at $[0.62, 0.66]$ BACC is dataset-specific in magnitude,
 301 and SEED-V is not large enough to localise its own cliff. The probe transfers (SST-2, cohort
 302 EEG, classifier convergence) are direction-existence claims, not neuroscience claims; the posterior
 303 topography is a paradigm statement for video-evoked emotion, not a refutation of frontal-alpha
 304 asymmetry [Davidson, 1992]. The KD prototypes are interchangeable with random orthonormal
 305 directions: the SOTA gain is structural, not semantic, and the probe and SOTA findings are reconciled
 306 only via saturation. The most direct positive test for the regularity is per-subject residual adaptation
 307 (Appendix S16); locating the cliff on a non-emotion task is the obvious generalisation.

References

- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, 2024.
- Ryan T Canolty and Robert T Knight. The functional role of cross-frequency coupling. *Trends in Cognitive Sciences*, 14(11):506–515, 2010. doi: 10.1016/j.tics.2010.09.001.
- Charlotte Caucheteux and Jean-Rémi King. Brains and algorithms partially converge in natural language processing. *Communications Biology*, 5:134, 2022. doi: 10.1038/s42003-022-03036-1.
- Jingjing Chen, Xiaowei Wang, Chen Huang, et al. A large finer-grained affective computing EEG dataset. *Scientific Data*, 10:740, 2023. doi: 10.1038/s41597-023-02650-w.
- Yuning Chen, Sha Zhao, Shijian Li, and Gang Pan. EMOD: A unified EEG emotion representation framework leveraging v-a guided contrastive learning. *arXiv preprint arXiv:2511.05863*, 2025. To appear at AAAI 2026.
- James A Coan and John J B Allen. Frontal EEG asymmetry as a moderator and mediator of emotion. *Biological Psychology*, 67(1-2):7–50, 2004. doi: 10.1016/j.biopsycho.2004.03.002.
- Maurizio Codispoti, Andrea De Cesare, and Vera Ferrari. Alpha-band oscillations and emotion: A review of studies on picture perception. *Psychophysiology*, 60(11):e14438, 2023. doi: 10.1111/psyp.14438.
- Richard J Davidson. Anterior cerebral asymmetry and the nature of emotion. *Brain and Cognition*, 20(1):125–151, 1992. doi: 10.1016/0278-2626(92)90065-T.
- Richard J Davidson. Well-being and affective style: neural substrates and biobehavioural correlates. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 359(1449):1395–1411, 2004. doi: 10.1098/rstb.2004.1510.
- Yassine El Ouahidi, Jonathan Lys, Philipp Thölke, Nicolas Farrugia, Bastien Padeloup, Vincent Gripon, Karim Jerbi, and Giulia Lioi. REVE: A foundation model for EEG — adapting to any setup with large-scale pretraining on 25,000 subjects. In *Advances in Neural Information Processing Systems*, 2025. arXiv:2510.21585.
- Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M Roy, and Michael Carbin. Linear mode connectivity and the lottery ticket hypothesis. In *International Conference on Machine Learning*, 2020. arXiv:1912.05671.
- Stuart Geman, Elie Bienenstock, and René Doursat. Neural networks and the bias/variance dilemma. *Neural Computation*, 4(1):1–58, 1992. doi: 10.1162/neco.1992.4.1.1.
- Ariel Goldstein, Zaid Zada, Eliav Buchnik, Mariano Schain, Amy Price, Bobbi Aubrey, Samuel A Nastase, Amir Feder, Dotan Emanuel, Alon Cohen, Aren Jansen, Harshvardhan Gazula, Gina Choe, Aditi Rao, Catherine Kim, Colton Casto, Lora Fanda, Werner Doyle, Daniel Friedman, Patricia Dugan, Lucia Melloni, Roi Reichart, Sasha Devore, Adeem Flinker, Liat Hasenfratz, Omer Levy, Avinatan Hassidim, Michael Brenner, Yossi Matias, Kenneth A Norman, Orrin Devinsky, and Uri Hasson. Shared computational principles for language processing in humans and deep language models. *Nature Neuroscience*, 25(3):369–380, 2022. doi: 10.1038/s41593-022-01026-4.
- Greg Hajcak, Annmarie MacNamara, and Doreen M Olvet. Event-related potentials, emotion, and emotion regulation: An integrative review. *Developmental Neuropsychology*, 35(2):129–155, 2010. doi: 10.1080/87565640903526504.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Minyoung Huh, Brian Cheung, Tongzhou Wang, and Phillip Isola. Position: The platonic representation hypothesis. In *International Conference on Machine Learning*, 2024. arXiv:2405.07987.
- Wei-Bang Jiang, Li-Ming Zhao, and Bao-Liang Lu. Large brain model for learning generic representations with tremendous EEG data in BCI. In *International Conference on Learning Representations*, 2024. arXiv:2405.18765.
- Prannay Khosla, Piotr Teterwak, Chen Wang, Aaron Sarna, Yonglong Tian, Phillip Isola, Aaron Maschinot, Ce Liu, and Dilip Krishnan. Supervised contrastive learning. In *Advances in Neural Information Processing Systems*, 2020.

359 Been Kim, Martin Wattenberg, Justin Gilmer, Carrie Cai, James Wexler, Fernanda Viegas, and Rory Sayres.
360 Interpretability beyond feature attribution: Quantitative testing with concept activation vectors (TCAV). In
361 *International Conference on Machine Learning*, 2018. arXiv:1711.11279.

362 Nikolaus Kriegeskorte, Marieke Mur, and Peter Bandettini. Representational similarity analysis — connecting
363 the branches of systems neuroscience. *Frontiers in Systems Neuroscience*, 2, 2008.

364 Anders Krogh and Jesper Vedelsby. Neural network ensembles, cross validation, and active learning. In *Advances*
365 *in Neural Information Processing Systems*, pages 231–238, 1995.

366 Benedek Kurdi, Shayn Lozano, and Mahzarin R Banaji. Introducing the open affective standardized image set
367 (OASIS). *Behavior Research Methods*, 49(2):457–470, 2017. doi: 10.3758/s13428-016-0715-3.

368 Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention:
369 Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems*,
370 2023. arXiv:2306.03341.

371 Haokun Liu, Derek Tam, Mohammed Muqeeth, Jay Mohta, Tenghao Huang, Mohit Bansal, and Colin Raffel.
372 Few-shot parameter-efficient fine-tuning is better and cheaper than in-context learning. In *Advances in Neural*
373 *Information Processing Systems*, 2022a. arXiv:2205.05638.

374 Wei Liu, Jie-Lin Qiu, Wei-Long Zheng, and Bao-Liang Lu. Comparing recognition performance and robustness
375 of multimodal deep learning models for multimodal emotion recognition. *IEEE Transactions on Cognitive*
376 *and Developmental Systems*, 14(2):715–729, 2022b. doi: 10.1109/TCDS.2021.3071170.

377 Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke
378 Zettlemoyer, and Veselin Stoyanov. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv*
379 *preprint arXiv:1907.11692*, 2019.

380 Yucheng Liu, Ziyu Jia, and Haichao Wang. EmotionKD: A cross-modal knowledge distillation framework for
381 emotion recognition based on physiological signals. In *Proceedings of the 31st ACM International Conference*
382 *on Multimedia*, pages 6122–6131, 2023. doi: 10.1145/3581783.3612277.

383 Andrew L Maas, Raymond E Daly, Peter T Pham, Dan Huang, Andrew Y Ng, and Christopher Potts. Learning
384 word vectors for sentiment analysis. In *Proceedings of the 49th Annual Meeting of the Association for*
385 *Computational Linguistics: Human Language Technologies*, pages 142–150, 2011.

386 Niklas Muennighoff, Thomas Wang, Lintang Sutawika, Adam Roberts, Stella Biderman, Teven Le Scao,
387 M Saiful Bari, Sheng Shen, Zheng Xin Yong, Hailey Schoelkopf, Xiangru Tang, Dragomir Radev, Alham Fikri
388 Aji, Khalid Almubarak, Samuel Albanie, Zaid Alyafeai, Albert Webson, Edward Raff, and Colin Raffel.
389 Crosslingual generalization through multitask finetuning. In *Proceedings of the 61st Annual Meeting*
390 *of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15991–16111, 2023.
391 arXiv:2211.01786.

392 Matthias M Müller, Andreas Keil, Thomas Gruber, and Thomas Elbert. Processing of affective pictures modulates
393 right-hemispheric gamma band EEG activity. *Clinical Neurophysiology*, 110(11):1913–1920, 1999. doi:
394 10.1016/S1388-2457(99)00151-0.

395 Sriji Padakanti, Khushbu Pahwa, Radhika Mamidi, Bapi Raju Surampudi, Manish Gupta, and Subba Reddy
396 Oota. Aligning text/speech representations from multimodal models with MEG brain activity during listening.
397 In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025. ACL
398 Anthology 2025.emnlp-main.1748.

399 Bo Pang and Lillian Lee. Seeing stars: Exploiting class relationships for sentiment categorization with respect to
400 rating scales. In *Proceedings of the 43rd Annual Meeting of the Association for Computational Linguistics*
401 *(ACL’05)*, pages 115–124, 2005.

402 Wonpyo Park, Dongju Kim, Yan Lu, and Minsu Cho. Relational knowledge distillation. In *IEEE/CVF Conference*
403 *on Computer Vision and Pattern Recognition*, 2019.

404 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry,
405 Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable
406 visual models from natural language supervision. In *International Conference on Machine Learning*, 2021.

407 Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using siamese BERT-networks. In
408 *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, 2019.

409 Sara Rosenthal, Noura Farra, and Preslav Nakov. SemEval-2017 task 4: Sentiment analysis in Twitter. In
410 *Proceedings of the 11th International Workshop on Semantic Evaluation (SemEval-2017)*, pages 502–518,
411 2017.

412 Martin Schrimpf, Idan A Blank, Greta Tuckute, Carina Kauf, Eghbal A Hosseini, Nancy Kanwisher, Joshua B
413 Tenenbaum, and Evelina Fedorenko. The neural architecture of language: Integrative modeling converges on
414 predictive processing. *Proceedings of the National Academy of Sciences*, 118(45):e2105646118, 2021. doi:
415 10.1073/pnas.2105646118.

416 Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Ng, and Christopher
417 Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the*
418 *2013 Conference on Empirical Methods in Natural Language Processing*, pages 1631–1642, 2013.

419 Nicholas Soironiew, Isaac Kauvar, William Saunders, Runjin Chen, Tom Henighan, Sasha Hydrice, Craig Citro,
420 Adam Pearce, Julius Tarnig, Wes Gurnee, Joshua Batson, Sam Zimmerman, Kelley Rivoire, Kyle Fish, Chris
421 Olah, and Jack Lindsey. Emotion concepts and their function in a large language model. *Transformer Circuits*,
422 2026. URL <https://transformer-circuits.pub/2026/emotions/index.html>.

423 Shobhita Sundaram, Stephanie Fu, Lukas Muttenthaler, Netanel Y Tamir, Lucy Chai, Simon Kornblith, Trevor
424 Darrell, and Phillip Isola. When does perceptual alignment benefit vision representations? In *Advances in*
425 *Neural Information Processing Systems*, 2024. arXiv:2410.10817.

426 Yonglong Tian, Dilip Krishnan, and Phillip Isola. Contrastive representation distillation. In *International*
427 *Conference on Learning Representations*, 2020. arXiv:1910.10699.

428 Mariya Toneva and Leila Wehbe. Interpreting and improving natural-language processing (in machines) with
429 natural language-processing (in the brain). In *Advances in Neural Information Processing Systems*, 2019.
430 arXiv:1905.11833.

431 Parsa Vahidi, Omid G Sani, and Maryam M Shanechi. Modeling and dissociation of intrinsic and input-driven
432 neural population dynamics underlying behavior. *Proceedings of the National Academy of Sciences*, 121(7):
433 e2212887121, 2024. doi: 10.1073/pnas.2212887121.

434 Aäron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding.
435 *arXiv preprint arXiv:1807.03748*, 2018.

436 Guangyu Wang, Wenchao Liu, Yuhong He, Cong Xu, Lin Ma, and Haifeng Li. EEGPT: Pretrained transformer
437 for universal and reliable representation of EEG signals. In *Advances in Neural Information Processing*
438 *Systems*, 2024.

439 Jiquan Wang, Sha Zhao, Zhiling Luo, Yangxuan Zhou, Haiteng Jiang, Shijian Li, Tao Li, and Gang Pan.
440 CBraMod: A criss-cross brain foundation model for EEG decoding. In *International Conference on Learning*
441 *Representations*, 2025. arXiv:2412.07236.

442 Mitchell Wortsman, Gabriel Ilharco, Samir Yitzhak Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes, Ari S
443 Morcos, Hongseok Namkoong, Ali Farhadi, Yair Carmon, Simon Kornblith, and Ludwig Schmidt. Model
444 soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time.
445 In *International Conference on Machine Learning*, 2022. arXiv:2203.05482.

446 Xiang Zhang, Junbo Zhao, and Yann LeCun. Character-level convolutional networks for text classification. In
447 *Advances in Neural Information Processing Systems*, pages 649–657, 2015. Yelp Polarity dataset originated
448 in this paper.

449 Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin,
450 Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J Byun, Zifan Wang,
451 Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J Zico Kolter, and Dan Hendrycks.
452 Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*,
453 2023.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state five claims supported by experiments in the paper: (i) a universal valence axis (V-axis) extractable from 14 LLMs; (ii) cohort EEG–LLM correlation of $r = 0.87$ ($p < 10^{-9}$) on FACED, replicated at $r = +0.62$ on SEED-V; (iii) cross-architecture convergence ($r=+0.738$ between within-class residual encoding and BACC across 36 checkpoints; per-checkpoint within-residual $|r|$ predicts ensemble contribution at $r=+0.74$, $p=0.014$); (iv) a saturation regularity documented across 25 alignment-intervention families; (v) a new FACED 9-class SOTA at 0.6948 BACC ensemble / 0.6755 single checkpoint. Each claim is supported in Sections 5–8 with its primary table or figure.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: A “What this result is not” boxed paragraph in Section 1 names the empirical-not-mathematical scope of the saturation regularity, the cohort-not-within-subject scope of the brain analyses, and the FACED-specific saturation threshold. Section 9 expands these into a numbered limitations subsection: (a) two EEG datasets (FACED deep, SEED-V replicates 5/5 claims; V-axis-as-supervision threshold is FACED-specific); (b) cohort-side analysis with explicit Simpson’s-paradox dynamic; (c) saturation threshold dataset-specific; (d) negative results require strong-baseline interpretation; (e) cross-architecture class-PC1 correlation softened by 93.5th-percentile null-direction control ($n = 1000$). Theta-gamma PAC null and arousal-axis brain-side failure are documented.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.

- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper is empirical. The “saturation regularity” (Section 3) is stated explicitly as an empirical pattern, not a mathematical theorem; we do not claim a proof. The boxed Statement of the Regularity in Section 3 writes out what would be the formal claim and the empirical pillars supporting it (16 / 25 statistically significant negatives, 5 monotonic-destruction families, unidirectional saturation transition, direct mechanism check). The Residual Contribution Regularity in Section 4 is an empirical regression claim (slope $\hat{\beta} = +0.74$, $p = 0.014$, $n = 10$), not a theorem.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The V-axis extraction protocol (story prompts, paraphrase generation, hidden-state extraction, PCA-PC1 with Joy-positive orientation) is documented in Section 5 and Appendix S2. EEG model training (architecture, optimisation, augmentation, KD, ensemble construction, exact seed list) is in Appendix S14. Statistical methods (bootstrap, random-direction null, paired tests) are in Appendix S15. Dataset splits (subjects 0–79 / 80–99 / 100–122 train/val/test on FACED) are in Appendix S1. The full intervention table with seeds and JSON pointers is in Appendix S10.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.

- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [\[Yes\]](#)

Justification: Code, configs, model checkpoints, and figure-generation scripts will be released at the camera-ready stage under an MIT licence (Appendix S19). The two datasets used (FACED, SEED-V) are publicly available from their original authors; we release our preprocessing and split scripts. The 10-checkpoint ensemble SOTA is reproducible from a single SLURM array job documented in the appendix. Anonymous-friendly access at submission time will be a placeholder URL.

Guidelines:

- The answer [\[N/A\]](#) means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [\[No\]](#) is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Optimiser (AdamW, $\text{lr} = 10^{-3}$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, weight decay 10^{-2}), schedule (cosine + 5-epoch warmup), batch size 128, training length $e \in \{100, 150\}$ epochs, depth $d \in \{3, 6\}$, filter dimension $f = 128$, seeds $\{42, 123, 456, 789, 2025\}$, augmentation parameters (Gaussian $\sigma = 0.05 \cdot \text{std}$, channel dropout $p = 0.15$, temporal masking $\leq 5\%$, all at $p = 0.6$), KD ($\lambda_{\text{KD}} = 0.5$, $T = 1.0$, rand9 9-D orthonormal teacher), and ensemble construction (10 ckpts split as $5 \times e = 100 + 5 \times e = 150$, uniform softmax averaging) are all specified in Appendix S14.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: We report standard deviation across 5 seeds in every recipe-cascade row (Section 8); Pearson r with p -values for every cross-architecture and brain correlation (Sections 6, 4); paired t -tests for V-axis interventions vs. matched-recipe baselines (Section 3); $B = 10,000$ subject-resampled bootstrap CIs and Fisher- z p -values for cohort EEG correlations (Appendix S15); 200-direction matched-norm random-direction null distributions for the cohort EEG-LLM correlation; 1000-direction null for the cross-architecture V-axis correlation; permutation p -values for the directional ablation (Section 4, $z = +10.87$, $p < 10^{-26}$ on the smoke checkpoint).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix S20 reports approximately 4,500 V100 GPU-hours total, broken down by experiment family: ~600 for the recipe ablation cascade, ~1,200 for the 25 V-axis-supervision interventions, ~800 for the 36-checkpoint cross-architecture analysis, ~1,200 for ensemble construction and val-test rank diagnostics, ~700 for cross-dataset (SEED-V) experiments. The 10-checkpoint ensemble SOTA result is reproducible in ~10 GPU-hours per seed. V-axis extraction is CPU-cheap (~10 minutes per LM on a single 80 GB GPU). The figure includes preliminary and failed experiments not in the paper.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We use only publicly released datasets (FACED [Chen et al., 2023], SEED-V [Liu et al., 2022b]) collected by other groups under their respective IRB-approved protocols. We collect no new human data. We use only publicly released LLMs under their published licences.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Improved EEG emotion classifiers have applications in mental-health monitoring and brain-computer interfaces (positive); they also pose risks of affective inference without consent in surveillance, workplace, or advertising contexts (negative). Section 9 flags this dual-use surface and notes that the V-axis extraction protocol applied to EEG could in principle reduce the calibration-data burden for affective inference systems — a feature that cuts both ways. We do not release human EEG data; only feature-side and training-side code.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: Our trained EEG-emotion checkpoints and V-axis extraction code do not introduce abuse vectors beyond what is already public via FACED, SEED-V, and the listed pre-trained LLMs. We release neither human EEG data nor LLM weights; only training configs and feature-side code. We assess the dual-use risk in the broader-impacts statement.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: FACED is cited [Chen et al., 2023] and used under its public release license; SEED-V is cited [Liu et al., 2022b]. All 14 LLMs (Qwen3 family of six dense sizes from 0.6B to 32B, Llama-4-Scout-17B, Mistral-7B, Phi-2-2.7B, Pythia-1.4B, BLOOM-560M, TinyLlama-1.1B, Gemma-27B, Gemma-4-31B) are cited and used under their respective Hugging Face / model-card licenses. The CBraMod [Wang et al., 2025] and EMOD [Chen et al., 2025] backbones are cited and used under their public licenses. CLIP [Radford et al., 2021] is cited.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The V-axis extraction scripts, V-axis-aligned EEG checkpoints, and figure-generation pipelines will be released with documentation (README, environment.yml, SLURM templates, per-experiment JSON results, plot-generation scripts) at the camera-ready stage. Anonymized URL placeholder will be in the supplementary materials at submission time.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: We do not collect new human data; all human EEG comes from the publicly released FAGED [Chen et al., 2023] and SEED-V [Liu et al., 2022b] datasets, originally collected by their respective groups under their own IRB-approved protocols.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: No new human-subjects research was conducted. The FACED [Chen et al., 2023] and SEED-V [Liu et al., 2022b] datasets were collected under their original IRB approvals, which we cite.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: LLMs are central to the V-axis extraction protocol. The 14 language models from which we extract the V-axis are named in Section 5 and Appendix S2: Qwen3 at all six dense sizes (lead model Qwen3-4B; also 0.6B/1.7B/8B/14B/32B), Llama-4-Scout-17B, Mistral-7B, Phi-2-2.7B, Pythia-1.4B, BLOOM-560M, TinyLlama-1.1B, Gemma-27B, Gemma-4-31B. The penultimate-layer hidden state ($L = L_{\max} - 1$, e.g., $L = 35$ of 36 for Qwen3-4B) is the source of all class centroids. We also use a generic instruction-tuned Qwen LM (temperature 0.7) as a paraphrase generator in pre-processing; the prompt is in Appendix S2. CLIP [Radford et al., 2021] (image encoder) is used for the cross-modal vision V-axis.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

Supplementary Material

S1 Datasets, Splits, and Preprocessing

FACED. 123 subjects, 28 emotional video clips covering 9 emotion categories (Anger, Disgust, Fear, Sadness, Neutral, Amusement, Inspiration, Joy, Tenderness) at 3–4 stimuli per class. EEG: 32 channels, 250 Hz sampling, 30-second clip duration. We use the official preprocessing protocol distributed with the dataset release [Chen et al., 2023]: bandpass 0.5–47 Hz, notch 50 Hz, ICA-based artefact rejection, common-average re-reference. Train / validation / test split: subjects 0–79 / 80–99 / 100–122 (no subject overlap), matching the protocol used by CBraMod [Wang et al., 2025] and EMOD [Chen et al., 2025] for fair comparison.

SEED-V. 16 subjects, 3 sessions, 5-class emotion (Disgust, Fear, Sad, Neutral, Happy), 62-channel EEG at 1000 Hz downsampled to 200 Hz. Standard SEED-V split with subject-disjoint train / test. We replicate the SEED-V CBraMod recipe at $d=6$ for the cross-architecture generality and five-claim re-derivation reported in Appendix S8.

DE features. For each (subject, stimulus, channel) we compute differential entropy in five canonical bands (δ 0.5–4, θ 4–8, α 8–13, β 13–30, γ 30–47 Hz) per second over the 30-second clip, then mean over time to obtain a 32×5 feature matrix per (subject, stim) pair. Cohort-level analyses average across the 123 subjects to give the $28 \times 32 \times 5$ tensor used in Section 6.

S2 V-Axis Extraction Protocol

Story prompts. Nine emotion-evocative stories (one per FACED class), each 1–3 sentences. Stories were authored once, reviewed by two annotators blind to the LLM evaluation results, and held fixed across all 14 LLMs.

Paraphrase generation. For each story, we generate $n = 50$ paraphrases via independent calls to a generic instruction-tuned Qwen LLM (temperature 0.7), with the prompt: “Rewrite the following so it preserves all emotional content but uses different words and sentence structures: {story}”. Paraphrases are filtered for length (5–80 tokens) and minimum cosine distance from the source (≥ 0.10 on Sentence-BERT embeddings).

Hidden-state extraction. Each paraphrase is tokenised and passed through the target LM. We take the last-token hidden state at layer L (default: $L_{\max} - 1$, the penultimate layer; for Qwen3-4B this is $L = 35$ of 36). Class centroids are $c_k = \frac{1}{50} \sum_{i=1}^{50} h_{k,i}^{(L)}$ for $k = 1, \dots, 9$.

V-axis as PCA-PC1. Stack the nine centroids into a $9 \times D$ matrix and run mean-centred PCA. The V-axis is the unit vector along PC1, oriented so that $\langle c_{\text{Joy}}, v \rangle > 0$. PC1 explains 41–58% of the across-class variance across the 14 models tested (median 49%).

Robustness. Story rephrasing was verified by sampling 5 alternative prompt sets per class and comparing per-stimulus V-axis projections; mean cross-prompt $r > 0.93$ across 14 models. Paraphrase count was verified at $n \in \{1, 5, 25, 50\}$; SST-2 zero-shot AUC 0.74 at $n = 1$, rising to the canonical 0.832 at $n = 50$ for Qwen3-4B.

S3 Per-LLM Cross-Architecture Convergence: Deep Dive

Sentiment benchmark table.

Three regimes against behavioural valence. The 14-LLM universality is not uniform: models cluster into three regimes by per-stim correlation against a behavioural valence reference ($n = 28$ FACED stim, mean of 123 subjects’ valence ratings).

- **Top tier** (per-stim $r > 0.85$, $n = 7$): all six Qwen3 sizes (Qwen3-14B $r = 0.944$, Qwen3-4B 0.944, Qwen3-8B 0.941, Qwen3-32B 0.937, Qwen3-0.6B 0.933, Qwen3-1.7B 0.930)

Benchmark	V-axis AUC	Supervised LR (5k)
SST-2 binary [Socher et al., 2013]	0.832	0.837
IMDB [Maas et al., 2011]	0.857	—
Yelp Polarity [Zhang et al., 2015]	0.938	—
Twitter SemEval-2017 [Rosenthal et al., 2017]	0.927	—
Rotten Tomatoes (MR) [Pang and Lee, 2005]	0.799	—
<i>SST-2 zero-shot reference baselines</i>		
SBERT prototype cosine [Reimers and Gurevych, 2019]	0.793	—
RoBERTa-large-MNLI entailment [Liu et al., 2019]	0.912	—

Table 3: Zero-shot V-axis sentiment results. SST-2 (0.832) is the verified Qwen3-4B lead-model number; the IMDB/Yelp/TweetEval/Rotten Tomatoes rows come from a concurrent multi-benchmark extraction with the V-axis recipe and may include earlier Qwen-family extractions. SST-2 alone is within 0.005 of a 5k-sample supervised LR baseline (0.837): the V-axis carries the structure of the supervised problem without seeing sentiment labels. For zero-shot context on SST-2, an all-MiniLM-L6-v2 prototype-cosine baseline (a generic sentence encoder, not specifically trained for sentiment) reaches 0.793, while RoBERTa-large-MNLI used as a zero-shot entailment classifier (a 355M-parameter model fine-tuned on a large supervised NLI corpus) reaches 0.912. Nine LLM stories therefore sit above an off-the-shelf sentence encoder and within 0.08 of a 355M-param NLI-supervised model on the same zero-shot evaluation.

901 plus Mistral-7B (0.935). This is the *modern-LLM manifold*; within Qwen3 it is exceptionally
902 tight (mean $r_{\text{behav}} = 0.938 \pm 0.005$, mean $r_{\text{eeg}} = 0.796 \pm 0.011$).
903 • **Middle tier** ($0.6 \leq r \leq 0.85$, $n = 4$): Llama-4-Scout-17B (0.826), Gemma-27B (0.807),
904 Gemma-4-31B (0.715), Phi-2-2.7B (0.651).
905 • **Out-of-manifold** ($r < 0.4$, $n = 3$): Pythia-1.4B (0.276), TinyLlama-1.1B (0.128),
906 BLOOM-560M (0.031).

907 The within-Qwen3 scaling is essentially flat ($r = +0.568$ between $\log(\text{hidden dim})$ and per-stim
908 r_{behav} , $p = 0.24$, $n = 6$): once a model has joined the modern-LLM manifold, convergence is *not*
909 a size-monotonic phenomenon. The Qwen3 family r_{eeg} range is $[+0.781, +0.812]$ with standard
910 deviation ± 0.011 , and the behaviourally best-aligned models are mid-size Qwen3 (4B–14B).

911 **V-shape across architectures.** The same V-shape emerges across five LM families (Qwen, Pythia,
912 BLOOM, Llama-3, Phi-3-mini), and across bidirectional encoders BERT, RoBERTa, and DeBERTa,
913 with the peak at layer $L_{\text{max}} - 1$ in each case. The cosine between input-layer V-axis and final-layer
914 V-axis is only 0.291, with mid-layer L14 at 0.067. The Gemma-4 family is the lone outlier: text
915 V-axis peaks at $L = 1$ rather than the penultimate layer; brain alignment survives at mid-layers
916 (Gemma-4-E2B at $L=17$ predicts EEG at $r = 0.819$; Gemma-4-E4B at $L=21$ at $r = 0.714$).

917 S4 Concept Library and Compositionality


918 **Compositional axis arithmetic.** Combining the V-axis of a stylistic concept (politeness) with
919 that of an emotional concept (happiness) predicts 4-quadrant categorical labels at 79% accuracy on
920 a held-out generation set, $3.16\times$ above chance, with uniform per-cell precision in the $[0.68, 0.86]$
921 range. Subtractive orthogonalisation cleanly separates the two: happy-AUC $0.912 \rightarrow 0.988$ on
922 the politeness-orthogonalised axis, while polite-AUC drops from $0.793 \rightarrow 0.560$. Downstream on
923 natural text the same composition predicts GoEmotions, SST-5, and Yelp quadrants at 41.8%–46.7%
924 (documented scope refinement).

925 **Recipe generality across 20 concepts.** We authored 9 stories per concept for 20 unrelated dimen-
926 sions (emotional, stylistic, abstract) and applied the identical extraction protocol. **20 of 20 concepts**
927 **reach binary AUC > 0.65 on held-out test stories, 17 of 20 reach ≥ 0.95 , 11 are perfect at 1.00**
928 **(Figure 5).**

929 **Cross-concept transfer matrix.** The 20×20 matrix has self-AUC mean 0.97 on the diagonal and
930 off-diagonal mean 0.83 (gap $+0.13$): the recipe extracts *concept-specific* directions rather than a

The 9-story V-axis recipe is concept-generic across 20 unrelated dimensions

Binary pole-vs-pole AUC, sorted; coloured by tier. Toxicity (Jigsaw) is the only substantive failure.



Source: reports/r6_concept_library_synthesis.md (Qwen3-4B, penultimate layer). ρ = Spearman monotonicity over 9 ordered classes.

Figure 5: Binary pole-vs-pole AUC across 20 unrelated concepts (sorted descending), plus the out-of-scope toxicity case (Jigsaw, AUC 0.59).

single global affect axis (Figure 6). Most general source axes: sarcasm (0.89), fear (0.88), envy and shame (0.87). Most specific: decisiveness (0.78), complexity (0.78). Most captured: urgency (0.97), fear (0.95), concreteness (0.94). Most concept-specific: specificity (0.65), sarcasm (0.67).

Toxicity: out-of-scope. Applied to the Jigsaw toxic-comment corpus, the toxicity axis attains AUC 0.59. Two non-mutually-exclusive hypotheses: (i) toxicity is more distributed in residual-stream geometry than a single late-layer PC1 captures (low-rank $k > 1$ subspace); (ii) Qwen-generated toxic stories, filtered through alignment training, may be too mild to span the real Jigsaw distribution. The recipe scope is sharply delineated: it covers valence and 19 other concepts, explicitly not toxicity.

S5 Multilingual and Vision Generality


Multilingual transfer. We translated the nine emotion stories into Japanese, Arabic, and Russian, and re-extracted the V-axis from four multilingual encoders. Direct application of an English-centric causal LM (Qwen3-4B) to non-English stories collapses to chance — a clean English-bias diagnostic. Switching to mT0-base [Muennighoff et al., 2023] (277M-param multilingual encoder-decoder) fully recovers the SST-2 result and *exceeds* the English baseline in every language tested (Figure 7).

Per-language V-axis cosines are moderate ($r_{en-es} = 0.47$, $r_{en-fr} = 0.16$, $r_{es-fr} = 0.46$): different languages share substantial direction but not entirely, consistent with concept-level rather than token-level encoding.

Cross-modal generality (CLIP vision). The same protocol applied to nine emotion-evocative *images* via CLIP [Radford et al., 2021] yields a V-axis that ranks the OASIS image dataset [Kurdi et al., 2017] on valence at Pearson $r = 0.869$ and on arousal at $r = 0.803$. Both *beat* supervised Ridge trained on the OASIS labels themselves (valence 0.836, arousal 0.789), making the zero-shot V-axis the strongest known image-affect predictor on this benchmark. The valence and arousal CLIP axes are essentially orthogonal ($|\cos \theta| = 0.013$).

S6 Specificity Controls (Nonce, Random, Arousal)

Nonce-word ablation. Substituting all content words in the nine emotion stories with phonologically valid English nonce words (preserving syntax and function words) drops SST-2 AUC from above 0.83 to chance (≈ 0.52) — the second decimal. The single most decisive control: had the V-axis



Source: reports/r6_concept_transfer_matrix.json (Owen3-4B, penultimate layer, 12 train + 8 test stories per concept). Cell (i, j) = source concept i 's V-axis classifying target concept j 's pole-vs-pole benchmark.

Figure 6: Cross-concept transfer matrix (best-orientation AUC). Diagonal: self-classification (white-bordered cells). Off-diagonal: source-axis-applied-to-target-benchmark AUC. Sorted top-to-bottom by mean off-diagonal. Mean diagonal 0.97, mean off-diagonal 0.83, gap +0.13.

958 been a syntactic or prompt-template artefact, the nonce variant would have retained template-level
959 structure and produced a non-trivial axis. The signal is encoded in semantic content.

960 **Random-direction baseline.** Random Gaussian directions in the LM's residual stream, L_2 -matched
961 to the V-axis, give SST-2 AUC ≈ 0.51 and lexicon $|r| \approx 0.10$ at every benchmark. Bootstrap 95% CIs
962 are tight: SST-2 AUC $[0.844, 0.890]$, EmoBank V $|r| \in [0.485, 0.512]$, Hu & Liu $|r| \in [0.740, 0.776]$
963 (1,000-sample bootstrap; statistical methods in Appendix S15).

Benchmark	V-axis (valence)	A-axis (arousal)
SST-2 binary AUC	0.832	—
EmoBank V $ r $ / EmoBank A $ r $	0.49	0.01–0.06
Warriner V $ r $ / Warriner A $ r $	0.703	≤ 0.18
NRC valence / arousal lexicon recovery	$\geq 76\%$	$\leq 13\%$
OASIS image valence (r) / arousal (r)	0.869	0.803
FACED EEG cohort r	0.870 (9-stim)	$r \in [0.18, 0.41]$

Table 4: Valence transfers across modalities; arousal does not in text but does in vision. The same recipe applied to an arousal axis yields the strongest known zero-shot OASIS arousal predictor ($r = 0.803$, beating supervised ridge 0.789) but collapses on all text benchmarks. Cross-language A-axes are essentially orthogonal.




Figure 7: Cross-lingual V-axis recovery. Direct causal-Qwen transfer collapses (English-bias); mT0-base recovers and exceeds the English ceiling at 5× smaller parameter count across three typologically distant languages.

Arousal asymmetry: where the recipe fails in text. The asymmetry — *valence transfers, arousal does not in text* — rules out the trivial explanation that “any sufficiently rich PCA recipe gives any sufficiently graded affective axis.” Arousal in text appears to be encoded along a more distributed or non-linear subspace that the late-layer principal-component recipe does not access; the same recipe *does* access it in vision.

S7 Brain-Side Deep Dive

Per-stimulus reliability. Per-stimulus split-half reliability (odd/even subjects, $n \approx 62$ each) is $r = 0.988$, indicating a clean stimulus-level signal. Trial-level correlations ($n = 3,444$ subject-stim pairs) range from $r = 0.17$ – 0.21 ($p < 10^{-23}$); class-level ranking ($n = 9$ centroids) hits $r = +0.886$ vs. behavioural valence. The cohort $r = 0.87$ headline is reliability-bounded.

14-LLM brain forest. Per-LLM brain-prediction quality across all 14 models is shown in Figure 8: the top-tier 13 LLMs all clear their per-LLM random-direction nulls, while the out-tier sits at the median.

Time-locked dynamics. Mid-late α/β peak at $t = 18$ – 21 s coincides with the LPP / sustained-attention window of Hajcak & Foti [Hajcak et al., 2010] and Codispoti et al. [Codispoti et al., 2023]; early γ peak at $t = 3$ – 6 s coincides with Müller’s affective-picture window [Müller et al., 1999]. Per-second cohort $|r|$ in the alpha band arrives at $t \approx 21$ s, in beta at $t \approx 18$ s, in gamma at $t \approx 3$ s; best individual channel $|r| = 0.68$ at $t = 21$ s.

Per-subject peak heterogeneity. Standard deviation of per-subject alpha peak times is $\sigma \approx 9$ s; only 17–24% of the 123 subjects have their individual alpha or beta peak inside the cohort 18–21 s window. The cohort signal is therefore an envelope of the distribution rather than a tight common attractor (NF4).

Gemma-4 boundary check. Brain alignment survives the text-domain $L = 1$ peak anomaly: Gemma-4-E2B at $L = 17$ predicts EEG at $r = 0.819$ ($p \approx 10^{-7}$), Gemma-4-E4B at $L = 21$ at $r = 0.714$ ($p \approx 2 \times 10^{-5}$). The brain captures Gemma-4’s V-axis at the layer where the model first stably encodes it, not at a fixed late-layer offset. Cross-architecture convergence refines to a property of the late-but-not-necessarily-penultimate residual stream geometry.

14 language models predict cohort EEG; the rank tracks text-domain V-axis quality




Figure 8: Per-LLM brain-prediction quality. Top-tier 7 LLMs all exceed the 95th percentile of their per-LLM 200-direction null ($p < 0.05$ each); out-tier 3 LLMs sit at the median or below. Restricted to top-tier the log-hidden-dim $\rightarrow r$ scaling is $r = +0.125$, $p = 0.684$; within-Qwen3 $r = +0.568$, $p = 0.24$, $n = 6$. Family-membership $\gg$ scale.

S8 SEED-V Replication Full Numerics

We re-test all five FACED claims with a SEED-V-derived V-axis. To remove FACED leakage, we re-build the V-axis from scratch using the same Section 5 protocol applied to the 5 SEED-V emotion classes (Qwen3-4B at $L=35$, 50 stories per class, PCA on the 5 centred centroids).

Step 1: SEED-V cohort EEG-LLM alignment. We compute DE features per (channel, band, second) across all 16 subjects and aggregate to a $45 \times 62 \times 5$ cohort tensor. For each (channel, band) cell we compute Pearson with the SEED-V V-axis projection.

- **Best cohort cell:** P1 / θ , $r = +0.6159$.
- **Region-mean $|r|$:** occipital 0.329, parietal 0.347, central 0.324, frontal 0.253.
- **Posterior dominance:** occipital – frontal = $+0.076$.
- **Random-direction null** (200 directions matched in L_2): empirical $|r|_{\text{max}} = 0.6159$ vs. null 95th percentile = 0.478, $p < 0.005$.

Step 2: SEED-V brain topography. The Davidson FAA probe (positive vs. negative valence, right – left) on SEED-V replicates the FACED FAA topography:

Electrode pair (right – left, alpha)	(pos – neg) DE Δ
F4 – F3	–0.0018
F8 – F7	+0.0050
FP2 – FP1	–0.0094
AF4 – AF3	–0.0039

Table 5: Davidson FAA on SEED-V. Sign and magnitude consistent with the FACED replication.

Step 3: SEED-V cross-architecture convergence. Penultimate (200-D) features from 15 stock CBraMod SEED-V-5s checkpoints (3 protocols $\times$ 5 seeds: vanilla baseline, augmentation, KD-midlayer) aggregated per-class on the test set. Class-level PC1/best-PC scores correlated with SEED-V V-axis class projection give $r(\text{BACC, class-best-PC } r) = +0.601$ ($p = 0.018$, $n = 15$) and $r(\text{BACC, trial-best-PC } r) = +0.632$ ($p = 0.012$). The unsigned class-PC1 correlation is $+0.358$

($p = 0.19$); the SEED-V BACC dynamic range is much narrower than FACED, so the strongest convergence shows up at the best-PC level.

Step 4: SEED-V V-axis as supervision. Adding a V-axis topo-MSE auxiliary loss ($\lambda_{\text{vax}}=0.05$ on {PO3, PO4, POz, Oz, O1, O2, P3, P4}) on top of the same CBraMod-5s recipe, run with 5 baseline + 5 V-axis seeds:

Variant	BACC (mean $\pm$ s.e. across 5 seeds)	vs. baseline Δ
Baseline (no V-axis loss)	0.4300 ± 0.0079	—
+ V-axis topo loss	0.4325 ± 0.0040	+0.0025

Table 6: SEED-V V-axis-as-supervision (within seed noise of zero). Two factors likely explain: (1) the SEED-V class-level V-axis is sharply 5-valued (cohort class-level $r = +0.989$), so adding a class-level V-axis target on top of one-hot supervision is largely redundant; (2) the SEED-V BACC dynamic range (0.43–0.45) is much narrower than FACED (0.55–0.58), so a small effect is hard to detect with five seeds.

Step 5: SEED-V directional ablation (Arditi-style). At inference time on each of the 15 SEED-V CBraMod checkpoints (5 baseline + 5 augmented + 5 KD-midlayer) we project penultimate features onto the orthogonal complement of the trial-level best-PC V-axis direction (no retraining; the class-level direction is degenerate at $k=5$, so we use the trial-level analogue, mirroring the trial-best-PC $r = +0.632$ estimator from Step 3). Every checkpoint shows a negative $\Delta\text{BACC} \in [-0.0291, -0.0093]$ (population mean -0.0173 ± 0.0061); a matched random-direction control ($n = 20$ uniform on $\mathbb{S}^{D-1}$) is centred at 0 ($\bar{\Delta}_{\text{rand}} = +0.00013 \pm 0.0008$). Per-checkpoint z -scores range $[+9.65, +44.92]$ (mean $+23.6$, all $p < 0.001$ parametric one-sided); empirically every V-direction ablation exceeds every one of the 20 random ablations on every checkpoint (15/15 at empirical $p < 0.05$). The SEED-V directional effect is proportionally larger than FACED (mean $\Delta/\text{baseline} \approx 3.9\%$ vs. 2.4% on FACED) and the mean z is $3\times$ FACED’s. As on FACED, V-axis amplification (rather than ablation) is not directionally consistent (mean $\Delta \approx 0$, mean $z \approx -2$) — expected, since the SEED-V V-axis already aligns with class-discriminative variance and doubling it does not help. Per-checkpoint table and the diagnostic information for both class-level (degenerate) and trial-level (used) residual directions are stored in `seedv_causal_ablation_full.json`.

S9 Brain Topography Deep Dive

Davidson FAA full pairs. Figure 9 compares all three Davidson asymmetry pairs to the posterior occipital region-mean reference.

Davidson frontal-alpha asymmetry replicates qualitatively, but at $\sim 10\times$ smaller magnitude than posterior V-axis encoding




Figure 9: Frontal-alpha asymmetry replicates in direction with smaller magnitude than posterior V-axis encoding for these video stimuli. (a) Anatomical scalp diagram with the three Davidson asymmetry pairs and their r values. (b) Per-pair alpha and delta asymmetry bars. (c) Direct comparison to the occipital region-mean $|r| = 0.21$.

Stimulus subset (out of 28 FACED stimuli)	Cohort r at PO3/ γ
All 28 stimuli (full)	+0.478
Anger only ($n=3$)	not estimable
Anger + Amusement + Tenderness ($n=9$)	+0.870
Excluding Anger ($n=25$)	+0.327
Excluding Anger + Amusement + Tenderness ($n=19$)	-0.015

Table 7: The cohort V-axis effect is essentially an *anger-versus-warm-positive* contrast across nine stimuli.

9-stim contrast and Simpson's paradox.

- Cohort r at fixed top-8 channels: +0.021
- Mean per-subject r at fixed top-8 channels: -0.062
- Per-subject best-channel oracle (1 cell): $|r| = 0.551$
- Per-subject best (channel, band) oracle: $|r| = 0.616$

The all-band oracle ($|r| = 0.616$) exceeds the cohort-fixed top-8 by +0.39 in absolute $|r|$, indicating considerable per-subject signal that the cohort summary suppresses. The Simpson's-paradox dynamic that Sani et al. [Vahidi et al., 2024] flag for cross-subject EEG-emotion modelling more broadly (Figure 10).




Figure 10: The cohort V-axis is a 9-stimulus contrast and a between-subject phenomenon. Top: drop-class subset bars, per-emotion drop- Δ , PO3/ γ stim-level scatter colour-coded by emotion. Bottom: per-subject r distribution at fixed top-8 channels (centred near 0, cohort line at +0.02); per-subject best-channel oracle distribution (centred at $|r| \approx 0.55$); summary callout.

Time-resolved peaks. Figure 11 traces per-second cohort $|r|$ across all five frequency bands.

- **Alpha** (8–13 Hz): peak at $t = 21$ s, cohort $r = 0.40$; best individual stimulus $r = 0.68$.
- **Beta** (13–30 Hz): peak at $t = 18$ s, cohort $r = 0.41$; best stim $r = 0.62$.
- **Gamma** (30–47 Hz): peak at $t = 3$ –6 s, cohort $r = 0.21$; best stim $r = 0.61$.
- **Delta/Theta** (0.5–8 Hz): peak at $t = 12$ s, cohort $r = 0.18$.




Figure 11: Time-resolved V-axis encoding across the 30-second clip. **(a)** Per-second cohort $|r|$ across all 28 stimuli, one trace per frequency band; the early γ peak ($t \approx 2$ s) coincides with Müller’s affective-picture window and the sustained α/β peak ($t \approx 17$ – 20 s) coincides with the LPP sustained-attention window. **(b)** Same trace restricted to the 9-stimulus emotional-pole subset where the cohort effect lives; α reaches $|r| \approx 0.78$. **(c)** Per-subject distribution of α -peak times: only $\sim 24\%$ of subjects peak inside the cohort 18–21 s window, evidence that the cohort signal is a between-subject phenomenon rather than a tight common attractor.

1047 **Functional connectivity.** The eight V-axis-aligned channels (PO3, F7, O1, P3, Oz, O2, P4, PO4)
 1048 do not act independently: they form a tight functional network in the gamma band. Pairwise γ -DE
 1049 correlation has mean off-diagonal $r = 0.675$ across the 8 channels — substantially above the cohort-
 1050 average of $r \approx 0.32$ in this band. Structure is dominantly a posterior-occipital-parietal cluster (PO3,
 1051 PO4, P3, P4, O1, O2, Oz; pairwise mean $r \approx 0.71$) with F7 as a single *frontal hub* (mean correlation
 1052 to the seven posterior nodes $r \approx 0.47$; Figure 12).




Figure 12: V-axis-aligned channels form a coordinated network with F7 as frontal hub. Left: scalp graph with 8 V-axis channels as red nodes and edges weighted by pairwise γ -DE correlation. Right: hierarchically clustered 8×8 correlation matrix; F7 auto-isolates as the frontal hub.

1053 **Mutual information vs. linear correlation.** We computed mutual information between the V-axis
 1054 projection and cohort EEG features (nearest-neighbour MI estimator on quantile-normalised inputs),
 1055 alongside the same 200 random-direction controls used for the linear test. Observed MI 0.112; null
 1056 mean 0.036; $p_{\text{MI}} = 0.115$ (n.s.) vs. linear $p_r = 0.020$. The relationship is essentially linear: a
 1057 non-linear estimator does not detect signal above what linear regression captures. We use ridge
 1058 regressions throughout.

1059 **Theta-gamma coupling does not mediate the V-axis.** Tort modulation index per channel at
1060 the theta-phase / gamma-amplitude pairing, correlated with V-axis encoding strength $|r|$ across 32
1061 channels: $\rho(\text{PAC}, Vr) = +0.082$, $p = 0.667$. Channels with strong theta-gamma CFC are not
1062 the channels that encode the V-axis, and vice versa — the V-axis encoding is dissociable from the
1063 phase-amplitude-coupling channel [Canolty and Knight, 2010].

1064 S10 Saturation: Full Intervention Table and Forest

#	Family / variant	Δ	p	Verdict
1	Frontal-mask $\lambda=0.5$	-0.052	0.0015	sig. negative
2	Frontal-mask $\lambda=0.1$	-0.009	0.21	n.s.
3	FAA $\lambda=0.5$	-0.044	0.006	sig. negative
4	FAA $\lambda=0.1$	-0.018	0.063	borderline
5	Anger-weighted $\lambda=0.5$	-0.054	0.0003	sig. negative
6	Occipital $\lambda=0.1$	-0.022	0.007	sig. negative
7	Topo-optimal $\lambda=0.1$	-0.013	0.039	sig. negative
8	Topo-optimal $\lambda=0.05$	+0.0021	0.50	n.s. (NULL)
9	Procrustes $\lambda=0.05$	+0.0012	0.89	n.s. (NULL)
10	RSA $\lambda=5.0$	-0.093	$< 10^{-4}$	sig. negative
11	RSA $\lambda=1.0$	-0.057	$< 10^{-3}$	sig. negative
12	Distance-CE $\tau=5.0$	-0.397	$< 10^{-4}$	catastrophic
13	Multi-V $\lambda=0.5$	-0.043	$< 10^{-3}$	sig. negative
14	EEG-AUX MSE $\lambda=0.5$	-0.043	$< 10^{-3}$	sig. negative
15	EEG-AUX MSE $\lambda=0.1$	-0.016	0.04	sig. negative
16	EMODSTYLE class $\lambda=1.0$	-0.010	0.18	n.s.
17	EMODSTYLE stim $\lambda=0.5$	+0.0066	0.22	n.s. (sweet spot)
18	PEFT fullhead $\lambda=0.1$	-0.009	0.069	borderline
19	PEFT LoRA $\lambda=0.1$	-0.010	0.12	n.s.
20	PEFT IA3 $\lambda=0.1$	-0.016	0.05	sig. negative
21	Pretrain-FT (frozen)	-0.401	$< 10^{-4}$	catastrophic
22	Pretrain-FT (unfrozen)	-0.190	$< 10^{-4}$	sig. negative
23	XEEG (FACED $\rightarrow$ SEED-V)	+0.0015	0.97	n.s. (zero)
24	SCALING (full data) $\lambda=0.5$	-0.045	< 0.01	sig. negative
25	Uncertainty multi-task	-0.067	$< 10^{-3}$	sig. negative
SOTA recipe + Topo $\lambda=0.1$		-0.015	0.020	sig. negative (cliff)
SOTA recipe + EMODSTYLE		-0.024	0.001	sig. negative (cliff)

Table 8: Full V-axis-as-supervision intervention results. 16 / 25 families produce a statistically significant negative (14 sig. negative + 2 catastrophic); the remaining 9 are not significantly different from zero. None reach $p < 0.05$ in the positive direction at the strong baseline. Below the line are the SOTA-recipe replications confirming the saturation cliff. Forest plot in Figure 13.

1065 S11 Saturation Extras: Anger Paradox, Mechanism Check, Path B

Recipe	n seeds	Vanilla / + V-axis	Δ
CBraMod baseline + Topo $\lambda=0.05$	5	0.572/0.567	-0.005 (n.s.)
CBraMod baseline + EMODSTYLE $\lambda=0.5$	4	0.572/0.577	+0.005 (n.s.)
EMOD $d=3$ vanilla + Topo $\lambda=0.05$	5	0.6194/0.6260	+0.007 (n.s.)
EMOD $d=6$ + LS + KD + aug + Topo (full)	partial	0.6581/0.6432	-0.015 ($p < 0.05$)
EMOD $d=6$ + LS + KD + aug + EMODSTYLE	partial	0.6581/0.6346	-0.024 ($p < 0.01$)

Table 9: V-axis supervision is within seed noise at weak baselines and statistically significantly *harms* the strong SOTA recipe. The transition is unidirectional: V-axis goes from neutral to actively harmful as the base recipe strengthens, with no intermediate “helps” regime in our 5-seed runs.

1066 Saturation transition table.




Figure 13: Forest plot of all V-axis-as-supervision interventions (Table 8) ranked by Δ vs. baseline, with significance tier coloured. None of the 25 reach $p < 0.05$ in the positive direction at the strong baseline.

1067 Saturation cliff figure (LF6).

1068 **Why anger-weighting hurts despite the analytical ceiling.** Section 7 showed that the cohort
1069 EEG signal is carried by a 9-stimulus emotional-pole contrast (Anger + Amusement + Tenderness,
1070 $r = 0.870$ at $\text{PO3}/\gamma$). Weighting the V-axis loss to emphasise these three classes pushes the analytical
1071 ceiling on the V-axis projection from $r = 0.478$ (uniform) to $r = 0.714$ (anger-weighted). By every
1072 classical analysis, this should be the strongest possible V-axis supervision. It is the worst single
1073 intervention we tested: $\Delta = -0.054$, $p = 0.0003$. This is the cleanest argument for saturation: the
1074 model has already absorbed exactly the structure the loss is trying to inject; optimising the loss must
1075 therefore perturb the absorbed structure in a counterproductive direction.

1076 **Direct mechanism check.** For each of 8 V-axis-supervised variants (Topo at $\lambda=0.05$, Procrustes at
1077 $\lambda=0.05$, EMODSTYLE-stim at $\lambda=0.5$, full SOTA recipe + Topo/EMODSTYLE at $e=100/e=150$),
1078 we measure the change in two quantities relative to the matched-recipe baseline: Δr_{PC1} (class-mean
1079 V-axis encoding) and Δr_{resid} (within-class V-axis residual encoding). V-axis training raises class-PC1
1080 $|r|$ by $+0.01$ to $+0.36$ across the 8 variants, but moves the within-class residual encoding by only
1081 $\sim 10^{-7}$ in absolute value — numerical zero. The accuracy decrement at strong recipes therefore
1082 comes from *noise injected into the class-PC1 basin*, not from any beneficial transfer of residual
1083 structure.

1084 **Path B: ensembling-in V-axis checkpoints.** We extend the 10-vanilla SOTA ensemble with 0–15
1085 V-axis-trained Topo checkpoints and 0–10 V-axis-trained EMODSTYLE checkpoints, comparing
1086 5-seed ensembles in each configuration. Adding 15 Topo V-axis ckpts to the 10-vanilla pool gives
1087 ensemble BACC $\Delta = -0.0145$ ($p < 10^{-3}$); adding all 25 V-axis ckpts gives $\Delta = -0.0193$
1088 ($p < 10^{-3}$). At every point on this curve, V-axis-trained ensembles are strictly worse than the




Figure 14: Δ BACC from V-axis supervision plotted against base-recipe BACC across ~ 25 interventions. The sign of the effect transitions in the interval $[0.62, 0.66]$.

1089 matched vanilla ensemble. This rules out the diversity-gain defence: V-axis checkpoints disagree, but
 1090 their disagreements are aligned with the wrong subspace.

1091 S12 Convergence and Ensemble Theory: Extras

1092 **Per-architecture breakdown of the 36-checkpoint correlation.** CBraMod checkpoints (lower
 1093 BACC, ~ 0.57) cluster at low V-axis encoding strength (mean class-PC1 $|r| = 0.21$). EMOD-vanilla
 1094 d6 (~ 0.65 BACC) sits in the middle (mean 0.67); EMOD d6 e150 (~ 0.66 BACC) clusters at the
 1095 top (mean 0.69). The trend spans the full BACC range $[0.57, 0.69]$ in a single approximately linear
 1096 band, with the inter-architecture gap dominating the within-architecture spread.

1097 **Random-direction null detail.** We sample 1000 random Gaussian directions $w \in \mathbb{R}^{28}$
 1098 each matched in L_2 norm to v_{CLIP} , and recompute the 36-checkpoint correlation $\rho_w =$
 1099 $\text{Pearson}(\text{BACC}, |\text{corr}(\text{PC}_1(H(m)), w)|)$. The distribution of ρ_w has mean -0.03 and standard
 1100 deviation 0.62 (range $[-0.95, +0.95]$); the wide null is a consequence of the low intrinsic dimension-
 1101 ality of the class-PC1 subspace ($D_{\text{eff}} = 9$). The empirical V-axis correlation $+0.885$ sits at the 93.5th
 1102 percentile of this null ($p_{\text{one}} = 0.066$). We treat the within-class residual ($D_{\text{resid}} \approx D_m - 9$) as the
 1103 statistically robust signal.

1104 **Class-PC1 basin saturation.** The 10 SOTA-pool checkpoints exhibit a saturated class-PC1 V-axis
 1105 range: $|r| \in [0.60, 0.77]$ across all 10 seeds (mean 0.69, std 0.05). Single-seed variation in class-
 1106 mean V-axis structure is small. The ensemble does not gain from variance in this subspace; gains
 1107 come exclusively from the orthogonal residual.

1108 **Directional ablation figure (LF16).**

1109 S13 Ensemble Generality, Mega-Pool, and Best-Single

1110 **Recipe cascade and two-tier scatter.** Figure 16 traces the seven recipe steps from CBraMod’s
 1111 0.572 to our 0.6948 ensemble. Figure 17 shows the per-checkpoint within-class residual–ensemble-
 1112 contribution correlation that motivates the ensembling strategy.

1113 **Generality of the ensemble mechanism.** Figure 18 extends the two-tier picture across four
 1114 benchmarks (FACED, SEED-V, CIFAR-10, MNIST).

1115 As individual accuracy approaches the Bayes optimum, the within-class residual variance — the
 1116 source of ensemble gain — shrinks toward zero. The two-tier picture predicts this: ensemble gain is

V-axis ablation hurts every checkpoint; encoding strength does not predict drop magnitude (n=10)
Pearson $r=+0.168$ ($p=0.643$, $n=10$) slope= $+4.69\%$ / unit $|r|$




Figure 15: Directional ablation of v_{resid} on all 10 SOTA-pool checkpoints. *Top*: per-checkpoint BACC drops (red) compared to matched-norm random-direction ablation (grey); the V-axis effect is well above the null on every checkpoint (mean $z \approx 7.7$, all $p < 0.001$). *Bottom*: per-checkpoint ablation ΔBACC vs. within-class residual encoding $|r|$; the population-mean drop ($\bar{\Delta} = -0.0157$) confirms the residual-encoding mechanism across the pool.




Figure 16: Recipe ablation cascade from the EMOD baseline at 0.6287 (previous FACED-9 SOTA [Chen et al., 2025]) to our 10-checkpoint ensemble at 0.6948 (+10.5% relative). Each bar is one recipe component or ensembling step. Older baselines (CBraMod 0.572, EmotionKD 0.628) shown for reference.

1117 a property of the within-class V-axis residual variance reduction subspace, which has measure zero
1118 at the dataset ceiling. SEED-V at 0.37 has the largest residual-variance ceiling; we observe $+0.034$
1119 gain. MNIST at 0.985 has essentially none; we observe $+0.001$ gain.

1120 **Mega-ensemble null.** Two mechanisms explain the null. First, deeper variants ($d=8, d=10$) stay in
1121 the same class-PC1 V-axis basin (mean class-PC1 $|r|$ across $d \in \{4, 6, 8, 10\}$: 0.43, 0.43, 0.39, 0.38)
1122 and their within-class residuals overlap more than $d=6$ seeds do. Second, LLM-KD variants encode
1123 the residual along a different axis from rand9-9D KD vanilla checkpoints, but in a way that does
1124 not transfer to the test distribution. Cross-architecture mixing ($d \in \{4, 6, 8, 10\}$) gives 0.6791,




Figure 17: Two-tier ensemble theory. Per-checkpoint within-class V-axis residual $|r|$ predicts leave-one-out ensemble contribution ($r = +0.74$, $p = 0.014$, $n = 10$). Top-7 vs. bottom-7 split highlighted.




Figure 18: Two-tier ensemble gain across four benchmarks. Δ BACC is largest where single-model accuracy is furthest from the dataset ceiling, and shrinks to near-zero on MNIST where the saturated regime leaves no within-class residual variance to cancel. (a) Scatter of Δ vs. headroom $1 - \text{Acc}$, with linear fit slope ~ 0.05 . (b) Paired single-seed-mean vs. 5-seed-ensemble bars per benchmark, sorted by headroom.

essentially tied with within-arch x-seed at 0.6798. The diversity that helps is intra-architecture, training-length-driven.

Cohen κ disagreement structure. Single-seed accuracy at $e=100$ (0.6581) and $e=150$ (0.6581) is identical, but Cohen κ within- $e=100$ is 0.694, within- $e=150$ 0.679, cross-group 0.702. The $e=150$ checkpoints disagree more amongst themselves and more across groups, which is why mixing $e=100$ and $e=150$ gives +0.014 over either group alone. Cross-group disagreement on test trials is 13.1%; the mixed ensemble gains +10.3 percentage points exactly on the disagreement subset.

Best single checkpoint without ensembling. For deployment scenarios where ensembling is impractical, our best single-checkpoint result is BACC = 0.6755 on the $d=6$ $e=150$ recipe, seed 789, val-selected. This is the top of a 25-checkpoint val-test rank correlation analysis (Spearman 0.825, $n = 25$) and exceeds CBraMod by +0.103 and EMOD by +0.047 on a strict apples-to-apples test split.

S14 EEG Model Training Details

Architecture. EMOD axial transformer [Chen et al., 2025]: input 32×5 DE features, axial self-attention over channels and bands at depth $d \in \{3, 6\}$, filter dimension $f = 128$. The full SOTA recipe is $d = 6$, $f = 128$.

Ensemble pool	n	BACC	Δ vs. 10-ckpt
$10 \times d=6$ ($e=100+e=150$, our SOTA)	10	0.6948	—
$+5 \times d=8$	15	0.6909	−0.0039
$+5 \times d=8 + 5 \times d=10$	20	0.6883	−0.0065
$+10 \times d=\{8, 10\} + 5 \times \text{LLM-KD}$	25	0.6831	−0.0117

Table 10: Mega-ensemble null. Adding architectural variants beyond $d=6$ produces monotonically lower ensemble BACC. The $d=6 + e=100/e=150$ pool is the optimum within the architectures evaluated.

Optimisation. AdamW with $\text{lr} = 10^{-3}$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, weight decay 10^{-2} . Cosine schedule with 5-epoch linear warmup. Batch size 128. Training length $e \in \{100, 150\}$ epochs. We use the validation BACC for checkpoint selection and report test BACC.

Augmentation. Per-trial Gaussian noise ($\sigma = 0.05 \cdot \text{std of feature distribution}$), random channel dropout ($p = 0.15$), and random temporal masking (up to 5% of seconds). All augmentations are applied with $p = 0.6$ during training only.

Knowledge distillation. Soft-target KD with rand9 9-D orthonormal teacher (no LLM content); $\lambda_{\text{KD}} = 0.5$, $T = 1.0$. The LLM-9-D and rand9-9-D variants give within-noise identical BACC, validating that KD provides architectural regularisation without semantic content.

Ensemble. 10 checkpoints split as 5 seeds at $e = 100 + 5$ seeds at $e = 150$. Seeds: $\{42, 123, 456, 789, 2025\}$ each. Ensemble prediction is uniform softmax averaging across the 10 checkpoints’ output probabilities, followed by argmax.

S15 Statistical Methods

Bootstrap CIs. Per-subject and per-stim 95% CIs are computed from $B = 10,000$ subject-resampled bootstraps. For cohort correlations, we report the bootstrap median, 2.5% / 97.5% percentile interval, and a Fisher- z p -value for $r = 0$.

Random-direction null. For cohort EEG–LLM correlations, we sample $N = 200$ random Gaussian directions in $\mathbb{R}^{28}$ matched in L_2 norm to the V-axis projection, recompute the cohort $|r|$ for each, and report the percentile of the empirical $|r|$ in the null distribution.

Cross-architecture null. For the 36-checkpoint cross-arch correlation, we sample $N = 1000$ random directions matched in norm to v_{CLIP} and recompute the BACC–V-axis correlation. The null is distributed broadly ($\sigma = 0.62$) due to the low intrinsic dimension of the class-PC1 subspace; we report the empirical correlation’s percentile rank.

Per-checkpoint LOO contribution. For the 10-vanilla $d=6$ ensemble, we compute $\Delta_k = \text{BACC}(\text{ens}) - \text{BACC}(\text{ens} \setminus k)$ for each k . Reported correlations between Δ_k and per-ckpt features are Pearson with $n = 10$.

V-axis intervention comparisons. For each intervention, Δ is computed seed-paired with the matched-recipe vanilla baseline; p -values are paired t -tests across 5 seeds.

S16 Discussion Extras and Future Work

Implications for affective neuroscience. For video-evoked emotion on FACED, the V-axis is encoded predominantly in posterior visual cortex ($|r|_{\text{occipital}} = 0.21$ vs. $|r|_{\text{frontal}} = 0.16$), and Davidson’s frontal-alpha asymmetry replicates in direction at smaller magnitude. We do not refute Davidson’s hypothesis — the frontal-alpha asymmetry is real, with F8-F7 alpha $r = +0.0155$ and Fp2-Fp1 alpha $r = +0.0116$, both in the hypothesised direction — but we add a posterior empirical account that is dominant for video paradigms. The 9-stimulus emotional-pole structure sharpens

the V-axis claim from “smooth gradient over all emotions” to “anger-versus-warm-positive contrast across nine clips”.

Per-subject V-axis adaptation does not transfer. The per-subject best-channel oracle reaches $|r| = 0.62$ (a +0.39 absolute headroom over the cohort-fixed top-8 of $|r| = 0.02$). No clean estimator we tried (V1 global top-K, V2 per-validation-subject majority vote, V3 TTA label-free profile) recovers more than +0.001 over the cohort top-K. The per-subject signal is real but not transferable through simple selection rules. A subject-conditioned channel-attention head trained on a small calibration set per subject is the natural follow-up.

Future work. Three lines: (i) per-subject V-axis adaptation closing the $|r| = 0.62$ oracle gap via subject-conditioned channel-attention; (ii) replicating the saturation regularity for other concept directions [Arditi et al., 2024] (arousal, dominance, formality, refusal) on their respective benchmarks; (iii) using the within-class residual $|r|$ as an online model-selection signal during training rather than as a post-hoc diagnostic.

S17 FACED Test Confusion Matrices

Figure 19 contrasts the best-single-checkpoint and 10-checkpoint-ensemble confusion matrices on the FACED 9-class test split.




Figure 19: FACED 9-class test confusion. Left: best single checkpoint (BACC 0.6755). Right: 10-checkpoint ensemble (BACC 0.6948). The ensemble’s gain is concentrated on the off-diagonal cells where single-seed disagreement is highest, consistent with the within-class residual mechanism of Section 8.

S18 Negative Results Catalogue

In the spirit of full disclosure, the following null or negative results bear mention beyond what was incorporated into the main paper:

- Theta-gamma cross-frequency coupling does not mediate the V-axis** (Section 7, Appendix S9): Tort modulation index $\rho(\text{PAC}, V_r) = +0.082$, $p = 0.667$.
- Mutual-information control is not significant** (Section 6): MI $p = 0.115$ vs. linear $p = 0.020$. The V-axis-EEG relationship is essentially linear.
- Random-direction null on cross-arch correlation** (Section 4): $r = +0.885$ at the 93.5th percentile of a 1000-direction null, $p_{\text{one}} = 0.066$. Within-class residual ($r = +0.74$, $p = 0.014$) is the statistically robust signal.
- Per-subject V-axis adaptation does not transfer.** The per-subject best-channel oracle reaches $|r| = 0.62$, but no clean estimator transfers more than +0.001 over the cohort top-K.

- 1204 5. **Cross-architecture ensembling does not add diversity.** Mixing $d \in \{4, 6, 8, 10\}$ check-
1205 points gives within-arch 0.6798 vs. cross-arch 0.6791 (n.s.).
- 1206 6. **Mega-ensembles plateau at 10 checkpoints.** 15-, 20-, 25-checkpoint pools all hit BACC
1207 ≤ 0.6948 (Appendix S13).
- 1208 7. **Stimulus-aggregation re-ranking does not help.** Operating at the 28-stimulus level instead
1209 of trial level gives BACC=1.0 by construction (label leak); the corresponding trial-level head
1210 does not transfer.
- 1211 8. **Test-time augmentation at K=5 does not improve the ensemble** (0.6753 $\rightarrow$ 0.6720).
1212 Averaging predictions over augmented test trials does not preserve the V-axis residual
1213 structure.
- 1214 9. **Toxicity: recipe scope.** The 9-prompt PCA recipe attains AUC 0.59 on Jigsaw toxic
1215 comments, below the 17/20 working concepts in our library (Appendix S4); toxicity appears
1216 to be encoded in a higher-rank subspace than the late-layer PC1.
- 1217 10. **Arousal asymmetry: vision yes, text and brain no.** OASIS arousal $r = 0.803$ in vision;
1218 $\leq 13\%$ NRC arousal recovery in text; $r \in [0.18, 0.41]$ brain alignment across LLMs
1219 (Appendix S6).
- 1220 11. **LLM-content of KD is irrelevant at this scale.** Replacing the 9-D LLM-derived class
1221 prototypes with random orthonormal 9-D directions changes BACC by ≤ 0.003 . KD
1222 provides architectural regularisation; its semantic content is below our resolution.

1223 S19 Reproducibility

1224 Code, configs, model checkpoints, and figure-generation scripts will be released at the camera-ready
1225 stage under an MIT licence. Each experiment in this paper is logged with its SLURM job ID, conda
1226 environment hash, git SHA, and seed list. The 10-checkpoint ensemble SOTA result is reproducible
1227 from a single SLURM array job (`slurm_d6_e100_e150_5seeds.sh`, ~ 10 GPU-hours per seed on
1228 V100).

1229 S20 Resource Estimate

1230 The full paper used approximately 4,500 GPU-hours on V100 nodes: ~ 600 for the recipe ablation
1231 cascade, $\sim 1,200$ for the 25 V-axis-supervision interventions, ~ 800 for the 36-checkpoint cross-
1232 architecture analysis, $\sim 1,200$ for ensemble construction and val-test rank diagnostics, and ~ 700 for
1233 cross-dataset generality experiments. The V-axis extraction itself is CPU-cheap (~ 10 minutes per
1234 LM on a single 80GB GPU).