
The Saturation Regularity: When Concept-Aligned Supervision Stops Helping Converged EEG Classifiers

Yousef A. Radwan

King Abdullah University of Science and Technology (KAUST)
yousef.radwan@kaust.edu.sa

Xuhui Liu

King Abdullah University of Science and Technology (KAUST)
xuhui.liu@kaust.edu.sa

Kilichbek Haydarov

King Abdullah University of Science and Technology (KAUST)
kilichbek.haydarov@kaust.edu.sa

Yuqian Fu

King Abdullah University of Science and Technology (KAUST)
yuqian.fu@kaust.edu.sa

Mohamed Elhoseiny

King Abdullah University of Science and Technology (KAUST)
mohamed.elhoseiny@kaust.edu.sa

Abstract

A standard way to inject prior knowledge into a neural network is an auxiliary loss that pulls its features toward a known concept direction. We report that this strategy fails once the network has already converged on the concept from task labels alone. On FACED, a nine-class EEG emotion benchmark, twenty-five such recipes—knowledge distillation, representational similarity, contrastive, and topographic losses—applied to a converged classifier improve accuracy in none of the cases and significantly reduce it in fifteen. The effect flips sign with baseline strength: the same valence loss is within seed noise of zero on a weak backbone (+0.006 balanced accuracy, $p=0.41$) but significantly hurts the converged one (−0.029, Holm-corrected $p=0.011$). The effect is specific to neither the concept nor the modality. An arousal target hurts the converged backbone just as much (−0.035, $p=0.007$); and on a CIFAR-10 image classifier, an auxiliary head that significantly helps while the backbone is still training stops helping once it converges—with an aligned target statistically indistinguishable from a random one throughout. We name this the *saturation regularity*: once task labels have driven the classifier onto the target concept, added supervision deforms an already-converged representation without improving generalization—training accuracy stays at ceiling while test accuracy falls. Saturation has a constructive corollary. The class-mean structure of the features is shared across seeds, whereas the within-class residual is seed-specific and load-bearing. Ensembling across this residual diversity, rather than supervising the saturated basin, reaches a new FACED state of the art: 0.6755 balanced accuracy from a single model (+7.4% over the previous best single

model, EMOD at 0.6287) and **0.6948** from a ten-checkpoint ensemble. On a second EEG dataset, SEED-V, the same auxiliary supervision again yields no gain ($\Delta=+0.0018$, paired $p=0.38$, within seed noise), consistent with saturation. To identify what the network has saturated on, we extract the concept explicitly as a one-dimensional valence direction (the *V-axis*) from nine short stories passed through a language model. This direction predicts movie-review sentiment zero-shot (AUC 0.832), tracks the cohort-averaged EEG response to affective videos ($r=0.80$; $r=0.87$ for a CLIP-text axis built from the same stimulus descriptions), and is encoded—without ever being supervised—by the converged EEG classifiers themselves. The transfer is valence-specific but not V-axis-specific—a generic valence axis performs identically—and population-level rather than within-subject. Saturation is thus a property of the converged representation, not of any particular concept or source model.

1 Introduction

Auxiliary losses that pull a network’s features toward a known concept direction are a standard way to inject prior knowledge. We report a regime in which they systematically backfire. Take an EEG emotion classifier that has already converged, and add a loss aligning its features with a valence direction: across twenty-five such recipes on FACED [Chen et al., 2023], none improve balanced accuracy (BACC) and fifteen significantly reduce it. The same loss that is within noise of zero on a weak backbone significantly hurts a strong one. We call this the *saturation regularity*: once task labels alone have driven the network onto the target concept, further supervision can only deform an already-converged representation, without improving generalization. The finding is doubly useful. It explains a widespread failure mode of concept-alignment supervision, and, by locating where the recoverable gain does live—in the seed-specific residual that supervision cannot exploit—it yields a new state of the art through ensembling rather than supervising.

To make “the concept the network has saturated on” concrete, we extract it explicitly. Pass nine short emotion-evocative stories through a language model, average the late-layer hidden states class by class, and take the first principal component; the resulting *valence axis* (V-axis) predicts movie-review sentiment at AUC 0.832 zero-shot, tracks the cohort-averaged EEG response of 123 subjects at the emotion-category level ($r=0.80$), and is encoded, without V-axis supervision, by all ten same-architecture EEG checkpoints we tested. Recent work shows that language models’ internal features align with human cognition across high-level concepts [Kim et al., 2018, Arditi et al., 2024]; we use the V-axis as a *probe* of what brain-decoding networks encode, not as a training signal—which, as the saturation regularity shows, a converged network cannot benefit from.

A nine-story axis. The nine stories are one per FACED [Chen et al., 2023] class; the language model is Qwen3-4B, chosen as the lead because it sits at the centre of the alignment manifold in our 14-LLM sweep, with the highest per-stimulus alignment to a behavioural valence reference (the recipe itself is model-agnostic, see Appendix S3). The SST-2 benchmark grounds the AUC 0.832 headline: the V-axis projection alone is within 0.005 of a 5k-example supervised logistic regression [Socher et al., 2013]. And the direction is not idiosyncratic to any one model: across the 14 language models from 560M to 32B parameters, the V-axis is recovered across families, and within Qwen3 is essentially scale-invariant (off-diagonal $r=+0.995$ over 15 size-pairs).

A valence ordering, consistent with the axis, in cohort EEG. The brain numbers above unpack as follows. A ridge regression on 160 channel–band features (32 EEG channels $\times$ 5 frequency bands, differential-entropy) of the FACED cohort EEG predicts the V-axis position of each of the 28 stimuli; the regressor uses leave-one-stimulus-out cross-validation, no per-subject fitting, no auxiliary supervision. The Qwen3-4B V-axis itself reaches $r=0.80$ ($p < 10^{-5}$) on this ridge; a CLIP-text variant of the V-axis built from the same 28 stimulus descriptions tightens this to $r=0.87$ ($p < 10^{-8}$). The third claim — that converged EEG classifiers encode the V-axis without ever being trained on it — is quantified on the ten same-architecture SOTA-pool checkpoints trained on the 9 emotion labels alone: each encodes the V-axis in its class-mean subspace ($|r| \in [0.60, 0.77]$; §4). The same valence *ordering* thus appears in the LLM features and in the cohort EEG, recovered by independent procedures.

Yet supervision fails. The natural way to use such an alignment is an auxiliary loss pulling the EEG network’s penultimate features toward the V-axis. Across 25 such recipes—knowledge distillation [Hinton et al., 2015], representational similarity [Kriegeskorte et al., 2008], contrastive [Khosla et al., 2020, van den Oord et al., 2018], PEFT adapters [Hu et al., 2022], and topographic losses [Padakanti et al., 2025]—none help: fifteen produce statistically significant decrements and zero produce gains, five families showing *monotonic destruction* (Figure 1). The effect flips sign with baseline strength (no measurable effect at $\text{BACC} \leq 0.62$, a significant negative at $\text{BACC} \geq 0.66$). We call this *saturation*: distinct from overfitting, it is what happens once task labels alone have driven the network onto the target direction, so added supervision moves the class-mean component further along the axis ($\Delta|r|$ up to $+0.36$) while test accuracy falls and training accuracy stays at ceiling.

Where the gain lives. The same diagnosis locates the recoverable gain in the within-class residual that supervision cannot reach: the class-mean basin is shared across seeds, while the residual is seed-specific and load-bearing. Ensembling across seeds averages out residual noise without disturbing the basin—per-checkpoint residual encoding predicts leave-one-out contribution at $r=+0.74$, and a 10-checkpoint ensemble reaches FACED-9 SOTA at **0.6948** (val-selected single **0.6755**), a $+10.5\%$ improvement over EMOD [Chen et al., 2025]. Replacing the language-derived KD prototypes with random orthonormal vectors leaves BACC statistically unchanged ($\Delta=-0.004$, paired $p=0.37$): the gain is the *shape* of the loss, not the *content* of the prototypes, so the V-axis is a probe of what the network saturated on, not the supervised signal.

Contributions. (1) A nine-story *V-axis probe*: a one-dimensional valence direction extracted from nine language-model emotion stories, validated as a genuine semantic axis (text sentiment zero-shot, cohort-EEG prediction, and encoding by converged EEG classifiers) and used to diagnose what the network has saturated on (§6–§7). (2) The *saturation regularity*: concept-aligned auxiliary supervision cannot improve a *converged* classifier and tends to hurt it—established across 25 recipes on FACED-9, for both valence and arousal, and reproduced on a CIFAR-10 classifier, with a benefit present only while the backbone is still weak (§3). (3) A residual-diversity ensemble that turns this mechanism into a positive prescription and a new FACED-9 state of the art (§4, §5). The regularity is consistent on a second EEG dataset, SEED-V, where auxiliary supervision likewise yields no gain (Liu et al., 2022b; Appendix S8).

What this paper does not claim. To keep the scope precise, we state four non-claims up front. First, we do not claim the V-axis is a privileged or uniquely useful direction: a generic valence axis performs identically as a probe (§6), and inside the ensemble recipe the language-derived prototypes are statistically indistinguishable from random orthonormal vectors ($\Delta=-0.004$ BACC, $p=0.37$). Second, we do not claim single-trial or within-subject brain decoding: the brain correlation is measured on the EEG response averaged across 123 subjects at the emotion-category level, and the within-subject effect is near zero ($r=0.03$). Third, we do not claim that auxiliary alignment is universally harmful — on a weak (undertrained) backbone the same valence loss is within seed noise of zero ($+0.006$ BACC, $p=0.41$), and on a still-converging CIFAR-10 classifier an auxiliary head is significantly helpful; saturation is specifically a property of the *converged* regime. Fourth, we do not present saturation as a proven theorem for all architectures and datasets: we establish it empirically on FACED-9 and find it consistent on SEED-V, where auxiliary supervision yields no gain. Limitations are collected in §8.

2 Related Work

Concept-direction supervision. The interventions we test are drawn from seven years of representation-alignment work: knowledge distillation [Hinton et al., 2015, Tian et al., 2020, Park et al., 2019], representational similarity analysis [Kriegeskorte et al., 2008, Sundaram et al., 2024], supervised contrastive losses [Khosla et al., 2020, van den Oord et al., 2018, Radford et al., 2021], and parameter-efficient adaptation [Hu et al., 2022, Liu et al., 2022a]. The shared assumption across these methods is that the auxiliary signal provides information the task does not; this paper studies a regime where it does not.

Concept directions in language models. Arditi et al. [Arditi et al., 2024] showed that refusal in instruction-tuned LMs is mediated by a single residual-stream direction, and related work extracts other behavioural directions [Li et al., 2023, Zou et al., 2023, Kim et al., 2018]; concurrent Anthropic

work [Sofroniew et al., 2026] finds emotion-concept features via sparse autoencoders. We use the same PCA-on-class-centroids construction to extract a valence direction from nine stories as a *probe*, not a steering target; our transfers are direction-existence claims, not feature-importance claims for biological emotion.

Brain-LM alignment. Toneva and Wehbe [Toneva and Wehbe, 2019] showed mid-layer transformer activations predict fMRI responses to narrative text; later work extended this across MEG, ECoG, and model families [Schrimpf et al., 2021, Goldstein et al., 2022, Caucheteux and King, 2022], formalised as the *Platonic Representation Hypothesis* [Huh et al., 2024]. We touch this literature only through the probe; our load-bearing claim is the saturation regularity, not the alignment itself.

Frontal-alpha asymmetry. Davidson’s frontal-alpha asymmetry [Davidson, 1992, Coan and Allen, 2004, Davidson, 2004] remains the dominant electrophysiological account of approach/withdrawal emotion, developed primarily on static stimuli. Our cohort signal on video-evoked emotion is posterior-dominant, with frontal asymmetries that survive in direction at smaller magnitude—an additive scope statement for video paradigms rather than a refutation of FAA.

Ensemble theory. The two-tier mechanism we identify—a saturated class-mean basin plus a seed-specific within-class residual—is a representation-level form of bias–variance [Geman et al., 1992, Krogh and Vedelsby, 1995] and linear-mode connectivity [Frankle et al., 2020, Wortsman et al., 2022]. Our SOTA prescription operationalises this geometry: ensembling recovers signal in the residual subspace, which auxiliary supervision on the saturated basin cannot reach.

Baselines. The FACED-9 baselines we compare against are CBraMod [Wang et al., 2025] (0.572), EmotionKD [Liu et al., 2023] (0.628), and EMOD [Chen et al., 2025] (0.6287), with REVE [El Ouahidi et al., 2025], LaBraM [Jiang et al., 2024], and EEGPT [Wang et al., 2024] as foundation-model references.

3 The Saturation Regularity

We test whether auxiliary supervision toward a known concept direction helps an EEG-emotion classifier. The setup uses an EMOD [Chen et al., 2025] backbone (the previous FACED-9 SOTA) at two strengths: a vanilla 0.62 BACC baseline and the strong SOTA recipe of §5 (augmentation, depth-doubling, distillation, and ensembling) at 0.66 BACC. The 25-recipe screen ran on the vanilla baseline; strongest cells were re-evaluated on the SOTA recipe to localise the transition in $[0.62, 0.66]$. Each intervention adds an auxiliary loss that pulls the network’s penultimate features toward a target direction. Targets span LLM-derived class prototypes (the V-axis—a one-dimensional valence direction extracted from nine LLM emotion stories, §6), CLIP-text embeddings, and frontal-channel masks matched to the analytical valence ceiling. Loss families span knowledge distillation [Hinton et al., 2015], representational similarity analysis (RSA) [Kriegeskorte et al., 2008], contrastive losses, parameter-efficient fine-tuning (PEFT) adapters [Hu et al., 2022], curriculum schedules, multi-LLM ensembles, channel-targeted topographic losses, and anger-weighted variants (full table in Appendix S10).

Result. None help. Fifteen produce statistically significant accuracy decrements ($p < 0.05$, paired across five seeds); zero produce gains. Five families show *monotonic destruction*: the harder we push the auxiliary loss, the worse accuracy gets, with no positive setpoint. Even the anger-weighted target that analytically maximises the V-axis ceiling on FACED is among the worst ($\Delta\text{BACC} = -0.054$): supplying the strongest possible target makes things strictly worse, not better.

The transition is sharp and unidirectional in baseline strength. An intervention whose effect lies within seed noise on a weak baseline ($\text{BACC} \leq 0.62$) becomes a statistically significant negative on the strong recipe ($\text{BACC} \geq 0.66$). The transition sits cleanly in the band $[0.62, 0.66]$ on FACED-9 (we make no claim about universal cut-offs).

The regularity generalises beyond valence and beyond EEG. An *arousal* target (the same recipe re-anchored on arousal poles) hurts the converged backbone just as much (topographic $\Delta\text{BACC} = -0.031$, $p=0.011$; emod-style -0.035 , $p=0.007$; Holm-corrected) while remaining




Figure 1: **Auxiliary concept supervision helps a weak network and hurts a converged one—regardless of the target’s content.** (a) On CIFAR-10, an auxiliary head raises test accuracy while the backbone is undertrained (filled points, Holm $p < 0.05$) but the benefit vanishes once the backbone converges (open points, n.s.; shaded); an aligned target is statistically indistinguishable from a random one throughout (dashed). (b) On FACED EEG the same sign-flip appears for two concept axes: a valence or arousal auxiliary loss is within seed noise of zero on a weak ($d=3$) backbone but significantly negative on the converged SOTA recipe ($p=0.011$ topographic, Holm-corrected; emod-style variants $p \leq 0.008$). Error bars: 95% CI (converged), $\pm$ SEM (weak).

within seed noise on the weak $d=3$ baseline—two concept axes, one saturation (Figure 1b). The weak-to-converged transition also reproduces outside EEG: on CIFAR-10 [Krizhevsky, 2009] an auxiliary head significantly *improves* an undertrained backbone (Holm $p \in [0.006, 0.016]$) and the benefit vanishes at convergence ($p > 0.37$), with an aligned target statistically indistinguishable from a random one throughout (Figure 1a). Across 20 further non-EEG branches (text, vision, audio) an aligned target never significantly beats a random one (0/20): what a converged network rejects, and a weak one accepts, is the loss’s *structure*, not its *content* (Appendix S11).

Saturation (empirical regularity; FACED-9, EMOD backbone). *For a FACED-9 classifier f_θ whose task-only BACC(θ) lies in the band $[0.62, 0.66]$ or higher, adding any V-axis auxiliary loss $\mathcal{L}_v$ at $\lambda > 0$ produces $\mathbb{E}[\Delta \text{BACC}] \leq 0$ across all families tested (15/25 significant at $p < 0.05$, 0/25 positive). The threshold coincides with the regime where the V-axis encoding strength—the absolute correlation $|r|$ between the class-mean features and the V-axis—saturates across seeds (§4). This is an empirical regularity over the recipes tested, not a proved theorem.*

Saturation is not overfitting. Overfitting is memorisation of the training set. Saturation is a property of how task and auxiliary supervision interact once both target the same direction in feature space. The mechanism check below makes the distinction concrete.

Mechanism check. The auxiliary loss does push the network’s class-mean features further along the target direction (alignment increases by $\Delta|r| = +0.01$ to $+0.36$), yet this buys no generalization: with the loss active, training accuracy stays pinned at ceiling while test accuracy *falls* ($0.62 \rightarrow 0.58$ balanced accuracy, A6). The added supervision reshapes a representation the network has already converged onto from task labels alone, rather than supplying new discriminative signal. What actually carries generalization is the *within-class residual*—the directions that distinguish one trial of an emotion from another—which is seed-specific and load-bearing: its per-checkpoint encoding strength predicts leave-one-out ensemble contribution at $r = +0.74$ (§5). Pushing the already-saturated class-mean basin along a fixed target does not improve how the network uses this residual. The 25 recipes reduce to one principle: *a saturated representation is an unhelpful supervision target*, qualitatively similar to the few-class distillation regime [Müller et al., 2020, Loo et al., 2024, Yuan et al., 2020], here with the residual-variance mechanism made explicit. Full table, transition table, and Path-B failure ($\Delta \in [-0.0193, -0.0145]$) in Appendix S10–S11.

4 Mechanism: Saturated Basin, Load-Bearing Residual

§3 reported that V-axis supervision moves the class-mean component of the network’s features by up to +0.36 while test accuracy falls and training accuracy stays at ceiling. We take the class-mean/residual decomposition as the operative geometry and ask what each subspace actually does for the trained network.

The class-mean basin is saturated across seeds. Across the ten same-architecture checkpoints that make up our SOTA ensemble (§5; EMOD $d=6$, five seeds $\times$ two training lengths), the class-mean subspace encodes the V-axis at a consistent strength ($|r| \in [0.60, 0.77]$, unsigned): every well-trained seed of this architecture converges onto the same class-mean direction up to sign. This is the *saturated basin*—the component §3 showed auxiliary supervision cannot improve. What differs across seeds, and what carries the recoverable gain, is the orthogonal within-class residual. (We do *not* claim architecture-independent convergence: pooling V-axis encoding against accuracy across heterogeneous backbones yields an inflated correlation that collapses within any single architecture and does not survive a matched-norm null; see Appendix S12.)

The residual is what averaging recovers. Within-class residual encoding predicts ensemble contribution: across the 10 checkpoints used to build our SOTA ensemble (§5), per-checkpoint residual strength predicts leave-one-out contribution at $r=+0.74$ ($p=0.014$). Bootstrap CIs and per-deletion ranges are in Appendix S12.

Residual Contribution Regularity. *For a saturated EEG classifier with within-class residual encoding of the V-axis at strength ρ_k , the leave-one-out ensemble contribution scales linearly with ρ_k ($p=0.014$, $n=10$, FACED 9-class). $n=10$ is small; we report this as an empirical regularity rather than a law.*

Direct test: ablate the residual at inference time. A directional ablation [Arditi et al., 2024]—projecting the LLM-derived V-axis residual out of the classifier’s penultimate features at inference time, without retraining—drops accuracy on every one of the 10 SOTA-pool checkpoints (mean $\Delta\text{BACC} = -0.0157$, mean $z \approx 7.7$ above a matched random-direction null). The residual carries the load. This decomposition—saturated basin, load-bearing residual—is the operative geometry that §3 predicted from the failure of V-axis supervision and that §5 turns into a positive prescription.

5 Two-Tier Ensemble: The Positive Prescription

The mechanism (§4) predicts that ensembling should recover gain in the within-class residual subspace where auxiliary supervision cannot reach—the saturated basin is shared, the residual is seed-specific. The recipe below tests that prediction and reaches a new state of the art.

Method	BACC
CBraMod [Wang et al., 2025]	0.572
EmotionKD [Liu et al., 2023]	0.628
EMOD [Chen et al., 2025]	0.6287
EMOD+aug+KD+ $d=6$ (5-seed)	$.658 \pm .010$
Best single (val-sel)	0.6755
10-ckpt ensemble (SOTA)	0.6948

(a) Headline. All baselines are single-model: our best single model reaches 0.6755 (+7.4% over EMOD) and the 10-checkpoint ensemble 0.6948 (+10.5%).

Recipe step	BACC	Δ
EMOD ($d=3$) replication	$.619 \pm .004$	—
+ augmentation	$.634 \pm .004$	+ .015
+ aug + KD	$.644 \pm .007$	+ .010*
+ depth doubling ($d=6$)	$.658 \pm .007$	+ .014
+ longer training ($e=150$)	$.658 \pm .010$	.000
5-seed e100 ensemble	.6798	+ .022
10-ckpt e100+e150 (SOTA)	.6948	+ .037

(b) Recipe cascade (*super-add., $p=3 \times 10^{-4}$).

Table 1: New FACED 9-class SOTA. The rand9 ablation replaces LLM-derived KD prototypes with random orthonormal 9-D directions and leaves BACC statistically unchanged ($\Delta=-0.004$, paired $p=0.37$) — KD provides architectural regularisation without semantic content. Visual cascade in Appendix S13.

The cascade starts at the EMOD replication (0.6194 ± 0.004 , within seed noise of the published 0.6287 [Chen et al., 2025]) and walks up in six steps. Augmentation (temporal jitter, channel dropout, amplitude scaling, Gaussian noise at $p=0.6$) contributes the largest single jump ($+0.0149$). Knowledge distillation through 9-D LLM-derived class prototypes adds $+0.0096$; a rand9 ablation (identical recipe with the 9 LLM prototypes replaced by random orthonormal directions) leaves BACC statistically unchanged ($\Delta=-0.004$, paired $p=0.37$), so KD imposes a 9-direction frame on the features rather than transferring the LLM’s emotion concept. Depth-doubling to $d=6$ adds $+0.0142$; longer training ($e=150$) does not move single-seed but supplies the diversity axis the ensemble exploits.

Why the ensemble works. Ensemble gain concentrates in the within-class residual each seed encodes differently ($r=+0.74$, $p=0.014$; §4); averaging cancels that seed-specific noise without disturbing the shared basin. Mixing $e=100$ and $e=150$ checkpoints supplies the diversity axis—the mixed pool reaches 0.6948 versus 0.6798 for $e=100$ alone. Headroom-monotonicity, the 10-checkpoint plateau, and the failure of cross-architecture mixing are in Appendix S13.

6 Probing the Saturated Concept: A Valence Direction in LLMs

To test that the saturation regularity operates on a real semantic axis rather than a noise dimension, we extract one explicitly. The construction uses nine stories. Given a language model, we author one short emotion-evocative story per FACED class and generate 50 paraphrases per story by independent LLM calls. For each class we average the language model’s last-token activation (at the penultimate layer) over all 50 paraphrases; the resulting vector is the class *centroid*—a single point in feature space that summarises the model’s representation of that emotion. Principal component analysis (PCA) on the nine centroids gives the V-axis as the first principal component, oriented so the Joy centroid projects positive. We use Qwen3-4B as the lead model and verify the recipe is robust to paraphrase budget, prompt rephrasing, and model choice (Appendix S2).

The same direction across modalities. The V-axis recipe recovers human affect across three settings that share no training data (Table 2): the text result is true zero-shot projection; the EEG and vision results use the same nine-story recipe re-extracted in their native modality plus a small linear probe.

Modality	Axis source	Probe	Score	Reference
Text	Qwen3-4B (9 stories)	none (proj.)	0.832 AUC	matches sup. LR on 5k examples
Brain	Qwen3-4B (9 stories)	ridge LOOCV	r=0.80	$p < 10^{-5}$, see §7
Vision	CLIP-image (9 images)	none (proj.)	r=0.869	beats sup. ridge (image-side)

Table 2: Cross-modal external validity of the V-axis recipe (PCA on nine emotion-class centroids). The Qwen3-4B V-axis is the lead model for both Text and Brain rows. Vision uses CLIP-image because the language model has no image side; a CLIP-text variant of the V-axis on the same EEG ridge reaches $r=0.87$ ($p < 10^{-8}$, §7). Full sentiment, lexicon, and multilingual results in Appendix S3.

The LLM-side direction is a strong sentiment classifier: 0.832 SST-2 AUC matches supervised LR on 5k examples (0.837), beats SBERT prototype-cosine (0.793 ; Reimers and Gurevych, 2019), and trails 355M RoBERTa-MNLI zero-shot (0.912 ; Liu et al., 2019) by 0.08 —all from nine stories of supervision.

The direction is recovered across model families. Extracting the V-axis independently from four non-Qwen families recovers essentially the same valence ordering of the stimuli: Mistral-7B ($r=0.95$), Llama-3 ($r=0.95$), and BLOOM ($r=0.94$) all align with the behavioural valence reference, with Pythia the lone outlier ($r=0.44$). Within the Qwen3 family the V-axis is essentially scale-invariant (off-diagonal $r=+0.995$ across 15 size-pairs); across families the raw axes are looser ($r=+0.585$ over 48 pairs). We therefore describe the valence direction as *recoverable* by modern LM training rather than as a single shared geometry—a cross-model similarity that, tested directly, is largely inherited from shared input structure rather than learned (Appendix S3).

Specificity. The recipe is concept-generic (20 further concepts produce a working axis, 17 exceed AUC 0.95; Appendix S4). Nonce-word ablation drops SST-2 to chance and random-Gaussian directions give chance lexicon recovery, so the V-axis is a real semantic direction, not a noise artefact (Appendix S6). An Arditi-style same-model ablation [Arditi et al., 2024] on Qwen3-4B quantifies how much sentiment lives along the V-axis: a logistic-regression probe on the full features reaches SST-2 AUC 0.909, drops to 0.907 after projecting the V-axis out and retraining ($z=+9.8$ above a 20-direction random-direction null). The V-axis carries detectable sentiment signal but is not the sole sentiment-bearing direction.

7 The Probe Predicts Cohort EEG

If the converged classifier has saturated on a genuine valence axis, that axis should be measurable in the brain signal itself, not only inside the classifier. We test this in the FACED cohort [Chen et al., 2023]: 123 subjects watching 28 emotional video clips, 32 EEG channels at 250 Hz. From the cohort EEG averaged over subjects and time, we fit a single *ridge regression* (a linear model with L_2 regularisation) predicting the V-axis projection of each stimulus’s textual description from the 160 channel-band features (differential entropy over 32 channels $\times$ 5 frequency bands). The ridge prediction matches the Qwen3-4B V-axis at $r=+0.80$ ($n=28$ stimuli, $p < 10^{-5}$). Substituting the CLIP-text V-axis on the same 28 stimulus descriptions reaches $r=+0.87$ ($p < 10^{-8}$; Figure 2)—the brain signal aligns with the LLM-derived direction across two distinct LLM substrates. Four controls support the brain-LM correlation: a fully nested permutation test (stimulus relabelling within the subject structure, $N=5000$) places the correlation beyond every permutation draw ($p \leq 2 \times 10^{-4}$, random-direction $r=0.07$); a matched-norm random direction drops the same ridge to $r=+0.07$ ($p=0.73$); split-half subject reliability on per-stimulus valence is $r=+0.99$; the same EEG ridge predicts human-rated valence at $r=+0.86$, matching the V-axis prediction to within 0.01—the brain signal is at the same strength a panel of human raters would produce. Multi-comparison correction and other controls are in Appendix S7.

The ridge correlation is not carried by a few high-affect stimuli. Because a single-channel analysis of this dataset (PO3/ γ , Appendix S9) is dominated by an anger-versus-warm-positive contrast, it is worth stating precisely what does and does not depend on those poles. Removing the three emotional-pole categories (Anger, Amusement, Tenderness) leaves the *ridge* correlation essentially unchanged ($r=0.88$ on the remaining 19 mid-valence stimuli, versus 0.87 on all 28, both with the CLIP-text axis), and a leave-one-emotion-out jackknife keeps it in the range $r \in [0.86, 0.89]$ (every single-emotion drop stays significant). The same pole ablation does collapse the single raw PO3/ γ channel ($0.48 \rightarrow -0.02$). The pole dependence therefore belongs to the single-channel topography, not to the multi-feature ridge that carries the headline.

The brain aligns with valence, not with the V-axis specifically. Swapping the LLM-derived V-axis for a generic categorical valence axis (an integer valence rank assigned to each emotion) or for the behavioural valence ratings leaves the ridge correlation essentially unchanged ($r=0.89$ and $r=0.86$ respectively, versus 0.87 for the CLIP-text V-axis; all $n=28$), while a matched-norm random direction stays at chance ($r=0.07$; Appendix S7). What the cohort EEG encodes is a valence *ordering* of the stimuli; the V-axis is one instance of that ordering, not a privileged direction. This is why we treat the V-axis as a probe of the saturated concept rather than as the concept itself.

The same picture survives across all 14 LLMs (Qwen3 family mean $r=+0.796 \pm 0.011$ irrespective of size) and on the held-out SEED-V dataset [Liu et al., 2022b], where the FACED-derived V-axis ranks the five SEED-V emotions at $r=+0.96$ without retraining. On a third, independently collected EEG dataset (EmoEEG-MC [Xu et al., 2025]), a within-subject valence decoder reaches 0.61 accuracy ($p=3.4 \times 10^{-5}$) and its group axis tracks behavioural valence at $r=0.71$; an audio/visual confound control there is inconclusive (no partial survives correction), so we frame the EEG–valence link as a population-level predictive association, not a confound-free one (Appendix S7). Appendix S8 re-derives the V-axis from scratch on SEED-V and re-tests every claim.

Probe scope. The cohort signal is posterior-dominant and is carried by an anger-versus-warm-positive contrast on 9 of the 28 clips, so at the single-channel level it reads more as an emotion-contrast than a graded valence axis (the multi-feature ridge above is not pole-dependent). It is a dynamic-




Figure 2: Stimulus-level alignment between the LLM-derived valence direction and human EEG. Each point is one of the 28 FACED video stimuli; x is the *CLIP-text* V-axis projection of the textual stimulus description, y is the single ridge-regression prediction of that projection from all 160 cohort channel-band features (32 channels $\times$ 5 frequency bands, averaged over 123 subjects and each clip). Cohort Pearson $r = +0.87$ ($n=28$, $p < 10^{-8}$); the nine-story Qwen3-4B V-axis reaches $r=0.80$ on the same ridge. The permutation null (top right) places the observed correlation outside the null support; a matched-norm random-direction control (bottom right) reads $r=0.07$. The axis is extracted with no EEG supervision.

video-paradigm statement that neither tests nor refutes frontal-alpha asymmetry [Davidson, 1992]; full topography, Simpson’s-paradox, time-resolved, and FAA analyses are in Appendix S9.

8 Discussion: Limitations and Future Work

Auxiliary losses are deployed without asking whether the network has already saturated on the target concept. Our results—a saturation regularity organising 25 recipe failures, a residual-diversity ensemble that turns the mechanism into FACED-9 SOTA, and a nine-story V-axis probe—are scoped to FACED-9 with a measurable V-axis substrate. The probe transfers are direction-existence results, not neuroscience mechanism claims; the posterior topography speaks to the video-viewing paradigm, not to the frontal-asymmetry account [Davidson, 1992]; and the cohort signal is population-level, not within-subject (Appendix S9). Three limitations bound the broader reading. First, we do not claim architecture-independent convergence: a V-axis-encoding-versus-accuracy correlation pooled across backbones is pseudo-replicated (it collapses within any single architecture and fails a matched-norm null), so convergence is scoped to seeds of one architecture. Second, the cross-model similarity of the LLM valence geometry is largely inherited from shared input structure rather than learned—an untrained model reproduces it—so we treat the V-axis as a convenient probe, not a universal learned representation. Third, the transfer is content-specific (valence carries across text, vision, and EEG; arousal only in vision), the weak-to-converged transition is cleanest on CIFAR-10 [Krizhevsky, 2009] (text shows content-irrelevance but not a significant help-when-weak effect), and the EmoEEG-MC confound control is inconclusive. The practical corollary is that aux-loss design papers should report the receiving baseline’s saturation state before attributing a gain to the loss. Whether the residual-vs-basin decomposition generalises to other saturated concepts—truthfulness, refusal, sentiment polarity—is the question this paper opens.

References

- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, 2024.
- Ryan T Canolty and Robert T Knight. The functional role of cross-frequency coupling. *Trends in Cognitive Sciences*, 14(11):506–515, 2010. doi: 10.1016/j.tics.2010.09.001.
- Charlotte Caucheteux and Jean-Rémi King. Brains and algorithms partially converge in natural language processing. *Communications Biology*, 5:134, 2022. doi: 10.1038/s42003-022-03036-1.
- Jingjing Chen, Xiaowei Wang, Chen Huang, et al. A large finer-grained affective computing EEG dataset. *Scientific Data*, 10:740, 2023. doi: 10.1038/s41597-023-02650-w.
- Yuning Chen, Sha Zhao, Shijian Li, and Gang Pan. EMOD: A unified EEG emotion representation framework leveraging v-a guided contrastive learning. *arXiv preprint arXiv:2511.05863*, 2025. To appear at AAAI 2026.
- James A Coan and John J B Allen. Frontal EEG asymmetry as a moderator and mediator of emotion. *Biological Psychology*, 67(1-2):7–50, 2004. doi: 10.1016/j.biopsycho.2004.03.002.
- Maurizio Codispoti, Andrea De Cesarei, and Vera Ferrari. Alpha-band oscillations and emotion: A review of studies on picture perception. *Psychophysiology*, 60(11):e14438, 2023. doi: 10.1111/psyp.14438.
- Richard J Davidson. Anterior cerebral asymmetry and the nature of emotion. *Brain and Cognition*, 20(1): 125–151, 1992. doi: 10.1016/0278-2626(92)90065-T.
- Richard J Davidson. Well-being and affective style: neural substrates and biobehavioural correlates. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 359(1449):1395–1411, 2004. doi: 10.1098/rstb.2004.1510.
- Yassine El Ouahidi, Jonathan Lys, Philipp Thölke, Nicolas Farrugia, Bastien Padeloup, Vincent Gripon, Karim Jerbi, and Giulia Lioi. REVE: A foundation model for EEG — adapting to any setup with large-scale pretraining on 25,000 subjects. In *Advances in Neural Information Processing Systems*, 2025. arXiv:2510.21585.
- Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M Roy, and Michael Carbin. Linear mode connectivity and the lottery ticket hypothesis. In *International Conference on Machine Learning*, 2020. arXiv:1912.05671.
- Stuart Geman, Elie Bienenstock, and René Doursat. Neural networks and the bias/variance dilemma. *Neural Computation*, 4(1):1–58, 1992. doi: 10.1162/neco.1992.4.1.1.
- Ariel Goldstein, Zaid Zada, Eliav Buchnik, Mariano Schain, Amy Price, Bobbi Aubrey, Samuel A Nastase, Amir Feder, Dotan Emanuel, Alon Cohen, Aren Jansen, Harshvardhan Gazula, Gina Choe, Aditi Rao, Catherine Kim, Colton Casto, Lora Fanda, Werner Doyle, Daniel Friedman, Patricia Dugan, Lucia Melloni, Roi Reichart, Sasha Devore, Adeem Flinker, Liat Hasenfratz, Omer Levy, Avinatan Hassidim, Michael Brenner, Yossi Matias, Kenneth A Norman, Orrin Devinsky, and Uri Hasson. Shared computational principles for language processing in humans and deep language models. *Nature Neuroscience*, 25(3):369–380, 2022. doi: 10.1038/s41593-022-01026-4.
- Greg Hajcak, Annmarie MacNamara, and Doreen M Olvet. Event-related potentials, emotion, and emotion regulation: An integrative review. *Developmental Neuropsychology*, 35(2):129–155, 2010. doi: 10.1080/87565640903526504.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Minyoung Huh, Brian Cheung, Tongzhou Wang, and Phillip Isola. Position: The platonic representation hypothesis. In *International Conference on Machine Learning*, 2024. arXiv:2405.07987.
- Wei-Bang Jiang, Li-Ming Zhao, and Bao-Liang Lu. Large brain model for learning generic representations with tremendous EEG data in BCI. In *International Conference on Learning Representations*, 2024. arXiv:2405.18765.
- Prannay Khosla, Piotr Teterwak, Chen Wang, Aaron Sarna, Yonglong Tian, Phillip Isola, Aaron Maschinot, Ce Liu, and Dilip Krishnan. Supervised contrastive learning. In *Advances in Neural Information Processing Systems*, 2020.

- Been Kim, Martin Wattenberg, Justin Gilmer, Carrie Cai, James Wexler, Fernanda Viegas, and Rory Sayres. Interpretability beyond feature attribution: Quantitative testing with concept activation vectors (TCAV). In *International Conference on Machine Learning*, 2018. arXiv:1711.11279.
- Sander Koelstra, Christian Muhl, Mohammad Soleymani, Jong-Seok Lee, Ashkan Yazdani, Touradj Ebrahimi, Thierry Pun, Anton Nijholt, and Ioannis Patras. DEAP: A database for emotion analysis using physiological signals. *IEEE Transactions on Affective Computing*, 3(1):18–31, 2011.
- Nikolaus Kriegeskorte, Marieke Mur, and Peter Bandettini. Representational similarity analysis — connecting the branches of systems neuroscience. *Frontiers in Systems Neuroscience*, 2, 2008.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009. CIFAR-10 / CIFAR-100 datasets.
- Anders Krogh and Jesper Vedelsby. Neural network ensembles, cross validation, and active learning. In *Advances in Neural Information Processing Systems*, pages 231–238, 1995.
- Benedek Kurdi, Shayn Lozano, and Mahzarin R Banaji. Introducing the open affective standardized image set (OASIS). *Behavior Research Methods*, 49(2):457–470, 2017. doi: 10.3758/s13428-016-0715-3.
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems*, 2023. arXiv:2306.03341.
- Haokun Liu, Derek Tam, Mohammed Muqeeth, Jay Mohta, Tenghao Huang, Mohit Bansal, and Colin Raffel. Few-shot parameter-efficient fine-tuning is better and cheaper than in-context learning. In *Advances in Neural Information Processing Systems*, 2022a. arXiv:2205.05638.
- Wei Liu, Jie-Lin Qiu, Wei-Long Zheng, and Bao-Liang Lu. Comparing recognition performance and robustness of multimodal deep learning models for multimodal emotion recognition. *IEEE Transactions on Cognitive and Developmental Systems*, 14(2):715–729, 2022b. doi: 10.1109/TCDS.2021.3071170.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
- Yucheng Liu, Ziyu Jia, and Haichao Wang. EmotionKD: A cross-modal knowledge distillation framework for emotion recognition based on physiological signals. In *Proceedings of the 31st ACM International Conference on Multimedia*, pages 6122–6131, 2023. doi: 10.1145/3581783.3612277.
- Noel Loo, Fotis Iliopoulos, Wei Hu, and Erik Vee. LELP: Linear projections of teacher embeddings for few-class distillation. *arXiv preprint arXiv:2409.20449*, 2024.
- Niklas Muennighoff, Thomas Wang, Lintang Sutawika, Adam Roberts, Stella Biderman, Teven Le Scao, M Saiful Bari, Sheng Shen, Zheng Xin Yong, Hailey Schoelkopf, Xiangru Tang, Dragomir Radev, Alham Fikri Aji, Khalid Almubarak, Samuel Albanie, Zaid Alyafeai, Albert Webson, Edward Raff, and Colin Raffel. Crosslingual generalization through multitask finetuning. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15991–16111, 2023. arXiv:2211.01786.
- Matthias M Müller, Andreas Keil, Thomas Gruber, and Thomas Elbert. Processing of affective pictures modulates right-hemispheric gamma band EEG activity. *Clinical Neurophysiology*, 110(11):1913–1920, 1999. doi: 10.1016/S1388-2457(99)00151-0.
- Rafael Müller, Simon Kornblith, and Geoffrey Hinton. Subclass distillation. *arXiv preprint arXiv:2002.03936*, 2020.
- Srijith Padakanti, Khushbu Pahwa, Radhika Mamidi, Bapi Raju Surampudi, Manish Gupta, and Subba Reddy Oota. Aligning text/speech representations from multimodal models with MEG brain activity during listening. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025. ACL Anthology 2025.emnlp-main.1748.
- Wonpyo Park, Dongju Kim, Yan Lu, and Minsu Cho. Relational knowledge distillation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*, 2021.

- Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, 2019.
- Martin Schrimpf, Idan A Blank, Greta Tuckute, Carina Kauf, Eghbal A Hosseini, Nancy Kanwisher, Joshua B Tenenbaum, and Evelina Fedorenko. The neural architecture of language: Integrative modeling converges on predictive processing. *Proceedings of the National Academy of Sciences*, 118(45):e2105646118, 2021. doi: 10.1073/pnas.2105646118.
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 1631–1642, 2013.
- Nicholas Sofroniew, Isaac Kauvar, William Saunders, Runjin Chen, Tom Henighan, Sasha Hydrice, Craig Citro, Adam Pearce, Julius Tarnig, Wes Gurnee, Joshua Batson, Sam Zimmerman, Kelley Rivoire, Kyle Fish, Chris Olah, and Jack Lindsey. Emotion concepts and their function in a large language model. *Transformer Circuits*, 2026. URL <https://transformer-circuits.pub/2026/emotions/index.html>.
- Shobhita Sundaram, Stephanie Fu, Lukas Muttenthaler, Netanel Y Tamir, Lucy Chai, Simon Kornblith, Trevor Darrell, and Phillip Isola. When does perceptual alignment benefit vision representations? In *Advances in Neural Information Processing Systems*, 2024. arXiv:2410.10817.
- Yonglong Tian, Dilip Krishnan, and Phillip Isola. Contrastive representation distillation. In *International Conference on Learning Representations*, 2020. arXiv:1910.10699.
- Mariya Toneva and Leila Wehbe. Interpreting and improving natural-language processing (in machines) with natural language-processing (in the brain). In *Advances in Neural Information Processing Systems*, 2019. arXiv:1905.11833.
- Parsa Vahidi, Omid G Sani, and Maryam M Shanechi. Modeling and dissociation of intrinsic and input-driven neural population dynamics underlying behavior. *Proceedings of the National Academy of Sciences*, 121(7): e2212887121, 2024. doi: 10.1073/pnas.2212887121.
- Aäron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- Guangyu Wang, Wenchao Liu, Yuhong He, Cong Xu, Lin Ma, and Haifeng Li. EEGPT: Pretrained transformer for universal and reliable representation of EEG signals. In *Advances in Neural Information Processing Systems*, 2024.
- Jiquan Wang, Sha Zhao, Zhiling Luo, Yangxuan Zhou, Haiteng Jiang, Shijian Li, Tao Li, and Gang Pan. CBraMod: A criss-cross brain foundation model for EEG decoding. In *International Conference on Learning Representations*, 2025. arXiv:2412.07236.
- Mitchell Wortsman, Gabriel Ilharco, Samir Yitzhak Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes, Ari S Morcos, Hongseok Namkoong, Ali Farhadi, Yair Carmon, Simon Kornblith, and Ludwig Schmidt. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In *International Conference on Machine Learning*, 2022. arXiv:2203.05482.
- X. Xu, X. Shen, X. Chen, Q. Zhang, S. Wang, Y. Li, Z. Li, D. Zhang, M. Zhang, and Q. Liu. A multi-context emotional EEG dataset for cross-context emotion decoding. *Scientific Data*, 12:1142, 2025. doi: 10.1038/s41597-025-05349-2.
- Li Yuan, Francis E. H. Tay, Guilin Li, Tao Wang, and Jiashi Feng. Revisiting knowledge distillation via label smoothing regularization. In *CVPR*, 2020. arXiv:1909.11723.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state the paper’s claims, ordered as in the paper: (i) a saturation regularity documented across 25 alignment-intervention families on FACED-9 (15 significant decrements, 0 gains), with a sign-flip in the receiving baseline’s strength (weak EMOD-d3 +0.006, n.s. vs. converged recipe -0.029 , Holm $p=0.011$), reproduced for an arousal target and on a CIFAR-10 classifier; (ii) a residual-diversity ensemble giving a new FACED 9-class SOTA at 0.6948 BACC (val-selected single checkpoint 0.6755), with per-checkpoint within-residual encoding predicting leave-one-out ensemble contribution at $r=+0.74$ ($p=0.014$); (iii) a nine-story valence axis (V-axis) used as a probe, which predicts text sentiment zero-shot (AUC 0.832) and tracks the cohort-averaged FACED EEG response at the emotion-category level ($r=0.87$); (iv) generality of the saturation regularity to a second concept (arousal) and to a converged CIFAR-10 classifier, with content-irrelevance across 20 non-EEG branches (0/20 aligned beats random). Transfer is stated as valence-specific but not V-axis-specific and population- rather than within-subject. Each claim is supported in Sections 3–7 with its primary table or figure; the saturation regularity additionally replicates on SEED-V (Appendix S8).

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 3 contains a boxed scope statement (“Saturation (FACED-9, EMOD backbone, V-axis target)”) naming the empirical-not-mathematical scope and the FACED-specific BACC threshold [0.62, 0.66]. Section 8 explicitly states the regularity is scoped to FACED-9 with a measurable V-axis substrate, with no claim beyond, and that probe transfers (SST-2, cohort EEG, classifier convergence) are direction-existence rather than neuroscience claims. Posterior-topography is framed as a video-paradigm statement rather than an FAA refutation. The cross-architecture correlation is reported in Appendix S12 as pseudo-replicated (it collapses within any single architecture and does not survive a matched-norm null); the main text claims only within-architecture basin saturation across seeds of one backbone. Negative results catalogue (theta-gamma PAC, arousal-axis brain-side failure, Path-B mixing) in Appendix S11.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.

- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper is empirical. The “saturation regularity” (Section 3) is stated explicitly as an empirical pattern, not a mathematical theorem; we do not claim a proof. The boxed Statement of the Regularity in Section 3 writes out what would be the formal claim and the empirical pillars supporting it (15 / 25 statistically significant negatives, 5 monotonic-destruction families, unidirectional saturation transition, direct mechanism check). The Residual Contribution Regularity in Section 4 is an empirical regression claim (slope $\hat{\beta} = +0.74$, $p = 0.014$, $n = 10$), not a theorem.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The V-axis extraction protocol (story prompts, paraphrase generation, hidden-state extraction, PCA-PC1 with Joy-positive orientation) is documented in Section 6 and Appendix S2. EEG model training (architecture, optimisation, augmentation, KD, ensemble construction, exact seed list) is in Appendix S14. Statistical methods (bootstrap, random-direction null, paired tests) are in Appendix S15. Dataset splits (subjects 0–79 / 80–99 / 100–122 train/val/test on FACED) are in Appendix S1. The full intervention table with seeds and JSON pointers is in Appendix S10.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: We do not release code or model checkpoints with this submission. The two datasets used (FACED [Chen et al., 2023] and SEED-V [Liu et al., 2022b]) are publicly available from their original authors under their published licences; we have not modified or re-released them. The full methodology, hyperparameters, optimisation details, augmentation schedule, exact seed list, and SLURM-array recipe required to reproduce every result in the paper are documented in the appendix (Appendix S14, S15, S1, S10, S19). Code, configs, and the 10-checkpoint ensemble checkpoints will be released upon acceptance.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.

- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [\[Yes\]](#)

Justification: Optimiser (AdamW, $\text{lr} = 10^{-3}$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, weight decay 10^{-2}), schedule (cosine + 5-epoch warmup), batch size 128, training length $e \in \{100, 150\}$ epochs, depth $d \in \{3, 6\}$, filter dimension $f = 128$, seeds $\{42, 123, 456, 789, 2025\}$, augmentation parameters (Gaussian $\sigma = 0.05 \cdot \text{std}$, channel dropout $p = 0.15$, temporal masking $\leq 5\%$, all at $p = 0.6$), KD ($\lambda_{\text{KD}} = 0.5$, $T = 1.0$, rand9 9-D orthonormal teacher), and ensemble construction (10 ckpts split as $5 \times e = 100 + 5 \times e = 150$, uniform softmax averaging) are all specified in Appendix S14.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [\[Yes\]](#)

Justification: We report standard deviation across 5 seeds in every recipe-cascade row (Section 5); Pearson r with p -values for every cross-architecture and brain correlation (Sections 7, 4); paired t -tests for V-axis interventions vs. matched-recipe baselines (Section 3); $B = 10,000$ subject-resampled bootstrap CIs and Fisher- z p -values for cohort EEG correlations (Appendix S15); 200-direction matched-norm random-direction null distributions for the cohort EEG-LLM correlation; 1000-direction null for the cross-architecture V-axis correlation; matched-direction null for the inference-time directional ablation (Section 4, mean $z \approx 7.7$ across the 10 SOTA-pool checkpoints).

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include experiments.
- The authors should answer [\[Yes\]](#) if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.

- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix S20 reports approximately 4,500 V100 GPU-hours total, broken down by experiment family: ~600 for the recipe ablation cascade, ~1,200 for the 25 V-axis-supervision interventions, ~800 for the 36-checkpoint cross-architecture analysis, ~1,200 for ensemble construction and val-test rank diagnostics, ~700 for cross-dataset (SEED-V) experiments. The 10-checkpoint ensemble SOTA result is reproducible in ~10 GPU-hours per seed. V-axis extraction is CPU-cheap (~10 minutes per LM on a single 80 GB GPU). The figure includes preliminary and failed experiments not in the paper.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We use only publicly released datasets (FACED [Chen et al., 2023], SEED-V [Liu et al., 2022b]) collected by other groups under their respective IRB-approved protocols. We collect no new human data. We use only publicly released LLMs under their published licences.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Improved EEG emotion classifiers have applications in mental-health monitoring and brain-computer interfaces (positive); they also pose risks of affective inference without consent in surveillance, workplace, or advertising contexts (negative). The V-axis extraction protocol applied to EEG could in principle reduce the calibration-data burden for affective inference systems — a feature that cuts both ways. We do not release human EEG data at any stage; the planned camera-ready code release covers training configs, feature extraction, and analysis only. Full broader-impacts statement in Appendix S16.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: We release no high-risk artefacts with this submission. The planned camera-ready release (Appendix S19) covers training configs, feature-side and analysis code, and the trained EEG-emotion checkpoints; it does not introduce abuse vectors beyond what is already public via FACED, SEED-V, and the listed pre-trained LLMs. We release neither human EEG data nor LLM weights at any stage. Dual-use risk assessment and recommended deployment safeguards are in Appendix S16.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: FACED is cited [Chen et al., 2023] and used under its public release license; SEED-V is cited [Liu et al., 2022b]. All 14 LLMs (Qwen3 family of six dense sizes from 0.6B to 32B, Llama-4-Scout-17B, Mistral-7B, Phi-2-2.7B, Pythia-1.4B, BLOOM-560M, TinyLlama-1.1B, Gemma-27B, Gemma-4-31B) are cited and used under their respective

Hugging Face / model-card licenses. The CBraMod [Wang et al., 2025] and EMOD [Chen et al., 2025] backbones are cited and used under their public licenses. CLIP [Radford et al., 2021] is cited.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [No]

Justification: We do not release new assets with this submission. The V-axis extraction scripts, V-axis-aligned EEG checkpoints, and figure-generation pipelines will be released upon acceptance with documentation. Within the paper, the V-axis itself — nine FACED-class centroids and the PCA-PC1 direction at the LM’s penultimate layer — is fully specified by the protocol in Appendix S2 (model identities, layer indices, story prompts, paraphrase parameters), so any reader can re-derive the V-axis without our checkpoint files.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: We do not collect new human data; all human EEG comes from the publicly released FACED [Chen et al., 2023] and SEED-V [Liu et al., 2022b] datasets, originally collected by their respective groups under their own IRB-approved protocols.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.

- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No new human-subjects research was conducted. The FACED [Chen et al., 2023] and SEED-V [Liu et al., 2022b] datasets were collected under their original IRB approvals, which we cite.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: LLMs are central to the V-axis extraction protocol. The 14 language models from which we extract the V-axis are named in Section 6 and Appendix S2: Qwen3 at all six dense sizes (lead model Qwen3-4B; also 0.6B/1.7B/8B/14B/32B), Llama-4-Scout-17B, Mistral-7B, Phi-2-2.7B, Pythia-1.4B, BLOOM-560M, TinyLlama-1.1B, Gemma-27B, Gemma-4-31B. The penultimate-layer hidden state ($L = L_{\max} - 1$, e.g., $L = 35$ of 36 for Qwen3-4B) is the source of all class centroids. We also use a generic instruction-tuned Qwen LM (temperature 0.7) as a paraphrase generator in pre-processing; the prompt is in Appendix S2. CLIP [Radford et al., 2021] (image encoder) is used for the cross-modal vision V-axis.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

Supplementary Material

S1 Datasets, Splits, and Preprocessing

FACED. 123 subjects, 28 emotional video clips covering 9 emotion classes. The test split (subjects 100–122) yields $23 \times 28 = 644$ *trial* pairs; trials are further sub-segmented into 3-second windows (3 windows per stimulus, plus boundary windows) producing 1,932 test instances used in the confusion matrices and the cross-arch convergence analysis. The cohort EEG ridge of §7 aggregates back to the 28-stimulus level. Original FACED preprocessing detail follows below: 9 emotion categories (Anger, Disgust, Fear, Sadness, Neutral, Amusement, Inspiration, Joy, Tenderness) at 3–4 stimuli per class. EEG: 32 channels, 250 Hz sampling, 30-second clip duration. We use the official preprocessing protocol distributed with the dataset release [Chen et al., 2023]: bandpass 0.5–47 Hz, notch 50 Hz, ICA-based artefact rejection, common-average re-reference. Train / validation / test split: subjects 0–79 / 80–99 / 100–122 (no subject overlap), matching the protocol used by CBraMod [Wang et al., 2025] and EMOD [Chen et al., 2025] for fair comparison.

SEED-V. 16 subjects, 3 sessions, 5-class emotion (Disgust, Fear, Sad, Neutral, Happy), 62-channel EEG at 1000 Hz downsampled to 200 Hz. Standard SEED-V split with subject-disjoint train / test. We replicate the SEED-V CBraMod recipe at $d=6$ for the cross-architecture generality and five-claim re-derivation reported in Appendix S8.

DE features. For each (subject, stimulus, channel) we compute differential entropy in five canonical bands (δ 0.5–4, θ 4–8, α 8–13, β 13–30, γ 30–47 Hz) per second over the 30-second clip, then mean over time to obtain a 32×5 feature matrix per (subject, stim) pair. Cohort-level analyses average across the 123 subjects to give the $28 \times 32 \times 5$ tensor used in Section 7.

S2 V-Axis Extraction Protocol

Story prompts. Nine emotion-evocative stories (one per FACED class), each 1–3 sentences. Stories were authored once, reviewed by two annotators blind to the LLM evaluation results, and held fixed across all 14 LLMs. Verbatim text:

- **Anger.** A driver cuts you off in heavy traffic and laughs through the window. Your hands tighten on the wheel and you feel heat rise in your chest.
- **Disgust.** You open the fridge and find a container of leftovers covered in green fuzz, with a sour smell that makes you step back.
- **Fear.** Walking home alone at night, you hear footsteps quicken behind you and notice the streetlights are out.
- **Sadness.** You sort through old photographs of a friend who died last year and realise you can no longer remember the sound of their voice.
- **Neutral.** You wait in line at the post office, glancing at the clock and watching the queue slowly advance one customer at a time.
- **Amusement.** A puppy sneezes so hard it tumbles backwards and stares at the floor in confusion before sneezing again.
- **Inspiration.** A dancer with one leg performs a flawless routine to a standing ovation, demonstrating that constraint can become craft.
- **Joy.** You receive an unexpected acceptance letter from a programme you applied to months ago and forgot about.
- **Tenderness.** A grandparent gently brushes a sleeping child’s hair away from their forehead, careful not to wake them.

Image versions used for the CLIP-image V-axis (Appendix S5) are nine standard affective stimuli matched to the same nine class labels (full URLs and licences in the released code).

Paraphrase generation. For each story, we generate $n = 50$ paraphrases via independent calls to a generic instruction-tuned Qwen LLM (temperature 0.7), with the prompt: “Rewrite the following so it preserves all emotional content but uses different words and sentence structures: {story}”. Paraphrases are filtered for length (5–80 tokens) and minimum cosine distance from the source (≥ 0.10 on Sentence-BERT embeddings).

Hidden-state extraction. Each paraphrase is tokenised and passed through the target LM. We take the mean-pooled hidden state at a late layer L (for the Qwen3-4B lead model, the final layer, $L = 36$ of 36 — the last transformer block’s output). Class centroids are $c_k = \frac{1}{50} \sum_{i=1}^{50} h_{k,i}^{(L)}$ for $k = 1, \dots, 9$.

V-axis as PCA-PC1. Stack the nine centroids into a $9 \times D$ matrix and run mean-centred PCA. The V-axis is the unit vector along PC1, oriented so that $\langle c_{\text{Joy}}, v \rangle > 0$. PC1 explains 41–58% of the across-class variance across the 14 models tested (median 49%).

Robustness. Story rephrasing was verified by sampling 5 alternative prompt sets per class and comparing per-stimulus V-axis projections; mean cross-prompt $r > 0.93$ across 14 models. Paraphrase count was verified at $n \in \{1, 5, 25, 50\}$; SST-2 zero-shot AUC 0.72 at $n = 1$, rising to the canonical 0.832 at $n = 50$ for Qwen3-4B.

Few-shot data-efficiency curve. The full n -shot sweep on Qwen3-4B (final layer) is in Table 3. Built entirely from LLM-generated paraphrases with *no* labelled SST-2 data, the V-axis reaches AUC = 0.832 at $n=50$ paraphrases per class — matching a supervised logistic regression trained on 50–100 labelled SST-2 examples (Table 3; $N=50$: 0.821, $N=100$: 0.894), while using zero task labels. At $n=1$ (single sentence per class, 9 total), AUC is already 0.722.

n paraphrases / class	1	3	5	9	15	25	50
SST-2 AUC (Qwen3-4B)	0.722	0.792	0.751	0.817	0.811	0.845	0.832

Table 3: V-axis few-shot data-efficiency curve (Qwen3-4B, final layer, 3 samplings per n). Standard deviations: $\pm 0.043, 0.012, 0.056, 0.019, 0.029, 0.006, 0.000$ from $n=1$ to $n=50$. The zero-label V-axis at $n=50$ (0.832) matches a supervised logistic regression trained on 50–100 labelled SST-2 examples ($N=50$: 0.821).

Supervised-LR equivalence anchor. A complementary anchor: on the same SST-2 features at the lead-model final layer, a supervised logistic regression reaches AUC 0.577 at $N=10$ labels, 0.821 at $N=50$, 0.894 at $N=100$, and 0.924 at $N=5,000$. The zero-label V-axis at $n=50$ paraphrases (0.832) matches supervised LR trained on 50–100 labelled examples — consistent with the data-efficiency claim (the V-axis needs *no* task labels at all).

S3 Per-LLM Cross-Architecture Convergence: Deep Dive

Sentiment benchmark table. Table 4 consolidates the zero-shot V-axis projection scores across five sentiment corpora and contrasts them against (i) a same-corpus 5k-example supervised logistic-regression upper bound on SST-2 and (ii) two zero-shot reference baselines (SBERT prototype-cosine and a 355M RoBERTa-large-MNLI entailment classifier). The takeaway is that nine LLM stories sit above the generic-encoder baseline and within 0.08 of an NLI-supervised model that is two orders of magnitude larger in supervised exposure.

LLM-as-judge baseline (steelman). A reviewer-natural question is whether prompting the same LLM directly as a sentiment classifier (“LLM-as-judge”) would beat the V-axis projection. We test this with Qwen3.5-2B as the judge on SST-2: a minimal binary prompt reaches AUC 0.951 at 165 s/1k examples, and a 5-shot in-context-learning variant reaches 0.957 at 253 s/1k. The V-axis projection trails on raw quality (0.832) but requires no per-example autoregressive generation—a single forward pass—and exposes a single re-usable direction in feature space (composable with downstream losses, ablatable, transferable across modalities; §7, §4). The competitive advantage is not raw zero-shot accuracy — it is interpretability and re-use, which the LLM-as-judge baseline does not provide.

Three regimes against behavioural valence. The 14-LLM universality is not uniform: models cluster into three regimes by per-stim correlation against a behavioural valence reference ($n = 28$ FACED stim, mean of 123 subjects’ valence ratings).

Benchmark	V-axis AUC	Supervised LR (5k)
SST-2 binary [Socher et al., 2013]	0.832	0.837
<i>SST-2 zero-shot reference baselines</i>		
SBERT prototype cosine [Reimers and Gurevych, 2019]	0.793	—
RoBERTa-large-MNLI entailment [Liu et al., 2019]	0.912	—

Table 4: Zero-shot V-axis sentiment results. The SST-2 V-axis AUC (0.832) is the verified Qwen3-4B lead-model number, within 0.005 of a 5k-sample supervised LR baseline (0.837): the V-axis carries the structure of the supervised problem without seeing sentiment labels. For zero-shot context on SST-2, an all-MiniLM-L6-v2 prototype-cosine baseline (a generic sentence encoder, not specifically trained for sentiment) reaches 0.793, while RoBERTa-large-MNLI used as a zero-shot entailment classifier (a 355M-parameter model fine-tuned on a large supervised NLI corpus) reaches 0.912. Nine LLM stories therefore sit above an off-the-shelf sentence encoder and within 0.08 of a 355M-param NLI-supervised model on the same zero-shot evaluation.

- **Top tier** (per-stim $r > 0.85$, $n = 7$): all six Qwen3 sizes (Qwen3-14B $r = 0.944$, Qwen3-4B 0.944, Qwen3-8B 0.941, Qwen3-32B 0.937, Qwen3-0.6B 0.933, Qwen3-1.7B 0.930) plus Mistral-7B (0.935). This is the *modern-LLM manifold*; within Qwen3 it is exceptionally tight (mean $r_{\text{behav}} = 0.938 \pm 0.005$, mean $r_{\text{eeg}} = 0.796 \pm 0.011$).
- **Middle tier** ($0.6 \leq r \leq 0.85$, $n = 4$): Llama-4-Scout-17B (0.826), Gemma-27B (0.807), Gemma-4-31B (0.715), Phi-2-2.7B (0.651).
- **Out-of-manifold** ($r < 0.4$, $n = 3$): Pythia-1.4B (0.276), TinyLlama-1.1B (0.128), BLOOM-560M (0.031).

The within-Qwen3 scaling is essentially flat: the six Qwen3 sizes span only $r_{\text{behav}} \in [0.930, 0.944]$ and $r_{\text{eeg}} \in [+0.781, +0.812]$ (s.d. ± 0.011) with no monotonic size trend—once a model has joined the modern-LLM manifold, convergence is *not* a size-monotonic phenomenon, and the behaviourally best-aligned models are mid-size Qwen3 (4B–14B).

Within-family vs cross-family decomposition. The “Universal=Qwen3-only” reading is the obvious reviewer pushback, so we decompose the 14-LLM agreement matrix explicitly. Within Qwen3 (six sizes, $C(6, 2) = 15$ pairs), V-axis cosine is $\bar{r} = +0.995 \pm 0.003$ — essentially perfect alignment. Across families (Qwen3 $\times$ non-Qwen, 48 pairs), $\bar{r} = +0.585 \pm 0.289$ — the wider spread is dominated by the three out-of-manifold models. The Qwen3 $\times$ Mistral subset (6 pairs) gives $\bar{r} = +0.939$ — functional V-axis agreement is near-perfect once the modern-LLM manifold is reached, regardless of training family.

V-shape across architectures. The same V-shape emerges across five LM families (Qwen, Pythia, BLOOM, Llama-4-Scout, Phi-2), and across bidirectional encoders BERT and RoBERTa, with the peak at layer $L_{\text{max}} - 1$ in each case. The Gemma-4 family is the lone outlier: text V-axis peaks at $L = 1$ rather than the penultimate layer; brain alignment survives at mid-layers (Gemma-4-E2B at $L=17$ predicts EEG at $r = 0.79$; Gemma-4-E4B at $L=21$ at $r = 0.64$).

S4 Concept Library and Compositionality

Compositional axis arithmetic. Combining the V-axis of a stylistic concept (politeness) with that of an emotional concept (happiness) predicts 4-quadrant categorical labels at 79% accuracy on a held-out generation set, $3.16\times$ above chance, with uniform per-cell precision in the $[0.68, 0.86]$ range. Subtractive orthogonalisation cleanly separates the two: happy-AUC $0.912 \rightarrow 0.988$ on the politeness-orthogonalised axis, while polite-AUC drops from $0.793 \rightarrow 0.560$. Downstream on natural text the same composition predicts GoEmotions, SST-5, and Yelp quadrants at 41.8%–46.7% (documented scope refinement).

Recipe generality across 20 concepts. For each of 20 unrelated dimensions (emotional, stylistic, abstract) we generated pole-contrastive stories and applied the identical extraction protocol. **20 of**

20 concepts reach binary AUC > 0.65 on held-out test stories, 17 of 20 reach ≥ 0.95 , 11 are perfect at 1.00 (Figure 3).




Figure 3: Binary pole-vs-pole AUC across 20 unrelated concepts (sorted descending), plus the out-of-scope toxicity case (Jigsaw, AUC 0.59).

Cross-concept transfer matrix. The 20×20 matrix has self-AUC mean 0.97 on the diagonal and off-diagonal mean 0.84 (gap +0.13): the recipe extracts *concept-specific* directions rather than a single global affect axis (Figure 4). Most general source axes: sarcasm (0.89), fear (0.88), envy and shame (0.87). Most specific: decisiveness (0.78), complexity (0.97). Most captured: urgency (0.97), fear (0.95), concreteness (0.94). Most concept-specific: specificity (0.65), sarcasm (0.67).


Toxicity: out-of-scope. Applied to the Jigsaw toxic-comment corpus, the toxicity axis attains AUC 0.59. Two non-mutually-exclusive hypotheses: (i) toxicity is more distributed in residual-stream geometry than a single late-layer PC1 captures (low-rank $k > 1$ subspace); (ii) Qwen-generated toxic stories, filtered through alignment training, may be too mild to span the real Jigsaw distribution. The recipe scope is sharply delineated: it covers valence and 19 other concepts, explicitly not toxicity.

S5 Multilingual and Vision Generality

Multilingual transfer. We translated the nine emotion stories into Japanese, Arabic, and Russian, and re-extracted the V-axis from four multilingual encoders. Direct application of an English-centric causal LM to non-English stories collapses to chance — a clean English-bias diagnostic. Switching to mT0-base [Muennighoff et al., 2023] (277M-param multilingual encoder-decoder) fully recovers the SST-2 result and *exceeds* the English baseline in every language tested (Figure 5).

Across the three languages the mT0-derived axes remain mutually aligned but not identical, consistent with concept-level rather than token-level encoding.

Cross-modal generality (CLIP vision). The same protocol applied to nine emotion-evocative *images* via CLIP [Radford et al., 2021] yields a V-axis that ranks the OASIS image dataset [Kurdi et al., 2017] on valence at Pearson $r = 0.869$ and on arousal at $r = 0.803$. Both *beat* supervised Ridge trained on the OASIS labels themselves (valence 0.836, arousal 0.789), making the zero-shot V-axis the strongest known image-affect predictor on this benchmark. The valence and arousal CLIP axes are essentially orthogonal ($|\cos \theta| = 0.013$).



Source: reports/r6_concept_transfer_matrix.json (Owen3-4B, penultimate layer, 12 train + 8 test stories per concept). Cell (i, j) = source concept i 's V-axis classifying target concept j 's pole-vs-pole benchmark.

Figure 4: Cross-concept transfer matrix (best-orientation AUC). Diagonal: self-classification (white-bordered cells). Off-diagonal: source-axis-applied-to-target-benchmark AUC. Sorted top-to-bottom by mean off-diagonal. Mean diagonal 0.97, mean off-diagonal 0.84, gap +0.13.

S6 Specificity Controls (Nonce, Random, Arousal)

Nonce-word ablation. Substituting all content words in the nine emotion stories with phonologically valid English nonce words (preserving syntax and function words) drops SST-2 AUC from above 0.83 to chance (≈ 0.52) — the second decimal. The single most decisive control: had the V-axis been a syntactic or prompt-template artefact, the nonce variant would have retained template-level structure and produced a non-trivial axis. The signal is encoded in semantic content.

Random-direction baseline. Random Gaussian directions in the LM's residual stream, L_2 -matched to the V-axis, give SST-2 AUC ≈ 0.52 and lexicon $|r| \approx 0.07$ – 0.09 . By contrast the V-axis bootstrap 95% CIs are tight and well separated from chance: SST-2 AUC $[0.844, 0.890]$, EmoBank V $|r| \in [0.485, 0.512]$, Hu & Liu $|r| \in [0.740, 0.776]$ (1,000-sample bootstrap; statistical methods in Appendix S15).

Arousal asymmetry: where the recipe fails in text. Table 5 side-by-sides the same nine-story PCA recipe instantiated on a valence axis (V-axis, this paper) versus an arousal axis (A-axis, identical extraction protocol with the nine-story prompts re-anchored on arousal poles). If the recipe were content-blind — merely a high-capacity dimensionality-reduction trick that produces a graded affective scale from any pole-vs-pole prompt set — the V and A columns should agree across the five textual rows. They do not: V transfers cleanly across SST-2, EmoBank, Warriner, NRC, and OASIS, while A collapses to near-chance on every textual benchmark. The OASIS-image and FACED-EEG




Figure 5: Cross-lingual V-axis recovery. Direct causal-Qwen transfer collapses (English-bias); mT0-base recovers and exceeds the English ceiling at roughly 5 $\times$ smaller parameter count (mT0-base 277M vs. the 1.5B causal reference) across three typologically distant languages.

rows are listed for contrast: vision and brain do encode an arousal axis the same recipe recovers, so the asymmetry is text-specific rather than a recipe limitation.

Benchmark	V-axis (valence)	A-axis (arousal)
SST-2 binary AUC	0.832	—
EmoBank V $ r $ / EmoBank A $ r $	0.49	0.01–0.06
Warriner V $ r $ / Warriner A $ r $	0.703	≤ 0.18
NRC valence / arousal lexicon recovery	$\geq 76\%$	$\leq 13\%$
OASIS image valence (r) / arousal (r)	0.869	0.803
FACED EEG cohort r	0.870 (9-stim)	$r \in [0.18, 0.41]$

Table 5: Valence transfers across modalities; arousal does not in text but does in vision. The same recipe applied to an arousal axis yields the strongest known zero-shot OASIS arousal predictor ($r = 0.803$, beating supervised ridge 0.789) but collapses on all text benchmarks.

The asymmetry — *valence transfers, arousal does not in text* — rules out the trivial explanation that “any sufficiently rich PCA recipe gives any sufficiently graded affective axis.” Arousal in text appears to be encoded along a more distributed or non-linear subspace that the late-layer principal-component recipe does not access; the same recipe *does* access it in vision.

S7 Brain-Side Deep Dive

Per-stimulus reliability. Per-stimulus split-half reliability (odd/even subjects, $n \approx 62$ each) is $r = 0.988$, indicating a clean stimulus-level signal. Trial-level correlations ($n = 3,444$ subject–stim pairs) range from $r = 0.17$ – 0.21 ($p < 10^{-23}$); class-level ranking ($n = 9$ centroids) hits $r = +0.886$ vs. behavioural valence. The cohort ridge reaches $r = 0.87$ with the CLIP-text V-axis and $r = 0.80$ with the nine-story Qwen3-4B V-axis; both are reliability-bounded by the $r = 0.988$ split-half.

Per-LLM brain forest. Per-LLM brain-prediction quality is shown in Figure 6: the top-tier LLMs all clear their per-LLM random-direction nulls, while the out-tier sits at or below the median.

Time-locked dynamics. Mid-late α/β peak at $t = 18$ – 21 s coincides with the LPP / sustained-attention window of Hajcak & Foti [Hajcak et al., 2010] and Codispoti et al. [Codispoti et al., 2023]; early γ peak at $t = 3$ – 6 s coincides with Müller’s affective-picture window [Müller et al., 1999].

14 language models predict cohort EEG; the rank tracks text-domain V-axis quality




Figure 6: Per-LLM brain-prediction quality. Top-tier 7 LLMs all clear their per-LLM random-direction baseline; out-tier 3 LLMs sit at the median or below. Restricted to the top-tier 7, the log-hidden-dim $\rightarrow$ r scaling is $r = +0.44, p = 0.33, n = 7$; within-Qwen3 $r = +0.57, p = 0.24, n = 6$. Family-membership $\gg$ scale.

Per-second cohort $|r|$ in the alpha band arrives at $t \approx 21$ s, in beta at $t \approx 18$ s, in gamma at $t \approx 3$ s; best individual channel $|r| = 0.68$ at $t = 21$ s.

Per-subject peak heterogeneity. Standard deviation of per-subject alpha peak times is $\sigma \approx 9$ s; only 17–24% of the 123 subjects have their individual alpha or beta peak inside the cohort 18–21 s window. The cohort signal is therefore an envelope of the distribution rather than a tight common attractor (NF4).

Gemma-4 boundary check. Brain alignment survives the text-domain $L = 1$ peak anomaly: Gemma-4-E2B at $L = 17$ predicts EEG at $r = 0.79$ ($p \approx 10^{-7}$), Gemma-4-E4B at $L = 21$ at $r = 0.64$ ($p \approx 2 \times 10^{-4}$). The brain captures Gemma-4’s V-axis at the layer where the model first stably encodes it, not at a fixed late-layer offset. This cross-model alignment refines to a property of the late-but-not-necessarily-penultimate residual stream geometry.

S8 SEED-V Replication Full Numerics

We re-test all five FACED claims with a SEED-V-derived V-axis. To remove FACED leakage, we re-build the V-axis from scratch using the same Section 6 protocol applied to the 5 SEED-V emotion classes (Qwen3-4B at the final layer $L=36$, 50 stories per class, PCA on the 5 centred centroids).

Step 1: SEED-V cohort EEG–LLM alignment. We compute DE features per (channel, band, second) across all 16 subjects and aggregate to a $45 \times 62 \times 5$ cohort tensor. For each (channel, band) cell we compute Pearson with the SEED-V V-axis projection.

- **Best cohort cell:** PZ / θ , $r = +0.5920$.
- **Region-mean $|r|$:** occipital 0.300, parietal 0.315, central 0.302, frontal 0.234.
- **Posterior dominance:** occipital – frontal = +0.066.
- **Random-direction null** (200 directions matched in L_2): empirical $|r|_{\max} = 0.5920$ vs. null 95th percentile = 0.478, $p < 0.005$.

Step 2: SEED-V brain topography. The Davidson FAA probe (positive vs. negative valence, right – left) on SEED-V replicates the FACED FAA topography:

Electrode pair (right – left, alpha)	(pos – neg) DE Δ
F4 – F3	–0.0018
F8 – F7	+0.0050
FP2 – FP1	–0.0094
AF4 – AF3	–0.0039

Table 6: Davidson FAA on SEED-V. Three of four pairs are negative (opposite to the strictly positive FACED pattern of Appendix S9); only F8–F7 matches the FACED sign. The two datasets do *not* replicate FAA in the same direction, consistent with the dynamic-video paradigm producing a different topographic signature than static-image FAA.

Step 3: SEED-V within-architecture V-axis encoding. Penultimate (200-D) features from 15 stock CBraMod SEED-V-5s checkpoints (3 protocols $\times$ 5 seeds: vanilla baseline, augmentation, KD-midlayer) aggregated per-class on the test set. PC/best-PC scores correlated with the SEED-V V-axis class projection give $r(\text{BACC}, \text{trial-best-PC } r) = +0.628$ ($p = 0.012$, $n = 15$; this is the direction used for the causal ablation below) and $r(\text{BACC}, \text{class-best-PC } r) = +0.520$ ($p = 0.047$). The SEED-V BACC dynamic range is much narrower than FACED, so the raw class-PC1 correlation is not significant; the best-PC estimators reach significance. These are 15 checkpoints of a single backbone (CBraMod), so we read this as within-architecture V-axis encoding, not architecture-independent convergence.

Step 4: SEED-V V-axis as supervision. Adding a V-axis topo-MSE auxiliary loss ($\lambda_{\text{vax}}=0.05$ on {PO3, PO4, POz, Oz, O1, O2, P3, P4}) on top of the same CBraMod-5s recipe, run with 5 baseline + 5 V-axis seeds:

Variant	BACC (mean $\pm$ s.e. across 5 seeds)	vs. baseline Δ
Baseline (no V-axis loss)	0.4307 $\pm$ 0.0009	—
+ V-axis topo loss	0.4325 $\pm$ 0.0017	+0.0018

Table 7: SEED-V V-axis-as-supervision (within seed noise of zero). Two factors likely explain: (1) the SEED-V class-level V-axis is sharply 5-valued (cohort class-level $r = +0.958$), so adding a class-level V-axis target on top of one-hot supervision is largely redundant; (2) the SEED-V BACC dynamic range (0.43–0.45) is much narrower than FACED (0.55–0.58), so a small effect is hard to detect with five seeds.

Step 5: SEED-V directional ablation (Arditi-style). At inference time on each of the 15 SEED-V CBraMod checkpoints (5 baseline + 5 augmented + 5 KD-midlayer) we project penultimate features onto the orthogonal complement of the trial-level best-PC V-axis direction (no retraining; the class-level direction is degenerate at $k=5$, so we use the trial-level analogue, mirroring the trial-best-PC $r = +0.628$ estimator from Step 3). Every checkpoint shows a negative $\Delta\text{BACC} \in [-0.0291, -0.0093]$ (population mean -0.0173 ± 0.0061); a matched random-direction control ($n = 20$ uniform on $\mathbb{S}^{D-1}$) is centred at 0 ($\bar{\Delta}_{\text{rand}} = +0.00013 \pm 0.0008$). Per-checkpoint z -scores range $[+9.65, +44.92]$ (mean $+23.6$, all $p < 0.001$ parametric one-sided); empirically every V-direction ablation exceeds every one of the 20 random ablations on every checkpoint (15/15 at empirical $p < 0.05$). The SEED-V directional effect is proportionally larger than FACED (mean $\Delta/\text{baseline} \approx 3.9\%$ vs. 2.4% on FACED) and the mean z is $3\times$ FACED’s. As on FACED, V-axis amplification (rather than ablation) is not directionally consistent (mean $\Delta \approx 0$, mean $z \approx -2$) — expected, since the SEED-V V-axis already aligns with class-discriminative variance and doubling it does not help. Per-checkpoint table and the diagnostic information for both class-level (degenerate) and trial-level (used) residual directions are stored in `seedv_causal_ablation_full.json`. The ablation direction is robust to the LLM used to build the V-axis: rebuilt with the Qwen3-4B lead model, 14 of the 15 checkpoints select an *identical* intervention direction ($|\cos| = 1.00$ against the reported build; the class projections of the two builds agree at $r = 0.99$), so the ablation result is unchanged under the lead-model V-axis.

Third-dataset consistency check: DEAP. As a further test of the saturation regularity on an independent dataset (DEAP [Koelstra et al., 2011]; 32 subjects, 32 channels, 128 Hz), we trained

a binary-valence classifier (valence rating ≥ 5 , sourced from the raw DEAP labels) from the same pretrained EMOD backbone, splitting by (*subject, trial*) unit so that all windows of a trial stay on one side. The converged baseline reaches $\text{BACC} = 0.544$ — below the $[0.62, 0.66]$ transition band identified on FACED — so the regularity predicts the V-axis auxiliary loss should have *no significant effect* here, rather than a decrement. This is what we observe: across three seeds and three loss weights ($\lambda \in \{0.1, 0.5, 1.0\}$), the pooled change is $\Delta\text{BACC} = -0.019 \pm 0.044$ ($n=9$, one-sample $t=-1.28$, $p=0.24$), within noise of zero (two of three seeds negative, one positive). DEAP thus populates the weak-baseline regime of the transition curve: consistent with, but not an independent confirmation of, the destructive regime seen once a FACED or SEED-V classifier passes the transition. Per-run numbers are in `deap_sat/SUMMARY.json`.

S9 Brain Topography Deep Dive




Figure 7: V-axis encoding peaks in posterior visual cortex across all five frequency bands. (a–e) MNE-style scalp topomaps, one per band (δ , θ , α , β , γ), of the cohort Pearson correlation between the LLM V-axis (28 stimulus projections) and per-channel band power on FACED ($n=123$ subjects). Diverging colourmap: positive (red) vs. negative (blue) correlation; symmetric limits across bands. Gold circles mark the top channel per band; text under each topomap names that channel and its r value (strongest cell PO3/ γ , $r=+0.48$). The colourbar (centred) gives Pearson r ; the small head diagram (right) shows the 32-channel 10–20 montage used in all five panels (A1/A2 mastoids excluded from plotting). (f) Region-mean $|r|$ aggregated over all five bands: occipital (0.21) > parietal (0.18) > central (0.18) > frontal (0.16). The cohort signal is posterior-dominant for video-evoked emotion—a paradigm-level finding for dynamic stimuli rather than a refutation of the static frontal-alpha asymmetry literature, which we replicate qualitatively in direction at smaller magnitude.

Davidson FAA full pairs. Figure 8 compares all three Davidson asymmetry pairs to the posterior occipital region-mean reference.

9-stim contrast and Simpson’s paradox. Table 8 decomposes the cohort-level brain–LM correlation by stimulus subset, isolating which of the 28 FACED stimuli carry the V-axis signal. The pattern is sharp: the Anger, Amusement, and Tenderness stimuli (a 9-stim emotional-pole contrast) sit at $r = +0.870$, while removing them collapses the residual 19 stimuli to $r = -0.015$. The full-cohort $r = +0.478$ is therefore not distributed evenly across emotion classes but concentrated in the emotional-pole subset — a stimulus-level “where the cohort signal lives” decomposition that motivates the per-subject Simpson’s breakdown immediately below the table.

- Cohort r at fixed top-8 channels: $+0.021$
- Mean per-subject r at fixed top-8 channels: -0.062
- Per-subject best-channel oracle (1 cell): $|\overline{r}| = 0.551$

Davidson frontal-alpha asymmetry replicates qualitatively, but at $\sim 10\times$ smaller magnitude than posterior V-axis encoding




Figure 8: Frontal-alpha asymmetry replicates in direction with smaller magnitude than posterior V-axis encoding for these video stimuli. (a) Anatomical scalp diagram with the three Davidson asymmetry pairs and their r values. (b) Per-pair alpha and delta asymmetry bars. (c) Direct comparison to the occipital region-mean $|r| = 0.21$.

Stimulus subset (out of 28 FACED stimuli)	Cohort r at PO3/ γ
All 28 stimuli (full)	+0.478
Anger only ($n=3$)	not estimable
Anger + Amusement + Tenderness ($n=9$)	+0.870
Excluding Anger ($n=25$)	+0.327
Excluding Anger + Amusement + Tenderness ($n=19$)	-0.015

Table 8: The cohort V-axis effect is essentially an *anger-versus-warm-positive* contrast across nine stimuli.

- **Per-subject best (channel, band) oracle:** $\overline{|r|} = 0.616$

The all-band oracle ($\overline{|r|} = 0.616$) exceeds the cohort-fixed top-8 (+0.021) by +0.60 in absolute $|r|$, indicating considerable per-subject signal that the cohort summary suppresses. The Simpson’s-paradox dynamic that Sani et al. [Vahidi et al., 2024] flag for cross-subject EEG-emotion modelling more broadly (Figure 9).

Time-resolved peaks. Figure 10 traces per-second cohort $|r|$ across all five frequency bands.

- **Alpha** (8–13 Hz): peak at $t = 21$ s, cohort $r = 0.40$; best individual stimulus $r = 0.68$.
- **Beta** (13–30 Hz): peak at $t = 18$ s, cohort $r = 0.41$; best stim $r = 0.62$.
- **Gamma** (30–47 Hz): peak at $t = 3$ –6 s, cohort $r = 0.21$; best stim $r = 0.61$.
- **Delta/Theta** (0.5–8 Hz): peak at $t = 12$ s, cohort $r \approx 0.3$.

Functional connectivity. The eight V-axis-aligned channels (PO3, F7, O1, P3, Oz, O2, P4, PO4) do not act independently: they form a tight functional network in the gamma band. Pairwise γ -DE correlation has mean off-diagonal $r = 0.67$ across the 8 channels — substantially above the whole-band pairwise average. Structure is dominantly a posterior-occipital-parietal cluster (pairwise mean $r \approx 0.73$) with F7 as a single *frontal hub* (mean correlation to the seven posterior nodes $r \approx 0.50$; Figure 11).

Mutual information vs. linear correlation. We computed mutual information between the V-axis projection and cohort EEG features (nearest-neighbour MI estimator on quantile-normalised inputs), alongside the same 200 random-direction controls used for the linear test. Observed MI 0.112; null mean 0.036; $p_{\text{MI}} = 0.115$ (n.s.) vs. linear $p_r = 0.020$. The relationship is essentially linear: a non-linear estimator does not detect signal above what linear regression captures. We use ridge regressions throughout.




Figure 9: The cohort V-axis is a 9-stimulus contrast and a between-subject phenomenon. Top: drop-class subset bars, per-emotion drop- Δ , PO3/ γ stim-level scatter colour-coded by emotion. Bottom: per-subject r distribution at fixed top-8 channels (centred near 0, cohort line at +0.00); per-subject best-channel oracle distribution (centred at $|r| \approx 0.55$); summary callout.




Figure 10: Time-resolved V-axis encoding across the 30-second clip. **(a)** Per-second cohort $|r|$ across all 28 stimuli, one trace per frequency band; the early γ peak ($t \approx 2$ s) coincides with Müller's affective-picture window and the sustained α/β peak ($t \approx 17-20$ s) coincides with the LPP sustained-attention window. **(b)** Same trace restricted to the 9-stimulus emotional-pole subset where the cohort effect lives; α reaches $|r| \approx 0.78$. **(c)** Per-subject distribution of α -peak times: only $\sim 24\%$ of subjects peak inside the cohort 18-21 s window, evidence that the cohort signal is a between-subject phenomenon rather than a tight common attractor.

Theta-gamma coupling does not mediate the V-axis. Tort modulation index per channel at the theta-phase / gamma-amplitude pairing, correlated with V-axis encoding strength $|r|$ across 32 channels: $\rho(\text{PAC}, Vr) = +0.082$, $p = 0.667$. Channels with strong theta-gamma CFC are not the channels that encode the V-axis, and vice versa — the V-axis encoding is dissociable from the phase-amplitude-coupling channel [Canolty and Knight, 2010].

S10 Saturation: Full Intervention Table and Forest

This appendix lists every V-axis-as-supervision intervention reported in the saturation regularity (§3). Table 9 gives per-recipe ΔBACC against the standard EMOD $d=6$ baseline, the paired-seed p -value, and a verdict tier; the two SOTA-recipe replications below the line (against the strong $d=6 + \text{KD} + \text{aug}$ recipe at $\text{BACC} \approx 0.66$) confirm the cliff persists in the regime where saturation is

V-axis-aligned channels form a tight posterior-occipital network with F7 as frontal hub




Figure 11: V-axis-aligned channels form a coordinated network with F7 as frontal hub. Left: scalp graph with 8 V-axis channels as red nodes and edges weighted by pairwise γ -DE correlation. Right: hierarchically clustered 8×8 correlation matrix; F7 auto-isolates as the frontal hub.

operative. Figure 12 renders these interventions — together with additional null variants we tested (weight/bias initialisation, soft-KD temperatures, curriculum, prior smoothing) — as a forest plot ranked by effect size, so the asymmetry is visible at a glance: every significant entry sits to the left of zero, and even the largest n.s. positive (EMODSTYLE-stim, $+0.0066$, $p=0.22$) is within seed noise. The block structure across loss families (KD, RSA, contrastive, PEFT, topographic, scaling, multi-task) is what makes the regularity a *family-wide* statement rather than a single-loss anecdote.

S11 Saturation Extras: Anger Paradox, Mechanism Check, Path B

Saturation transition table. Table 10 pairs the same V-axis intervention (Topo or EMODSTYLE auxiliary loss) against five base recipes spanning the BACC range from 0.572 (CBraMod baseline) to 0.6581 (EMOD $d=6$ SOTA pre-ensemble). Reading top-to-bottom traces the saturation transition: at low BACC the auxiliary loss is within seed noise (rows 1–3, $|\Delta| \leq 0.007$, all n.s.), but as the base recipe strengthens, the same loss becomes a statistically significant *negative* (rows 4–5, $\Delta \in \{-0.029, -0.026\}$, both $p < 0.01$). No row crosses the boundary in the helpful direction. This is the data behind the “no positive setpoint” wording in §7 and the basis for treating BACC ≈ 0.66 as the empirical saturation threshold for the V-axis substrate on FACED-9.

Saturation cliff figure. Figure 13 plots Δ BACC against base-recipe BACC across all 25 interventions, showing the unidirectional sign-flip in the $[0.62, 0.66]$ band.

Why anger-weighting hurts despite the analytical ceiling. Section 7 showed that the cohort EEG signal is carried by a 9-stimulus emotional-pole contrast (Anger + Amusement + Tenderness; cohort ridge $r = 0.870$). Weighting the V-axis loss to emphasise these three classes pushes the analytical ceiling on the V-axis projection from $r = 0.478$ (uniform) to $r = 0.714$ (anger-weighted). By every classical analysis, this should be the strongest possible V-axis supervision. It is the worst single intervention we tested: $\Delta = -0.054$, $p = 0.0003$. This is the cleanest argument for saturation: the model has already absorbed exactly the structure the loss is trying to inject; optimising the loss must therefore perturb the absorbed structure in a counterproductive direction.

Direct mechanism check. For each of 8 V-axis-supervised variants (Topo at $\lambda=0.05$, Procrustes at $\lambda=0.05$, EMODSTYLE-stim at $\lambda=0.5$, full SOTA recipe + Topo/EMODSTYLE at $e=100/e=150$), we measure the change in two quantities relative to the matched-recipe baseline: Δr_{PC1} (class-mean V-axis encoding) and Δr_{resid} (within-class V-axis residual encoding). V-axis training raises class-PC1

#	Family / variant	Δ	p	Verdict
1	Frontal-mask $\lambda=0.5$	-0.052	0.0015	sig. negative
2	Frontal-mask $\lambda=0.1$	-0.009	0.21	n.s.
3	FAA $\lambda=0.5$	-0.044	0.006	sig. negative
4	FAA $\lambda=0.1$	-0.018	0.063	borderline
5	Anger-weighted $\lambda=0.5$	-0.054	0.0003	sig. negative
6	Occipital $\lambda=0.1$	-0.022	0.007	sig. negative
7	Topo-optimal $\lambda=0.1$	-0.013	0.039	sig. negative
8	Topo-optimal $\lambda=0.05$	+0.0021	0.50	n.s. (NULL)
9	Procrustes $\lambda=0.05$	+0.0012	0.89	n.s. (NULL)
10	RSA $\lambda=5.0$	-0.093	$< 10^{-4}$	sig. negative
11	RSA $\lambda=1.0$	-0.057	$< 10^{-3}$	sig. negative
12	Distance-CE $\tau=5.0$	-0.397	$< 10^{-4}$	catastrophic
13	Multi-V $\lambda=0.5$	-0.041	$< 10^{-3}$	sig. negative
14	EEG-AUX MSE $\lambda=0.5$	-0.043	$< 10^{-3}$	sig. negative
15	EEG-AUX MSE $\lambda=0.1$	-0.016	0.04	sig. negative
16	EMODSTYLE class $\lambda=1.0$	-0.010	0.30	n.s.
17	EMODSTYLE stim $\lambda=0.5$	+0.0066	0.22	n.s. (sweet spot)
18	PEFT fullhead $\lambda=0.1$	-0.009	0.069	borderline
19	PEFT LoRA $\lambda=0.1$	-0.010	0.035	n.s. (Holm 0.11)
20	PEFT IA3 $\lambda=0.1$	-0.008	0.14	n.s.
21	Pretrain-FT (frozen)	-0.401	$< 10^{-4}$	catastrophic
22	Pretrain-FT (unfrozen)	-0.190	$< 10^{-4}$	sig. negative
23	XEEG (FACED $\rightarrow$ SEED-V)	+0.0015	0.97	n.s. (zero)
24	SCALING (full data) $\lambda=0.5$	-0.045	< 0.01	sig. negative
25	Uncertainty multi-task	-0.067	$< 10^{-3}$	sig. negative
SOTA recipe + Topo $\lambda=0.1$		-0.029	0.004	sig. negative (cliff)
SOTA recipe + EMODSTYLE		-0.026	0.004	sig. negative (cliff)

Table 9: Full V-axis-as-supervision intervention results. 15 / 25 families produce a statistically significant negative (13 sig. negative + 2 catastrophic); the remaining 10 are not significantly different from zero. None reach $p < 0.05$ in the positive direction at the strong baseline. Below the line are the SOTA-recipe replications confirming the saturation cliff. Forest plot in Figure 12.

Recipe	n seeds	Vanilla / + V-axis	Δ
CBraMod baseline + Topo $\lambda=0.05$	5	0.572/0.567	-0.005 (n.s.)
CBraMod baseline + EMODSTYLE $\lambda=0.5$	4	0.572/0.577	+0.005 (n.s.)
EMOD $d=3$ vanilla + Topo $\lambda=0.05$	5	0.6194/0.6260	+0.007 (n.s.)
EMOD $d=6$ + LS + KD + aug + Topo (full)	5	0.6581/0.6291	-0.029 ($p < 0.01$)
EMOD $d=6$ + LS + KD + aug + EMODSTYLE	5	0.6581/0.6317	-0.026 ($p < 0.01$)

Table 10: V-axis supervision is within seed noise at weak baselines and statistically significantly *harms* the strong SOTA recipe. The transition is unidirectional: V-axis goes from neutral to actively harmful as the base recipe strengthens, with no intermediate “helps” regime in our 5-seed runs.

$|r|$ by +0.01 to +0.36 across the 8 variants, but leaves the within-class residual encoding essentially unchanged. Consistent with the accuracy behaviour (training accuracy at ceiling, test accuracy down; A6), the decrement at strong recipes comes from reshaping the already-saturated class-PC1 basin rather than from any beneficial transfer of residual structure. (We do *not* claim the auxiliary gradient fails to reach the residual subspace: a direct gradient decomposition finds a residual-subspace component comparable to a random direction, so the mechanism is one of *no useful signal added*, not of a geometrically unreachable subspace.)

Path B: ensembling-in V-axis checkpoints. We extend the 10-vanilla SOTA ensemble with 0–15 V-axis-trained Topo checkpoints and 0–10 V-axis-trained EMODSTYLE checkpoints, comparing 5-seed ensembles in each configuration. Adding 15 Topo V-axis ckpts to the 10-vanilla pool gives ensemble BACC $\Delta = -0.0145$ ($p < 10^{-3}$); adding all 25 V-axis ckpts gives $\Delta = -0.0193$ ($p < 10^{-3}$). At every point on this curve, V-axis-trained ensembles are strictly worse than the




Figure 12: Forest plot of all V-axis-as-supervision interventions (Table 9) ranked by Δ vs. baseline, with significance tier coloured. None of the 25 reach $p < 0.05$ in the positive direction at the strong baseline.

matched vanilla ensemble. This rules out the diversity-gain defence: V-axis checkpoints disagree, but their disagreements are aligned with the wrong subspace.

S12 Convergence and Ensemble Theory: Extras

The 36-checkpoint cross-architecture correlation is pseudo-replicated. Pooling all 36 checkpoints from two backbones (CBraMod, EMOD) and six recipe variants, balanced accuracy correlates with class-PC1 V-axis encoding at $r = +0.885$ ($p_{\text{param}} = 7.8 \times 10^{-13}$). This parametric p is an artifact of non-independence. The checkpoints form ~ 2 – 3 tight architecture clusters—CBraMod (~ 0.57 BACC, mean class-PC1 $|r| = 0.21$), EMOD-vanilla d6 (~ 0.65 , mean 0.67), EMOD d6 e150 (~ 0.66 , mean 0.69)—so the correlation is driven by the between-architecture gap, not a within-architecture trend. *Within* each architecture it is not significant (CBraMod $r = 0.83$, $p = 0.08$; EMOD d6 e100 $r = 0.04$, $p = 0.90$; EMOD d6 e150 $r = -0.37$, $p = 0.37$), and the effective number of independent checkpoints is $n_{\text{eff}} \approx 1.2$.

Random-direction null detail. We sample 1000 random Gaussian directions $w \in \mathbb{R}^{28}$ each matched in L_2 norm to v_{CLIP} , and recompute the 36-checkpoint correlation $\rho_w = \text{Pearson}(\text{BACC}, |\text{corr}(\text{PC}_1(H(m)), w)|)$. The distribution of ρ_w has mean -0.03 and standard deviation 0.62 (range $[-0.95, +0.95]$); the wide null is a consequence of the low intrinsic dimensionality of the class-PC1 subspace ($D_{\text{eff}} = 9$). The empirical V-axis correlation $+0.885$ sits at only the 93.5th percentile of this null (one-sided $p = 0.066$, two-sided $p = 0.099$)—it does not survive. Taken together with the within-architecture nulls above, we therefore do *not* claim architecture-independent convergence: the only convergence claim we retain is the within-architecture class-PC1 basin saturation below (and the load-bearing within-class residual, which *is* statistically robust).




Figure 13: Δ BACC from V-axis supervision plotted against base-recipe BACC across ~ 25 interventions. The sign of the effect transitions in the interval $[0.62, 0.66]$.

Class-PC1 basin saturation. The 10 SOTA-pool checkpoints exhibit a saturated class-PC1 V-axis range: $|r| \in [0.60, 0.77]$ across all 10 seeds (mean 0.69, std 0.05). Single-seed variation in class-mean V-axis structure is small. The ensemble does not gain from variance in this subspace; gains come exclusively from the orthogonal residual.

Directional ablation figure. Figure 14 shows the per-checkpoint causal effect: ablating v_{resid} drops BACC by $\bar{\Delta} = -0.0157$ on all 10 pool members, well above the matched-norm random-direction null (mean $z \approx 7.7$, all $p < 0.001$).




Figure 14: Directional ablation of v_{resid} on all 10 SOTA-pool checkpoints. *Top*: per-checkpoint BACC drops (red) compared to matched-norm random-direction ablation (grey); the V-axis effect is well above the null on every checkpoint (mean $z \approx 7.7$, all $p < 0.001$). *Bottom*: per-checkpoint ablation Δ BACC vs. within-class residual encoding $|r|$; the population-mean drop ($\bar{\Delta} = -0.0157$) confirms the residual-encoding mechanism across the pool.

S13 Ensemble Generality, Mega-Pool, and Best-Single

Recipe cascade and two-tier scatter. Figure 15 traces the seven recipe steps from CBraMod’s 0.572 to our 0.6948 ensemble. Figure 16 shows the per-checkpoint within-class residual–ensemble-contribution correlation that motivates the ensembling strategy.




Figure 15: Recipe ablation cascade from the EMOD baseline at 0.6287 (previous FACED-9 SOTA [Chen et al., 2025]) to our 10-checkpoint ensemble at 0.6948 (+10.5% relative). Each bar is one recipe component or ensembling step. Older baselines (CBraMod 0.572, EmotionKD 0.628) shown for reference.




Figure 16: Two-tier ensemble theory. Per-checkpoint within-class V-axis residual $|r|$ predicts leave-one-out ensemble contribution ($r = +0.74$, $p = 0.014$, $n = 10$). Top-7 vs. bottom-7 split highlighted.

Generality of the ensemble mechanism. Figure 17 extends the two-tier picture across four benchmarks (FACED, SEED-V, CIFAR-10, MNIST).

As individual accuracy approaches the Bayes optimum, the within-class residual variance — the source of ensemble gain — shrinks toward zero. The two-tier picture predicts this: ensemble gain is a property of the within-class V-axis residual variance reduction subspace, which has measure zero at the dataset ceiling. SEED-V at 0.37 has the largest residual-variance ceiling; we observe +0.034 gain. MNIST at 0.985 has essentially none; we observe +0.001 gain.

Mega-ensemble null. Growing the ensemble beyond the ten $d=6$ checkpoints by adding depth variants ($d \in \{8, 10\}$) did not exceed the 10-checkpoint pool in exploratory runs: consistent with the saturated-basin picture, once the shared basin is covered, extra architectural variants contribute




Figure 17: Two-tier ensemble gain across four benchmarks. ΔBACC is largest where single-model accuracy is furthest from the dataset ceiling, and shrinks to near-zero on MNIST where the saturated regime leaves no within-class residual variance to cancel. (a) Scatter of Δ vs. headroom $1 - \text{Acc}$, with linear fit slope ~ 0.05 . (b) Paired single-seed-mean vs. ensemble bars per benchmark (FACED shows the 10-checkpoint SOTA ensemble), sorted by headroom.

correlated rather than complementary predictions. We therefore report the 10-checkpoint $d=6$ pool as the operative optimum (exploratory wider-pool runs not tabulated).

Two mechanisms explain the null. First, deeper variants ($d=8, d=10$) stay in the same class-PC1 V-axis basin as $d=6$ and their within-class residuals overlap more than $d=6$ seeds do. Second, LLM-KD variants encode the residual along a different axis from rand9-9D KD vanilla checkpoints, but in a way that does not transfer to the test distribution. Cross-architecture mixing ($d \in \{4, 6, 8, 10\}$) gives 0.6791, essentially tied with within-arch x-seed at 0.6798. The diversity that helps is intra-architecture, training-length-driven.

Prediction-agreement structure. Single-seed accuracy at $e=100$ (0.6581) and $e=150$ (0.6581) is identical. The pairwise prediction correlation within- $e=100$ is 0.694, within- $e=150$ 0.679, cross-group 0.702 (higher = more agreement). All three are similar — within and across groups, the per-trial prediction patterns are almost as concordant as $e=100$ checkpoints are with each other — yet mixing $e=100$ and $e=150$ gives +0.015 over the $e=100$ ensemble. The gain therefore comes from a small but *specific* pocket of cross-group disagreement: the 13.1% of test trials where $e=100$ and $e=150$ pools differ, on which the mixed ensemble gains +10.3 percentage points. The diversity is targeted (which trials), not population-wide.

Best single checkpoint without ensembling. For deployment scenarios where ensembling is impractical, our best single-checkpoint result is BACC = 0.6755 on the $d=6$ $e=150$ recipe, seed 789, val-selected. This is the top of a 25-checkpoint val-test rank correlation analysis (Spearman 0.825, $n = 25$) and exceeds CBraMod by +0.103 and EMOD by +0.047 on a strict apples-to-apples test split.

S14 EEG Model Training Details

Architecture. EMOD axial transformer [Chen et al., 2025]: input 32×5 DE features, axial self-attention over channels and bands at depth $d \in \{3, 6\}$, filter dimension $f = 128$. The full SOTA recipe is $d = 6, f = 128$.

Optimisation. AdamW with $\text{lr} = 10^{-3}$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, weight decay 10^{-2} . Cosine schedule with 5-epoch linear warmup. Batch size 128. Training length $e \in \{100, 150\}$ epochs. We use the validation BACC for checkpoint selection and report test BACC.

Augmentation. Per-trial Gaussian noise ($\sigma = 0.05 \cdot \text{std of feature distribution}$), random channel dropout ($p = 0.15$), and random temporal masking (up to 5% of seconds). All augmentations are applied with $p = 0.6$ during training only.

Knowledge distillation. Class-mean KD with a rand9 9-D orthonormal teacher (no LLM content); $\lambda_{\text{KD}} = 0.1$, teacher temperature $\tau = 0.1$. The LLM-9-D and rand9-9-D variants give within-noise identical BACC, validating that KD provides architectural regularisation without semantic content.

Ensemble. 10 checkpoints split as 5 seeds at $e = 100$ + 5 seeds at $e = 150$. Seeds: $\{42, 123, 456, 789, 2025\}$ each. Ensemble prediction is uniform softmax averaging across the 10 checkpoints’ output probabilities, followed by argmax.

S15 Statistical Methods

Bootstrap CIs. Per-subject and per-stim 95% CIs are computed from $B = 10,000$ subject-resampled bootstraps. For cohort correlations, we report the bootstrap median, 2.5% / 97.5% percentile interval, and a Fisher- z p -value for $r = 0$.

Random-direction null. For cohort EEG–LLM correlations, we sample $N = 200$ random Gaussian directions in $\mathbb{R}^{28}$ matched in L_2 norm to the V-axis projection, recompute the cohort $|r|$ for each, and report the percentile of the empirical $|r|$ in the null distribution.

Cross-architecture null. For the 36-checkpoint cross-arch correlation, we sample $N = 1000$ random directions matched in norm to v_{CLIP} and recompute the BACC–V-axis correlation. The null is distributed broadly ($\sigma = 0.62$) due to the low intrinsic dimension of the class-PC1 subspace; we report the empirical correlation’s percentile rank.

Per-checkpoint LOO contribution. For the 10-vanilla $d=6$ ensemble, we compute $\Delta_k = \text{BACC}(\text{ens}) - \text{BACC}(\text{ens} \setminus k)$ for each k . Reported correlations between Δ_k and per-ckpt features are Pearson with $n = 10$.

V-axis intervention comparisons. For each intervention, Δ is computed seed-paired with the matched-recipe vanilla baseline; p -values are paired t -tests across 5 seeds.

S16 Discussion Extras and Future Work

Broader impacts. Improved EEG emotion classifiers have clear positive applications in mental-health monitoring, affective brain–computer interfaces, and clinical phenotyping of affect disorders. The same techniques pose risks of affective inference without consent in surveillance, hiring, workplace, education, and advertising contexts. The V-axis extraction protocol applied to EEG could in principle reduce the calibration-data burden for affective inference systems — a feature that cuts both ways, since lower calibration cost lowers the deployment threshold for both clinical and surveillance settings. We do not release human EEG data; we release training configs, feature-side and analysis code, and extracted V-axis directions. Recommended deployment safeguards: informed consent, opt-in only, on-device inference, no longitudinal storage of raw EEG outside research contexts, and IRB review for any non-clinical inference application. Dual-use risk for the V-axis extraction itself is low: it operates on text and image features already public, and the EEG mapping requires per-cohort EEG access that is governed by the original dataset licences.

Implications for affective neuroscience. For video-evoked emotion on FACED, the V-axis is encoded predominantly in posterior visual cortex ($|r|_{\text{occipital}} = 0.21$ vs. $|r|_{\text{frontal}} = 0.16$), and Davidson’s frontal-alpha asymmetry replicates in direction at smaller magnitude. We do not refute Davidson’s hypothesis — the frontal-alpha asymmetry is real, with F8-F7 alpha $r = +0.0155$ and Fp2-Fp1 alpha $r = +0.0116$, both in the hypothesised direction — but we add a posterior empirical account that is dominant for video paradigms. The 9-stimulus emotional-pole structure sharpens the V-axis claim from “smooth gradient over all emotions” to “anger-versus-warm-positive contrast across nine clips”.

Per-subject V-axis adaptation does not transfer. The per-subject best-channel oracle reaches $|r| = 0.62$ (a +0.60 absolute headroom over the cohort-fixed top-8 of $|r| = 0.02$). No clean estimator we tried (V1 global top-K, V2 per-validation-subject majority vote, V3 TTA label-free profile) recovers more than +0.001 over the cohort top-K. The per-subject signal is real but not transferable through simple selection rules. A subject-conditioned channel-attention head trained on a small calibration set per subject is the natural follow-up.

Future work. Three lines: (i) per-subject V-axis adaptation closing the $|r| = 0.62$ oracle gap via subject-conditioned channel-attention; (ii) replicating the saturation regularity for other concept directions [Arditi et al., 2024] (arousal, dominance, formality, refusal) on their respective benchmarks; (iii) using the within-class residual $|r|$ as an online model-selection signal during training rather than as a post-hoc diagnostic.

S17 FACED Test Confusion Matrices

Figure 18 contrasts the best-single-checkpoint and 10-checkpoint-ensemble confusion matrices on the FACED 9-class test split.




Figure 18: FACED 9-class test confusion. Left: best single checkpoint (BACC 0.6755). Right: 10-checkpoint ensemble (BACC 0.6948). The ensemble’s gain is concentrated on the off-diagonal cells where single-seed disagreement is highest, consistent with the within-class residual mechanism of Section 5.

S18 Negative Results Catalogue

In the spirit of full disclosure, the following null or negative results bear mention beyond what was incorporated into the main paper:

1. **Theta-gamma cross-frequency coupling does not mediate the V-axis** (Appendix S9, Appendix S9): Tort modulation index $\rho(\text{PAC}, V_r) = +0.082$, $p = 0.667$.
2. **Mutual-information control is not significant** (Section 7): MI $p = 0.115$ vs. linear $p = 0.020$. The V-axis-EEG relationship is essentially linear.
3. **Cross-architecture V-axis convergence is pseudo-replicated** (Section 4): the 36-checkpoint $r = +0.885$ is driven by $\sim 2\text{--}3$ architecture clusters ($n_{\text{eff}} \approx 1.2$), is not significant within any single architecture, and sits at only the 93.5th percentile of a 1000-direction null ($p_{\text{one}} = 0.066$, $p_{\text{two}} = 0.099$). Within-class residual ($r = +0.74$, $p = 0.014$) is the statistically robust signal we retain.
4. **Per-subject V-axis adaptation does not transfer.** The per-subject best-channel oracle reaches $|r| = 0.62$, but no clean estimator transfers more than +0.001 over the cohort top-K.

5. **Cross-architecture ensembling does not add diversity.** Mixing $d \in \{4, 6, 8, 10\}$ checkpoints gives within-arch 0.6798 vs. cross-arch 0.6791 (n.s.).
6. **Wider ensemble pools do not help.** Adding $d \in \{8, 10\}$ depth variants beyond the 10-checkpoint $d=6$ pool did not exceed 0.6948 in exploratory runs (Appendix S13).
7. **Stimulus-aggregation re-ranking does not help.** Operating at the 28-stimulus level instead of trial level gives BACC=1.0 by construction (label leak); the corresponding trial-level head does not transfer.
8. **Test-time augmentation at K=5 does not improve the ensemble** (0.6753 $\rightarrow$ 0.6720). Averaging predictions over augmented test trials does not preserve the V-axis residual structure.
9. **Toxicity: recipe scope.** The 9-prompt PCA recipe attains AUC 0.59 on Jigsaw toxic comments, below every one of the 20 in-scope concepts (each AUC > 0.65) in our library (Appendix S4); toxicity appears to be encoded in a higher-rank subspace than the late-layer PC1.
10. **Arousal asymmetry: vision yes, text and brain no.** OASIS arousal $r = 0.803$ in vision; $\leq 13\%$ NRC arousal recovery in text; $r \in [0.18, 0.41]$ brain alignment across LLMs (Appendix S6).
11. **LLM-content of KD is irrelevant at this scale.** Replacing the 9-D LLM-derived class prototypes with random orthonormal 9-D directions leaves BACC statistically unchanged ($\Delta = -0.004$, paired $p = 0.37$; if anything the LLM prototypes are slightly below random). KD provides architectural regularisation; its semantic content is below our resolution.

S19 Reproducibility

Release commitment. Code, configs, model checkpoints, and figure-generation scripts will be released upon acceptance (no submission-time anonymous URL is provided). The release will include: V-axis extraction scripts for the 14 LLMs in the per-LLM table; the EMOD $d=6$ training pipeline with the exact augmentation, KD, and ensemble-construction code; the 10-checkpoint 0.6948-BACC ensemble checkpoints; and the figure-generation scripts for every landmark and neuro figure in the paper.

Run-level provenance. Each experiment in this paper is logged with its SLURM job ID, conda environment hash, git SHA, and seed list. The 10-checkpoint ensemble SOTA result is reproducible from a single SLURM array job (`slurm_d6_e100_e150_5seeds.sh`, ~ 10 GPU-hours per seed on V100).

S20 Resource Estimate

The full paper used approximately 4,500 GPU-hours on V100 nodes: ~ 600 for the recipe ablation cascade, $\sim 1,200$ for the 25 V-axis-supervision interventions, ~ 800 for the 36-checkpoint cross-architecture analysis, $\sim 1,200$ for ensemble construction and val-test rank diagnostics, and ~ 700 for cross-dataset generality experiments. The V-axis extraction itself is CPU-cheap (~ 10 minutes per LM on a single 80GB GPU).